

SQL OVERVIEW

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 4

SQL

2

- SQL = Structured Query Language
- Original language was “SEQUEL”
 - ▣ IBM’s System R project (early 1970’s)
 - ▣ “Structured English Query Language”
- Caught on very rapidly
 - ▣ Simple, declarative language for writing queries
 - ▣ Also includes many other features
- Standardized by ANSI/ISO
 - ▣ SQL-86, SQL-89, SQL-92, SQL:1999, SQL:2003, SQL:2008, SQL:2011
 - ▣ Most implementations *loosely* follow the standards (plenty of portability issues)

SQL Features

3

- Data Definition Language (DDL)
 - ▣ Specify relation schemas (attributes, domains)
 - ▣ Specify a variety of integrity constraints
 - ▣ Access constraints on data
 - ▣ Indexes and other storage “hints” for performance
- Data Manipulation Language (DML)
 - ▣ Generally based on relational algebra
 - ▣ Supports querying, inserting, updating, deleting data
 - ▣ Very sophisticated features for multi-table queries
- Other useful tools
 - ▣ Defining views, transactions, etc.

SQL Basics

4

- SQL language is case-insensitive
 - ▣ both keywords and identifiers (for the most part)
- SQL statements end with a semicolon
- SQL comments have two forms:
 - ▣ Single-line comments start with two dashes
 - `-- This is a SQL comment.`
 - ▣ Block comments follow C style
 - `/*`
 - `* This is a block comment in SQL.`
 - `*/`

SQL Databases

5

- SQL relations are contained within a database
 - ▣ Each application usually works against its own database
 - ▣ Several applications may share the same database, too
- An example from MySQL:
 `CREATE DATABASE bank;`
 `USE bank;`
 - ▣ Creates a new, empty database called **bank**
 - ▣ **USE** statement makes **bank** the “default” database for the current connection
 - ▣ DDL and DML operations will be evaluated in the context of the connection’s default database

Creating a SQL Table

6

- In SQL, relations are called “tables”
 - ▣ Not *exactly* like relational model “relations” anyway

- Syntax:

```
CREATE TABLE t (  
    attr1 domain1,  
    attr2 domain2,  
    ... ,  
    attrN domainN  
);
```

- ▣ **t** is name of relation (table)
- ▣ **attr1, ...** are names of attributes (columns)
- ▣ **domain1, ...** are domains (types) of attributes

SQL Names

7

- Tables, columns, etc. require names
- Rules on valid names can vary dramatically across implementations
- Good, portable rules:
 - ▣ First character should be alphabetical
 - ▣ Remaining characters should be alphanumeric or underscore ‘_’
 - ▣ Use same the case in DML that you use in DDL

SQL Attribute Domains

8

- Some standard SQL domain types:

CHAR (N)

- A character field, fixed at N characters wide
- Short for **CHARACTER (N)**

VARCHAR (N)

- A variable-width character field, with maximum length N
- Short for **CHARACTER VARYING (N)**

INT

- A signed integer field (typically 32 bits)
- Short for **INTEGER**
- Also **TINYINT**, **SMALLINT**, **BIGINT**, etc.
- Also unsigned variants
 - Non-standard, only supported by some vendors

CHAR vs. VARCHAR

9

- Both **CHAR** and **VARCHAR** have a size limit
- **CHAR** is a fixed-length character field
 - ▣ Can store shorter strings, but storage layer pads out the value to the full size
- **VARCHAR** is a variable-length character field
 - ▣ Storage layer doesn't pad out shorter strings
 - ▣ String's length must also be stored for each value
- Use **CHAR** when all values are approximately (or *exactly*) the same length
- Use **VARCHAR** when values can be very different lengths

SQL Attribute Domains (2)

10

□ More standard SQL domain types:

NUMERIC (P , D)

- A fixed-point number with user-specified precision
- P total digits; D digits to right of decimal place
- Can exactly store numbers

DOUBLE PRECISION

- A double-precision floating-point value
- An approximation! Don't use for money! 😊
- **REAL** is sometimes a synonym

FLOAT (N)

- A floating-point value with at least N bits of precision

SQL Attribute Domains (3)

11

- Other useful attribute domains, too:
DATE, TIME, TIMESTAMP
 - For storing temporal data
- Large binary/text data fields
BLOB, CLOB, TEXT
 - Binary Large Objects, Character Large Objects
 - Large text fields
 - **CHAR, VARCHAR** tend to be very limited in size
- Other specialized types
 - ▣ Enumerations, geometric or spatial data types, etc.
 - ▣ User-defined data types

Choosing the Right Type

12

- Need to think carefully about what type makes most sense for your data values
- Example: storing ZIP codes
 - ▣ US postal codes for mail routing
 - ▣ 5 digits, e.g. 91125 for Caltech
- Does **INTEGER** make sense?
- **Problem 1:** Some ZIP codes have leading zeroes!
 - ▣ Many east-coast ZIP codes start with 0.
 - ▣ Numeric types won't include leading zeros.
- **Problem 2:** US mail also uses ZIP+4 expanded ZIP codes
 - ▣ e.g. 91125-8000
- **Problem 3:** Many foreign countries use non-numeric values

Choosing the Right Type (2)

13

- Better choice for ZIP codes?
 - ▣ A **CHAR** or **VARCHAR** column makes much more sense
- For example:
 - ▣ **CHAR (5)** or **CHAR (9)** for US-only postal codes
 - ▣ **VARCHAR (20)** for US + international postal codes
- Another example: monetary amounts
 - ▣ Floating-point representations cannot exactly represent all values
 - e.g. 0.1 is an infinitely-repeating binary decimal value
 - ▣ Use **NUMERIC** to represent monetary values

Example SQL Schema

14

□ Creating the account relation:

```
CREATE TABLE account (  
    acct_id      CHAR(10) ,  
    branch_name  CHAR(20) ,  
    balance      NUMERIC(12, 2)  
);
```

- ▣ Account IDs can't be more than 10 chars
- ▣ Branch names can't be more than 20 chars
- ▣ Balances can have 10 digits left of decimal, 2 digits right of decimal
 - Fixed-point, exact precision representation of balances

Inserting Rows

15

- Tables are initially empty
- Use **INSERT** statement to add rows

```
INSERT INTO account
VALUES ('A-301', 'New York', 350);
INSERT INTO account
VALUES ('A-307', 'Seattle', 275);
...
```

- ▣ String values are single-quoted
- ▣ (In SQL, double-quoted strings refer to column names)
- ▣ Values appear in same order as table's attributes

Inserting Rows (2)

16

- Can specify which attributes in **INSERT**

```
INSERT INTO account (acct_id, branch_name, balance)
VALUES ('A-301', 'New York', 350);
```

- ▣ Can list attributes in a different order

- ▣ Can exclude attributes that have a default value

- Problem: We can add multiple accounts with same account ID!

```
INSERT INTO account
VALUES ('A-350', 'Seattle', 800);

INSERT INTO account
VALUES ('A-350', 'Los Angeles', 195);
```


Primary Key Constraints

17

- The **CREATE TABLE** syntax also allows integrity constraints to be specified
 - ▣ Are often specified after all attributes are listed
- Primary key constraint:

```
CREATE TABLE account (  
    acct_id          CHAR(10) ,  
    branch_name     CHAR(20) ,  
    balance          NUMERIC(12, 2) ,  
  
    PRIMARY KEY (acct_id)  
);
```

- ▣ Database won't allow two rows with same account ID

Primary Key Constraints (2)

18

- A primary key can have multiple attributes

```
CREATE TABLE depositor (  
    customer_name  VARCHAR(30) ,  
    acct_id        CHAR(10) ,  
    PRIMARY KEY (customer_name, acct_id)  
);
```

- ▣ Necessary because SQL tables are multisets
- A table cannot have multiple primary keys
 - ▣ (obvious)
- *Many* other kinds of constraints too
 - ▣ Will cover in future lectures!

Removing Rows, Tables, etc.

19

- Can delete rows with **DELETE** command

- ▣ Delete bank account with ID A-307:

- ```
DELETE FROM account WHERE acct_id = 'A-307';
```

- ▣ Delete all bank accounts:

- ```
DELETE FROM account;
```

- Can drop tables and databases:

- ▣ Remove account table:

- ```
DROP TABLE account;
```

- ▣ Remove an entire database, including all tables!

- ```
DROP DATABASE bank;
```

Issuing SQL Queries

20

- SQL queries use the **SELECT** statement
- Very central part of SQL language
 - ▣ Concepts appear in all DML commands
- General form is:

```
SELECT  $A_1, A_2, \dots$   
      FROM  $r_1, r_2, \dots$   
      WHERE  $P;$ 
```

- ▣ r_i are the relations (tables)
- ▣ A_i are attributes (columns)
- ▣ P is the selection predicate

SELECT Operations

21

- **SELECT** A_1, A_2, \dots
 - ▣ Corresponds to a relational algebra project operation
$$\Pi_{A_1, A_2, \dots}(\dots)$$
 - ▣ Some books call σ “restrict” because of this name mismatch
- **FROM** r_1, r_2, \dots
 - ▣ Corresponds to Cartesian product of relations r_1, r_2, \dots
$$r_1 \times r_2 \times \dots$$

SELECT Operations (2)

22

□ WHERE P

- Corresponds to a selection operation

$\sigma_P(\dots)$

- Can be omitted. When left off, $P = \text{true}$

□ Assembling it all:

**SELECT A_1, A_2, \dots FROM r_1, r_2, \dots
WHERE P;**

- Equivalent to: $\Pi_{A_1, A_2, \dots}(\sigma_P(r_1 \times r_2 \times \dots))$

SQL and Duplicates

23

- Biggest difference between relational algebra and SQL is use of multisets
 - ▣ In SQL, relations are multisets of tuples, not sets
- Biggest reason is practical:
 - ▣ Removing duplicate tuples is time consuming!
- Must revise definitions of relational algebra operations to handle duplicates
 - ▣ Mainly affects set-operations: \cup , \cap , $-$
 - ▣ (Book explores this topic in depth)
- SQL provides ways to exclude duplicates for all operations

Example Queries

24

“Find all branches with at least one bank account.”

```
SELECT branch_name  
FROM account;
```

- Equivalent to typing:

```
SELECT ALL branch_name  
FROM account;
```

+	-----	+
	branch_name	
+	-----	+
	New York	
	Seattle	
	Los Angeles	
	New York	
	Los Angeles	
+	-----	+

- To eliminate duplicates:

```
SELECT DISTINCT branch_name  
FROM account;
```

+	-----	+
	branch_name	
+	-----	+
	New York	
	Seattle	
	Los Angeles	
+	-----	+

Selecting Specific Attributes

25

- Can specify one or more attributes to appear in result

“Find ID and balance of all bank accounts.”

```
SELECT acct_id, balance
FROM account;
```

acct_id	balance
A-301	350.00
A-307	275.00
A-318	550.00
A-319	80.00
A-322	275.00

- Can also specify * to mean “all attributes”

```
SELECT * FROM account;
```

- ▣ Returns all details of all accounts.

acct_id	branch_name	balance
A-301	New York	350.00
A-307	Seattle	275.00
A-318	Los Angeles	550.00
A-319	New York	80.00
A-322	Los Angeles	275.00

Computing Results

26

- The **SELECT** clause is a *generalized projection* operation

- ▣ Can compute results based on attributes

```
SELECT cred_id, credit_limit - balance
FROM credit_account;
```

- ▣ Computed values don't have a (standard) name!

- Many DBMSes name the 2nd column "credit_limit - balance"

- Can also name (or rename) values

```
SELECT cred_id,
       credit_limit - balance AS available_credit
FROM credit_account;
```

WHERE Clause

27

- The **WHERE** clause specifies a selection predicate
 - ▣ Can use comparison operators:
 - =, <> equals, not-equals (!= also usually supported)
 - <, <= less than, less or equal
 - >, >= greater than, greater or equal
 - ▣ Can refer to any attribute in **FROM** clause
 - ▣ Can include arithmetic expressions in comparisons

WHERE Examples

28

“Find IDs and balances of all accounts in the Los Angeles branch.”

```
SELECT acct_id, balance FROM account
WHERE branch_name = 'Los Angeles';
```

acct_id	balance
A-318	550.00
A-322	275.00

“Retrieve all details of bank accounts with a balance less than \$300.”

```
SELECT * FROM account
WHERE balance < 300;
```

acct_id	branch_name	balance
A-307	Seattle	275.00
A-319	New York	80.00
A-322	Los Angeles	275.00

Larger Predicates

29

- Can use **AND**, **OR**, **NOT** in **WHERE** clause

```
SELECT acct_id, balance FROM account
  WHERE branch_name = 'Los Angeles' AND
         balance < 300;
```

```
SELECT * FROM account
  WHERE balance >= 250 AND balance <= 400;
```

- SQL also has **BETWEEN** and **NOT BETWEEN** syntax

```
SELECT * FROM account
  WHERE balance BETWEEN 250 AND 400;
```

- ▣ Note that **BETWEEN** includes interval endpoints!

String Comparisons

30

- String values can be compared
 - ▣ Lexicographic comparisons
 - ▣ Default is often to ignore case!

```
SELECT 'HELLO' = 'hello'; -- Evaluates to true
```
- Can also do pattern matching with **LIKE** expression
string_attr LIKE pattern
 - ▣ **pattern** is a string literal enclosed in single-quotes
 - % (percent) matches a substring
 - _ (underscore) matches a single character
 - Can escape % or _ with a backslash \

String-Matching Example

31

“Find all accounts at branches with ‘le’ somewhere in the name.”

▣ Why? I don’t know...

```
SELECT * FROM account
WHERE branch_name LIKE '%le%';
```

acct_id	branch_name	balance
A-307	Seattle	275.00
A-318	Los Angeles	550.00
A-322	Los Angeles	275.00

String Operations

32

- Regular-expression matching is also part of the SQL standard (SQL:1999)
- String-matching operations tend to be expensive
 - ▣ Especially patterns with a leading wildcard, e.g. ' %abc '
- Try to avoid heavy reliance on pattern-matching

- If string searching is required, try to pre-digest text and generate search indexes
 - ▣ Some databases provide “full-text search” capabilities, but such features are vendor-specific!

FROM Clause

33

- Can specify one or more tables in **FROM** clause
 - If multiple tables:
 - ▣ Select/project against Cartesian product of relations
 - Produces a row for every combination
 - of input tuples.
- ```
SELECT * FROM borrower, loan;
```

| cust_name | loan_id | loan_id | branch_name   | amount  |
|-----------|---------|---------|---------------|---------|
| Anderson  | L-437   | L-419   | Seattle       | 2900.00 |
| Jackson   | L-419   | L-419   | Seattle       | 2900.00 |
| Lewis     | L-421   | L-419   | Seattle       | 2900.00 |
| Smith     | L-445   | L-419   | Seattle       | 2900.00 |
| Anderson  | L-437   | L-421   | San Francisco | 7500.00 |
| Jackson   | L-419   | L-421   | San Francisco | 7500.00 |
| Lewis     | L-421   | L-421   | San Francisco | 7500.00 |
| ...       |         |         |               |         |

# FROM Clause (2)

34

- If tables have overlapping attributes, use `tbl_name.attr_name` to distinguish

```
SELECT * FROM borrower, loan
WHERE borrower.loan_id = loan.loan_id;
```

| cust_name | loan_id | loan_id | branch_name   | amount  |
|-----------|---------|---------|---------------|---------|
| Jackson   | L-419   | L-419   | Seattle       | 2900.00 |
| Lewis     | L-421   | L-421   | San Francisco | 7500.00 |
| Anderson  | L-437   | L-437   | Las Vegas     | 4300.00 |
| Smith     | L-445   | L-445   | Los Angeles   | 2000.00 |

- All columns can be referred to by `tbl_name.attr_name`
- This kind of query is called an equijoin
- Databases optimize equijoin queries very effectively.

# SQL and Joins

35

- SQL provides several different options for performing joins across multiple tables
- This form is the most basic usage
  - ▣ Was in earliest versions of SQL
  - ▣ Doesn't provide natural joins
  - ▣ Can't do outer joins either
- Will cover other forms of SQL join syntax soon...

# Renaming Tables

36

- Can specify alternate names in **FROM** clause too

- ▣ Write: **table AS name**

- ▣ (The **AS** is optional, but it's clearer to leave it in.)

- Previous example:

“Find the loan with the largest amount.”

- ▣ Started by finding loans that have an amount smaller than some other loan's amount

- ▣ Used Cartesian product and rename operation

```
SELECT DISTINCT loan.loan_id
FROM loan, loan AS test
WHERE loan.amount < test.amount;
```

```
+-----+
| loan_id |
+-----+
| L-445 |
| L-419 |
| L-437 |
+-----+
```

# Renaming Tables (2)

37

- When a table is renamed in **FROM** clause, can use the new name in both **SELECT** and **WHERE** clauses
- Useful for long table names! 😊

```
SELECT c.cust_name, l.amount
 FROM customer AS c, borrower AS b,
 loan AS l
 WHERE c.cust_name = b.cust_name AND
 b.loan_id = l.loan_id;
```

# Set Operations

38

- SQL also provides set operations, like relational algebra
- Operations take two relations and produce an output relation
- Set-union:  
*select<sub>1</sub> UNION select<sub>2</sub> ;*
- Set-intersection:  
*select<sub>1</sub> INTERSECT select<sub>2</sub> ;*
- Set-difference:  
*select<sub>1</sub> EXCEPT select<sub>2</sub> ;*
- **Note:** *select<sub>i</sub>* are complete **SELECT** statements!

# Set-Operation Examples

39

- Find customers with an account or a loan:  

```
SELECT cust_name FROM depositor UNION
SELECT cust_name FROM borrower;
```

  - ▣ Database automatically eliminates duplicates
- Find customers with an account but not a loan:  

```
SELECT cust_name FROM depositor EXCEPT
SELECT cust_name FROM borrower;
```

  - ▣ Can also put parentheses around **SELECT** clauses for readability  

```
(SELECT cust_name FROM depositor)
EXCEPT
(SELECT cust_name FROM borrower);
```

# Set Operations and Duplicates

40

- By default, SQL set-operations eliminate duplicate tuples
  - ▣ Opposite to default behavior of **SELECT**!
- Can keep duplicate tuples by appending **ALL** to set operation:

*select<sub>1</sub>* UNION ALL *select<sub>2</sub>* ;

*select<sub>1</sub>* INTERSECT ALL *select<sub>2</sub>* ;

*select<sub>1</sub>* EXCEPT ALL *select<sub>2</sub>* ;



# How Many Duplicates?

41

- Need to define behavior of “set operations” on multisets
- Given two multiset relations  $r_1$  and  $r_2$ 
  - ▣  $r_1$  and  $r_2$  have same schema
  - ▣ Some tuple  $t$  appears  $c_1$  times in  $r_1$ , and  $c_2$  times in  $r_2$

$$\mathbf{r}_1 \cup_{\text{ALL}} \mathbf{r}_2$$

contains  $c_1 + c_2$  copies of  $t$

$$\mathbf{r}_1 \cap_{\text{ALL}} \mathbf{r}_2$$

contains  $\min(c_1, c_2)$  copies of  $t$

$$\mathbf{r}_1 -_{\text{ALL}} \mathbf{r}_2$$

contains  $\max(c_1 - c_2, 0)$  copies of  $t$

# Other Relational Operations

42

- Can actually update definitions of all relational operations to support multisets
- Necessary for using relational algebra to model execution plans
- Not terribly interesting though... 😊
- If you're curious, see book for details

# SQL Style Guidelines

43

- Follow good coding style in SQL!
- Some recommendations:
  - ▣ Use lowercase names for tables, columns, etc.
  - ▣ Put a descriptive comment above every table
  - ▣ Write all SQL keywords in uppercase
  - ▣ Follow standard indentation scheme
    - e.g. indent columns in table declarations by 2-4 spaces
  - ▣ Keep lines to 80 characters or less!
    - wrap lines in reasonable places
- **Note:** You will lose points for sloppy SQL.

# Next Time

44

- Sorting results
- Grouping and aggregate functions
- Nested queries and many more set operations
- How to update SQL databases