

DATABASE SCHEMA DESIGN

ENTITY-RELATIONSHIP MODEL

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 14

Designing Database Applications

2

- Database applications are large and complex
- A few of the many design areas:
 - ▣ Database schema (physical/logical/view)
 - ▣ Programs that access and update data
 - ▣ Security constraints for data access
- Also requires familiarity with the problem domain
 - ▣ Domain experts *must* help drive requirements

General Approach

3

- Collect user requirements
 - ▣ Information that needs to be represented
 - ▣ Operations to perform on that information
 - ▣ Several techniques for representing this info, e.g. UML
- Develop a conceptual schema of the database
 - ▣ A high-level representation of the database's structure and constraints
 - Physical *and* logical design issues are ignored at this stage
 - ▣ Follows a specific data model
 - ▣ Often represented graphically

Conceptual Schema

4

- Also need to create a specification of functional requirements
 - ▣ “What operations will be performed against the data?”
 - ▣ Updating data, adding data, deleting data, ...
- Designer can use functional requirements to verify the conceptual schema
 - ▣ Is each operation possible?
 - ▣ How complicated or involved is it?
 - ▣ Performance or scalability concerns?

Implementation Phases

5

- Once conceptual schema and functional requirements are verified:
 - ▣ Convert conceptual schema into an implementation data model
 - ▣ Want to have a simple mapping from conceptual model to implementation model
- Finally: any necessary physical design
 - ▣ Not always present!
 - ▣ Smaller applications have few physical design concerns
 - ▣ Larger systems usually need additional design and tuning (e.g. indexes, disk-level partitioning of data)

Importance of Design Phase

6

- Not all changes have the same impact!
- Physical-level changes have the least impact
 - ▣ (Thanks, relational model!)
 - ▣ Typically affect performance, scalability, reliability
 - ▣ Little to no change in functionality
- Logical-level changes are typically *much* bigger
 - ▣ Affects how to interact with the data...
 - ▣ Also affects what is even *possible* to do with the data
- Very important to spend time up front designing the database schema

Design Decisions

7

- Many different ways to represent data
- Must avoid two major problems:
 - ▣ Unnecessary redundancy
 - Redundant information wastes space
 - Greater potential for inconsistency!
 - Ideally: each fact appears in exactly one place
 - ▣ Incomplete representation
 - Schema must be able to fully represent all details and relationships required by the application

More Design Decisions

8

- Even with correct design, usually many other concerns
 - ▣ How easy/hard is it to access useful information? (e.g. reporting or summary info)
 - ▣ How hard is it to update the system?
 - ▣ Performance considerations?
 - ▣ Scalability considerations?
- Schema design requires a good balance between aesthetic and practical concerns
 - ▣ Frequently need to make compromises between conflicting design principles

Simple Design Example

9

- Purchase tracking database
 - ▣ Store details about product purchases by customers
 - ▣ Actual purchases tracked in database
- Can represent sales as relationships between customers and products
 - ▣ What if product price changes? Where to store product sale price? Will this affect other recent purchases?
 - ▣ What about giving discounts to some customers? May want to give different prices to different customers.
- Can also represent sales as separate entities
 - ▣ Gives much more flexibility for special pricing, etc.

The Entity-Relationship Model

10

- A very common model for schema design
 - ▣ Also written as “E-R model”
- Allows for specification of complex schemas in graphical form
- Basic concepts are simple, but can also represent very sophisticated abstractions
 - ▣ e.g. type hierarchies
- Can be mapped very easily to the relational model!
 - ▣ Simplifies implementation phase
 - ▣ Mapping process can be automated by design tools

Entities and Entity-Sets

11

- An entity is any “thing” that can be uniquely represented
 - e.g. a product, an employee, a software defect
 - ▣ Each entity has a set of attributes
 - ▣ Entities are uniquely identified by some set of attributes
- An entity-set is a named collection of entities of the same type, with the same attributes
 - ▣ Can have multiple entity-sets with same entity type, representing different (possibly overlapping) sets

Entities and Entity-Sets (2)

12

- An entity has a set of attributes
 - ▣ Each attribute has a name and domain
 - ▣ Each attribute also has a corresponding value
- Entity-sets also specify a set of attributes
 - ▣ Every entity in the entity-set has the same set of attributes
 - ▣ Every entity in the entity-set has its own value for each attribute


Diagramming an Entity-Set

13

Example: a *customer* entity-set

- Attributes:

- *cust_id*
- *name*
- *street_address*
- *city*



- Entity-set is denoted by a box
- Name of entity-set is given in the top part of box
- Attributes are listed in the lower part of the box

Relationships

14

- A relationship is an association between two or more entities
 - ▣ e.g. a bank loan, and the customer who owns it
- A relationship-set is a named collection of relationships of the same type
 - ▣ i.e. involving the same entities
- Formally, a relationship-set is a mathematical relation involving n entity-sets, $n \geq 2$
 - ▣ E_1, E_2, \dots, E_n are entity sets; e_1, e_2, \dots are entities in E_1, E_2, \dots
 - ▣ A relationship set R is a subset of:
 $\{ (e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n \}$
 - ▣ (e_1, e_2, \dots, e_n) is a specific relationship in R

Relationships (2)

15

- Entity-sets participate in relationship-sets
 - ▣ Specific entities participate in a relationship instance
- Example: bank loans
 - ▣ *customer* and *loan* are entity-sets
 - (555-55-5555, Jackson, Woodside) is a *customer* entity
 - (L-14, 1500) is a *loan* entity
 - ▣ *borrower* is a relationship-set
 - *customer* and *loan* participate in *borrower*
 - *borrower* contains a relationship instance that associates customer “Jackson” and loan “L-14”

Relationships and Roles

16

- An entity's role in a relationship is the function that the entity fills

Example: a *purchase* relationship between a *product* and a *customer*

- the product's role is that it was purchased
- the customer did the purchasing
- Roles are usually obvious, and therefore unspecified
 - Entities participating in relationships are distinct...
 - Names clearly indicate the roles of various entities...
 - In these cases, roles are left unstated.

Relationships and Roles (2)

17

- Sometimes the roles of entities are *not* obvious
 - ▣ Situations where entity-sets in a relationship-set are *not* distinct

Example: a relationship-set named *works_for*, specifying employee/manager assignments

- ▣ Relationship involves two entities, and both are *employee* entities
- Roles are given names to distinguish entities
 - ▣ The relationship is a set of entities ordered by role:
(*manager*, *worker*)
 - ▣ First entity's role is named *manager*
 - ▣ Second entity's role is named *worker*

Relationships and Attributes

18

- Relationships can also have attributes!
 - ▣ Called descriptive attributes
 - ▣ They describe a particular relationship
 - ▣ They *do not* identify the relationship!
- Example:
 - ▣ The relationship between a software defect and an employee can have a *date_assigned* attribute
- Note: this distinction between entity attributes and relationship attributes is not made by relational model
 - ▣ Entity-relationship model is a higher level of abstraction than the relational model

Relationships and Attributes (2)

19

- Specific relationships are identified *only* by the participating entities
 - ▣ ...not by any relationship attributes!
 - ▣ Different relationships are allowed to have the same value for a descriptive attribute
 - ▣ (This is why entities in an entity-set must be uniquely identifiable.)
- Given:
 - ▣ Entity-sets A and B , both participating in a relationship-set R
- Specific entities $a \in A$ and $b \in B$ can only have one relationship instance in R
 - ▣ Otherwise, we would require more than just the participating entities to uniquely identify relationships

Degree of Relationship Set


20

- Most relationships in a schema are binary
 - ▣ Two entities are involved in the relationship
- Sometimes there are ternary relationships
 - ▣ Three entities are involved
 - ▣ Far less common, but still useful at times
- The number of entity-sets that participate in a relationship-set is called its degree
 - ▣ Binary relationship: $\text{degree} = 2$
 - ▣ Ternary relationship: $\text{degree} = 3$

Diagramming a Relationship-Set

21

Example: the *borrower* relationship-set between the *customer* and *loan* entity-sets



- Relationship-set is a diamond
 - ▣ Connected to participating entity-sets by solid lines
- Relationship-set can have descriptive attributes
 - ▣ Listed in another box, connected with a dotted-line

Attribute Structure

22

- Each attribute has a domain or value set
 - ▣ Values come from that domain or value set
- Simple attributes are atomic – they have no subparts
 - ▣ e.g. *amount* attribute is a single numeric value
- Composite attributes have subparts
 - ▣ Can refer to composite attribute as a whole
 - ▣ Can also refer to subparts individually
 - ▣ e.g. *address* attribute, composed of *street*, *city*, *state*, *postal_code* attributes

Attribute Cardinality

23

- Single-valued attributes only store one value
 - ▣ e.g. a *customer*'s *cust_id* only has one value
- Multi-valued attributes store zero or more values
 - ▣ e.g. a *customer* can have multiple *phone_number* values
 - ▣ A multi-valued attribute stores a set of values, not a multiset
 - ▣ Different *customer* entities can have different sets of phone numbers
 - ▣ Lower and upper bounds can be specified too
 - Can set upper bound on *phone_number* to 2

Attribute Source

24

- Base attributes (aka source attributes) are stored in the database
 - ▣ e.g. the *birth_date* of a *customer* entity
- Derived attributes are computed from other attributes
 - ▣ e.g. the *age* of a *customer* entity would be computed from their *birth_date*

Diagramming Attributes

25

- Example: Extend customers with more detailed info
- Composite attributes are shown as a hierarchy of values
 - ▣ Indented values are components of the higher-level value
 - ▣ e.g. *name* is comprised of *first_name*, *middle_initial*, and *last_name*

<i>customer</i>
<i>cust_id</i>
<i>name</i> <ul style="list-style-type: none"><i>first_name</i><i>middle_initial</i><i>last_name</i>
<i>address</i> <ul style="list-style-type: none"><i>street</i><i>city</i><i>state</i><i>zip_code</i>

Diagramming Attributes (2)

26

- Example: Extend customers with more detailed info
- Multivalued attributes are enclosed with curly-braces
 - ▣ e.g. each customer can have zero or more phone numbers

<i>customer</i>
<i>cust_id</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>
{ <i>phone_number</i> }

Diagramming Attributes (3)

27

- Example: Extend customers with more detailed info
- Derived attributes are indicated by a trailing set of parentheses ()
 - ▣ e.g. each customer has a base attribute recording their date of birth
 - ▣ Also a derived attribute that reports the customer's current age

<i>customer</i>
<i>cust_id</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>
{ <i>phone_number</i> }
<i>birth_date</i>
<i>age</i> ()

Representing Constraints

28

- E-R model can represent different kinds of constraints
 - ▣ Mapping cardinalities
 - ▣ Key constraints in entity-sets
 - ▣ Participation constraints
- Allows more accurate modeling of application's data requirements
 - ▣ Constrain design so that schema can only represent valid information
- Enforcing constraints can impact performance...
 - ▣ Still ought to specify them in the design!
 - ▣ Can always leave out constraints at implementation time

Mapping Cardinalities

29

- Mapping cardinality represents:
 - “How many other entities can be associated with an entity, via a particular relationship set?”
- Example:
 - ▣ How many *customer* entities can the *borrower* relationship associate with a single *loan* entity?
 - ▣ How many *loans* can *borrower* relationship associate with a single *customer* entity?
 - ▣ Specific answer depends on what is being modeled
- Also known as the cardinality ratio
- Easiest to reason about with binary relationships

Mapping Cardinalities (2)


30

Given:

- Entity-sets A and B
- Binary relationship-set R associating A and B

One-to-one mapping (1:1)

- An entity in A is associated with *at most* one entity in B
- An entity in B is associated with *at most* one entity in A



Mapping Cardinalities (3)


31

One-to-many mapping (1:M)

- An entity in *A* is associated with *zero or more* entities in *B*
- An entity in *B* is associated with *at most one* entity in *A*

Many-to-one mapping (M:1)

- Opposite of one-to-many
- An entity in *A* is associated with *at most one* entity in *B*
- An entity in *B* is associated with *zero or more* entities in *A*




Mapping Cardinalities (4)

32

Many-to-many mapping

- An entity in A is associated with *zero or more* entities in B
- An entity in B is associated with *zero or more* entities in A



Mapping Cardinalities (5)

33

- Which mapping cardinality is most appropriate for a given relationship?
 - ▣ Answer depends on what you are trying to model!
 - ▣ Could just use many-to-many relationships everywhere, but that would be dumb.
- Goal:
 - ▣ Constrain the mapping cardinality to most accurately reflect what should be allowed
 - ▣ Database can enforce these constraints automatically
 - ▣ Good schema design reduces or eliminates the *possibility* of storing bad data

Example: *borrower* relationship between *customer* and *loan*

34

One-to-one mapping:

- ▣ Each customer can have only one loan
- ▣ Customers can't share loans
(e.g. with spouse or business partner)

One-to-many mapping:

- ▣ A customer can have multiple loans
- ▣ Customers still can't share loans

Many-to-one mapping:

- ▣ Each customer can have only one loan
- ▣ Customers can share loans

Many-to-many mapping:

- ▣ A customer can have multiple loans
- ▣ Customers can share loans too


Best choice for *borrower* :
many-to-many mapping

Handles real-world needs!

Diagramming Cardinalities

35


- In relationship-set diagrams:
 - ▣ an arrow towards an entity represents “one”
 - ▣ a simple line represents “many”
 - ▣ arrow is *always* towards the entity
- Many-to-many mapping between *customer* and *loan*:



Diagramming Cardinalities (2)

36

- One-to-many mapping between *customer* and *loan*:




- ▣ Each customer can have multiple loans
- ▣ A loan is owned by exactly one customer
 - (Actually, this is technically “at most one”. Participation constraints will allow us to say “exactly one.”)

Diagramming Cardinalities (3)

37

- One-to-one mapping between *customer* and *loan*:



- Each customer can have only one loan
- A loan is owned by exactly one customer

ENTITY-RELATIONSHIP MODEL II

Last Lecture

2

- Began to explore the Entity-Relationship Model
 - ▣ A visual representation of database schemas
 - ▣ Can represent entities and relationships
 - ▣ Can represent constraints in the schema
- Last time, left off with mapping cardinalities

Entity-Set Keys

3

- Entities in an entity-set must be uniquely distinguishable using their values
 - Entity-set: each entity is unique
- E-R model also includes the notion of keys:
 - Superkey: a set of one or more attributes that can uniquely identify an entity
 - Candidate key: a *minimal* superkey
 - Primary key: a candidate key chosen by DB designer as the primary means of accessing entities
- Keys are a property of the entity-set
 - They apply to *all* entities in the entity-set

Choosing Candidate Keys

4

- Candidate keys constrain the values of the key attributes
 - ▣ No two entities can have the same values for those attributes
 - ▣ Need to ensure that database can actually represent all expected circumstances
- Simple example: *customer* entity-set
 - ▣ Using customer name as a candidate key is bad design: different customers can have the same name

Choosing Primary Keys

5

- An entity-set may have multiple candidate keys
- The primary key is the candidate key most often used to reference entities in the set
 - ▣ In logical/physical design, primary key values will be used to represent relationships
 - ▣ External systems may also use primary key values to reference entities in the database
- The primary key attributes should never change!
 - ▣ If ever, it should be *extremely* rare.

Choosing Keys: Performance


6

- Large, complicated, or multiple-attribute keys are generally slower
 - ▣ Use smaller, single-attribute keys
 - (You can always generate them...)
 - ▣ Use faster, fixed-size types
 - e.g. **INT** or **BIGINT**
- Especially true for primary keys!
 - ▣ Values used in both database and in access code
 - ▣ Use something small and simple, if possible


Diagramming Primary Keys

7

- In an entity-set diagram, all attributes in the primary key have an underlined name



- Another example: a geocache *location* entity-set



Keys and Relationship-Sets

8

- Need to be able to distinguish between individual relationships in a relationship-set as well
 - ▣ Relationships aren't distinguished by their descriptive attributes
 - ▣ (They might not even have descriptive attributes)
- Relationships are identified by the *entities* participating in the relationship
 - ▣ Specific relationship instances are uniquely identified by the primary keys of the participating entities

Keys and Relationship-Sets (2)

9

- Given:
 - ▣ R is a relationship-set with no descriptive attributes
 - ▣ Entity-sets E_1, E_2, \dots, E_n participate in R
 - ▣ $primary_key(E_i)$ denotes set of attributes in E_i that represent the primary key of E_i
- A relationship instance in R is identified by $primary_key(E_1) \cup primary_key(E_2) \cup \dots \cup primary_key(E_n)$
 - ▣ This is a superkey
 - ▣ Is it a candidate key?
 - Depends on the *mapping cardinality* of the relationship set!

Keys and Relationship-Sets (3)

10

- If R also has descriptive attributes $\{a_1, a_2, \dots\}$, a relationship instance is described by:
$$\text{primary_key}(E_1) \cup \text{primary_key}(E_2) \cup \dots \cup \text{primary_key}(E_n) \cup \{a_1, a_2, \dots\}$$
- ▣ Not a minimal superkey!
- ▣ By definition, there can only be one relationship between $\{E_1, E_2, \dots, E_n\}$ in the relationship-set
 - i.e. the descriptive attributes do not identify specific relationships!
- Thus, just as before, this is also a superkey:
$$\text{primary_key}(E_1) \cup \text{primary_key}(E_2) \cup \dots \cup \text{primary_key}(E_n)$$

Relationship-Set Primary Keys


11

- What is the primary key for a binary relationship-set?
 - ▣ Must also be a candidate key
 - ▣ Depends on the mapping cardinalities
- Relationship-set R , involving entity-sets A and B
 - ▣ If mapping is many-to-many, primary key is:
$$\text{primary_key}(A) \cup \text{primary_key}(B)$$
 - ▣ Any given entity's primary-key values can appear multiple times in R
 - ▣ We need both entity-sets' primary key attributes to uniquely identify relationship instances

Relationship-Set Primary Keys (2)

12


- Relationship-set R , involving entity-sets A and B
 - ▣ Individual relationships are described by $primary_key(A) \cup primary_key(B)$
- If mapping is one-to-many:
 - ▣ Entities in B associated with *at most* one entity in A
 - ▣ A given $primary_key(A)$ value can appear in multiple relationships
 - ▣ Each value of $primary_key(B)$ can appear only once
 - ▣ Relationships in R are uniquely identified by $primary_key(B)$
 - ▣ $primary_key(B)$ is primary key of relationship-set



Relationship-Set Primary Keys (3)

13

- Relationship-set R , involving entity-sets A and B
- Many-to-one is exactly the opposite of one-to-many
 - ▣ $primary_key(A)$ uniquely identifies relationships in R



Relationship-Set Primary Keys (4)


14

- Relationship-set R , involving entity-sets A and B
- If mapping is one-to-one:
 - ▣ Entities in A associated with *at most* one entity in B
 - ▣ Entities in B associated with *at most* one entity in A
 - ▣ Each entity's key-value can appear only once in R
 - ▣ Either entity-set's primary key can be primary key of R
- For one-to-one mapping, $primary_key(A)$ and $primary_key(B)$ are both candidate keys
 - ▣ Make sure to enforce both candidate keys in the implementation schema!

Example

15

- What is the primary key for *borrower* ?



- *borrower* is a many-to-many mapping
 - ▣ Relationship instances are described by (*cust_id*, *loan_id*, *access_date*)
 - ▣ Primary key for relationship-set is (*cust_id*, *loan_id*)

Participation Constraints

16

- Given entity-set E , relationship-set R
 - ▣ How many entities in E participate in R ?
 - ▣ In other words, what is minimum number of relationships that each entity in E *must* participate in?
- If every entity in E participates in at least one relationship in R , then:
 - ▣ E 's participation in R is total
- If only some entities in E participate in relationships in R , then:
 - ▣ E 's participation in R is partial

Participation Constraints (2)


17

- Example: *borrower* relationship between *customer* and *loan*
- A customer might not have a bank loan
 - ▣ Could have a bank account instead
 - ▣ Could be a new customer
 - ▣ Participation of *customer* in *borrower* is partial
- Every loan definitely has at least one customer
 - ▣ Doesn't make any sense not to!
 - ▣ Participation of *loan* in *borrower* is total

Diagramming Participation

18


- Can indicate participation constraints in entity-relationship diagrams
 - ▣ Partial participation shown with a single line
 - ▣ Total participation shown with a double line



Numerical Constraints

19


- Can also state numerical participation constraints
 - ▣ Specifies how many different relationship instances each entity in the entity-set can participate in
 - ▣ Indicated on link between entity and relationship
- Form: lower..upper
 - ▣ * means “unlimited”
 - ▣ 1..* = one or more
 - ▣ 0..3 = between zero and three, inclusive
 - ▣ etc.



Numerical Constraints (2)

20

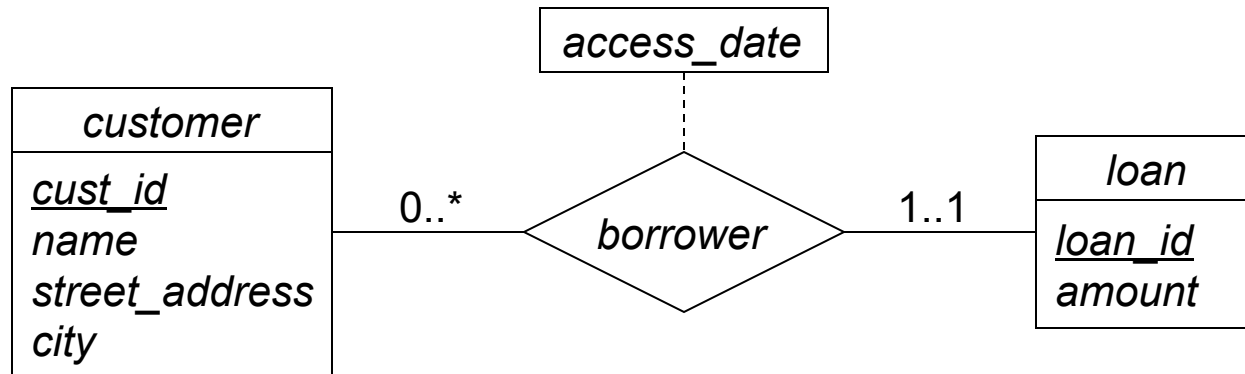
- Can also state mapping constraints with numerical participation constraints
- Total participation:
 - ▣ Lower bound at least 1
- Partial participation:
 - ▣ Lower bound is 0



Numerical Constraint Example

21

□ What does this mean?




- Each *customer* entity may participate in zero or more relationships in this relationship-set
 - A customer can have zero or more loans.
- Each *loan* entity must participate in exactly one relationship (no more, no less) in this relationship-set
 - Each loan must be owned by exactly one customer.

Numerical Constraint Example (2)

22

- What is the mapping cardinality of *borrower* ?



- From last slide:
 - A *customer* can have zero or more *loans*
 - Each *loan* must be owned by exactly one *customer*.
- This is a one-to-many mapping from *customer* to *loan*

Diagramming Roles

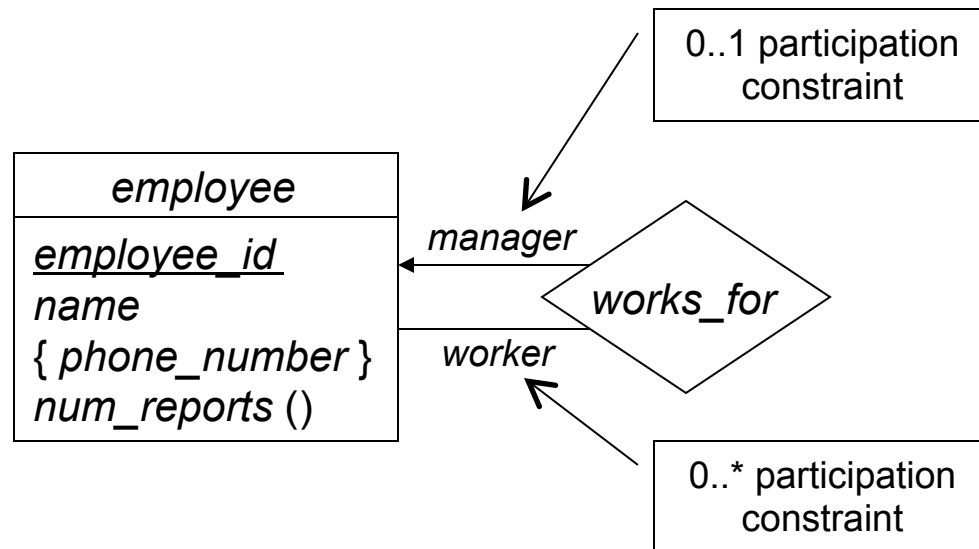
23

- Entities have roles in relationships
 - An entity's role indicates the entity's function in the relationship
 - e.g. role of customer in *borrower* relationship-set is that they own the loan
- Sometimes roles are ambiguous
 - e.g. when the same kind of entity is involved in a relationship multiple times
- Example: *works_for* relationship
 - Relationship is between two *employee* entities
 - One is the *manager*; the other is the *worker*

Diagramming Roles (2)

24

- If roles need to be indicated, put labels on the lines connecting entity to relationship



- *works_for* relationship-set is one-to-many from managers to workers

Weak Entity-Sets

25

- Sometimes an entity-set doesn't have distinguishing attributes
 - ▣ Can't define a primary key for the entity-set!
 - ▣ Called a weak entity-set
- Example:
 - ▣ Checking accounts have a unique account number
 - ▣ Checks have a check number
 - Unique for a given account, but not across all accounts!
 - Number only makes sense in context of a particular account
 - ▣ Want to store check transactions in the database

Weak Entity-Sets (2)

26

- Weak entity-sets *must* be associated with another (strong) entity-set
 - ▣ Called the identifying entity-set, or owner entity-set
 - ▣ The identifying entity-set owns the weak entity-set
 - ▣ Association called the identifying relationship
- Every weak entity *must* be associated with an identifying entity
 - ▣ Weak entity's participation in relationship-set is total
 - ▣ The weak entity-set is existence dependent on the identifying entity-set
 - ▣ If the identifying entity is removed, its weak entities should also cease to exist
 - ▣ (*this is where cascade-deletes may be appropriate...*)

Weak Entity-Set Keys

27

- Weak entity-sets don't have a primary key
 - Still need to distinguish between weak entities associated with a particular strong entity
- Weak entities have a discriminator
 - A set of attributes that distinguishes between weak entities associated with a strong entity
 - Also known as a partial key
- Checking account example:
 - The check number is the discriminator for check transactions

Weak Entity-Set Keys (2)


28

- Using discriminator, can define a primary key for weak entity-sets
- For a weak entity-set W , and an identifying entity-set S , primary key of W is:
$$\text{primary_key}(S) \cup \text{discriminator}(W)$$
- Checking account example:
 - *account_number* is primary key for checking accounts
 - *check_number* is discriminator (partial key) for checks
 - Primary key for check transactions would be $(\text{account_number}, \text{check_number})$

Diagramming Weak Entity-Sets

29


- Weak entity-sets drawn similarly to strong entity-sets
 - ▣ Difference: discriminator attributes are underlined with a dashed underline
- Identifying relationship to the owning entity-set is indicated with a double diamond
 - ▣ One-to-many mapping
 - ▣ Total participation on weak entity side



Common Attribute Mistakes

30

- Don't include entity-set primary key attributes on other entity-sets!
 - ▣ e.g. customers and loans, in a one-to-many mapping




- Even if every loan is owned by only one customer, this is still wrong
 - ▣ The association is recorded by the *relationship*, so specifying foreign key attributes on the entity-set is redundant

Common Attribute Mistakes (2)

31

- Don't include primary key attributes as descriptive attributes on relationship-set, either!
- This time, assume *borrower* is a 1:1 mapping
 - ▣ IDs used as descriptive attributes on *borrower*




- Again, this is implicit in the relationship

ENTITY-RELATIONSHIP MODEL III

N-ary Relationships

2

- Can specify relationships of degree > 2 in E-R model
- Example:




- Employees are assigned to jobs at various branches
- Many-to-many mapping: any combination of employee, job, and branch is allowed
- An employee can have several jobs at one branch

N-ary Mapping Cardinalities

3

- Can specify *some* mapping cardinalities on relationships with degree > 2
- Each combination of employee and branch can only be associated with one job:




- Each employee can have only one job at each branch

N-ary Mapping Cardinalities (2)

4


- For degree > 2 relationships, we only allow at most one edge with an arrow
- Reason: multiple arrows on N-ary relationship-set is ambiguous
 - ▣ (several meanings have been defined for this in the past)
- Relationship-set R associating entity-sets A_1, A_2, \dots, A_n
 - ▣ No arrows on edges A_1, \dots, A_i
 - ▣ Arrows are on edges to A_{i+1}, \dots, A_n
- Meaning 1 (the simpler one):
 - ▣ A particular combination of entities in A_1, \dots, A_i can be associated with at most one set of entities in A_{i+1}, \dots, A_n
 - ▣ Primary key of R is union of primary keys from set $\{A_1, A_2, \dots, A_i\}$



N-ary Mapping Cardinalities (3)

5

- Relationship-set R associating entity-sets A_1, A_2, \dots, A_n
 - ▣ No arrows on edges A_1, \dots, A_i ; arrows on edges to A_{i+1}, \dots, A_n
- Meaning 2 (the insane one):
 - ▣ For each entity-set A_k ($i < k \leq n$), a particular combination of entities from *all other* entity-sets can be associated with at most one entity in A_k
 - ▣ R has a candidate key for each arrow in N-ary relationship-set
 - ▣ For each k ($i < k \leq n$), another candidate key of R is union of primary keys from entity-sets $\{A_1, A_2, \dots, A_{k-1}, A_{k+1}, \dots, A_n\}$



Two candidate keys:
 $(a, b, c), (a, b, d)$

Meaning 2:

(a_1, b_2, c_3, d_5)
 (a_3, b_4, c_3, d_2)
 (a_1, b_2, c_1, d_4)
 (a_3, b_4, c_5, d_7)
 ~~(a_1, b_2, c_3, d_6)~~
 ~~(a_3, b_4, c_8, d_2)~~
...

} All disallowed
by meaning 1!

N-ary Mapping Cardinalities (4)

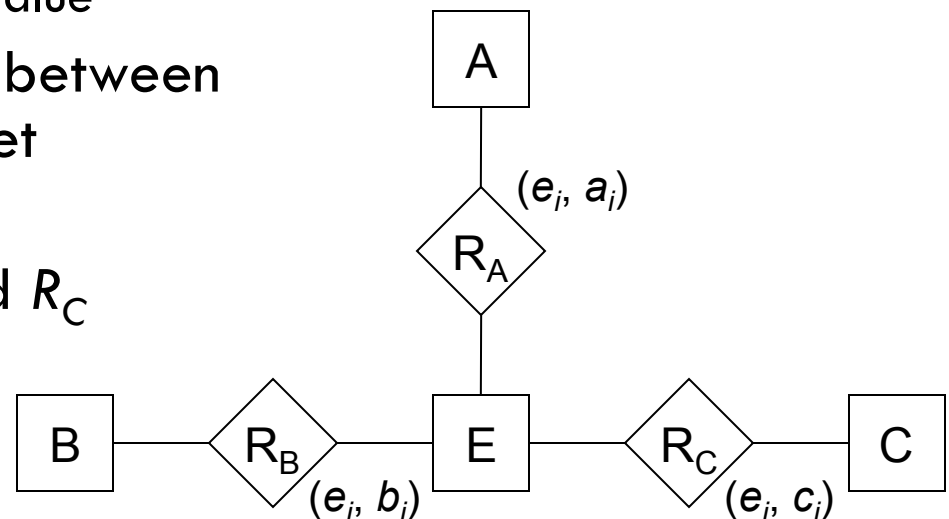

6

- Both interpretations of multiple arrows have been used in books and papers...
- If we only allow one edge to have an arrow, both definitions are equivalent
 - ▣ The ambiguity disappears

Binary vs. N-ary Relationships

7



- Often have only binary relationships in DB schemas
- For degree > 2 relationships, *could* replace with binary relationships
 - ▣ Replace N-ary relationship-set with a new entity-set E
 - Create an identifying attribute for E
 - e.g. an auto-generated ID value
 - ▣ Create a relationship-set between E and each other entity-set
 - ▣ Relationships in R must be represented in R_A , R_B , and R_C



Binary vs. N-ary Relationships (2)

8

- **Are these representations identical?**
- Example: Want to represent a relationship between entities a_5 , b_1 and c_2
 - ▣ How many relationships can we actually have between these three entities?
- Ternary relationship set:
 - ▣ Can only store one relationship between a_5 , b_1 and c_2 , due to primary key of R
- Alternate approach:
 - ▣ Can create many relationships between these entities, due to the entity-set E !
 - $(a_5, e_1), (b_1, e_1), (c_2, e_1)$
 - $(a_5, e_2), (b_1, e_2), (c_2, e_2)$
 - ...
 - ▣ Can't constrain in exactly the same ways



Binary vs. N-ary Relationships (3)

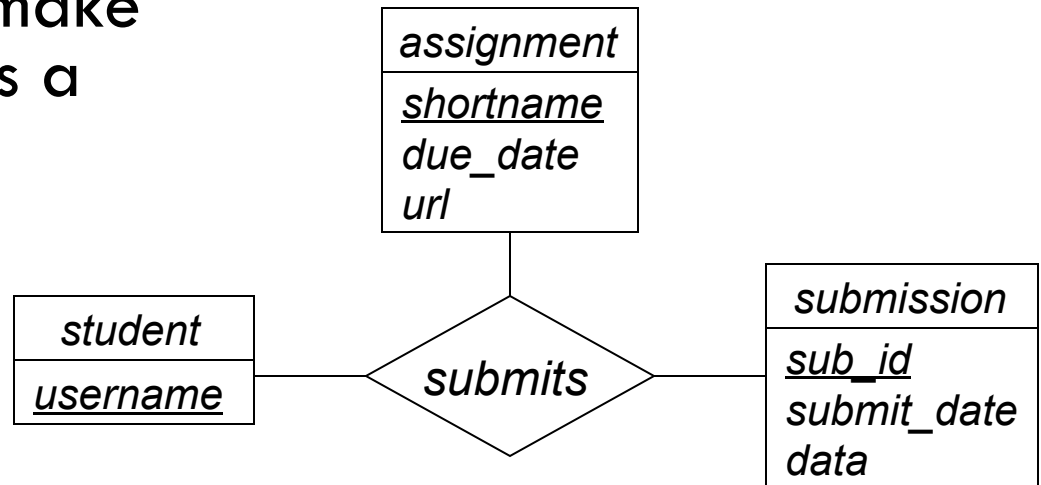
9

- Using binary relationships is sometimes more intuitive for particular designs
- Example: office-equipment inventory database
 - ▣ Ternary relationship-set *inventory*, associating *department*, *machine*, and *vendor* entity-sets
- What if vendor info is unknown for some machines?
 - ▣ For ternary relationship, must use *null* values to represent missing vendor details
 - ▣ With binary relationships, can simply not have a relationship between *machine* and *vendor*
- For cases like these, use binary relationships
 - ▣ If it makes sense to model as separate binary relationships, do it that way!

Course Database Example

10

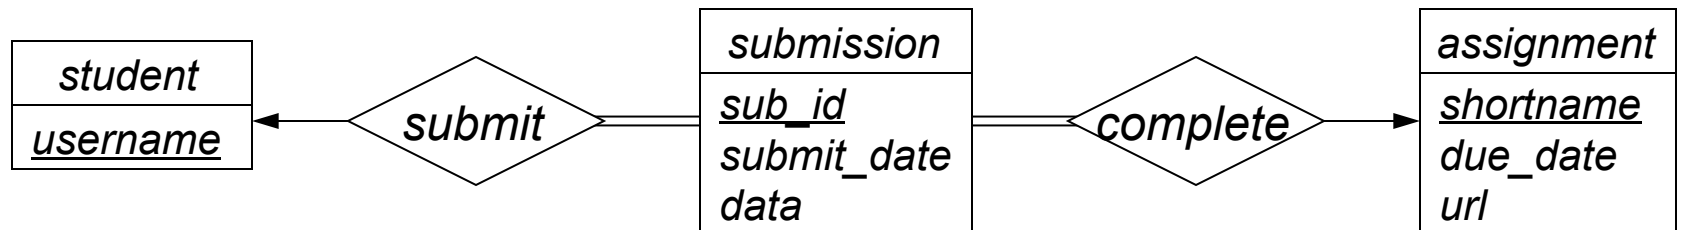
- What about this case:
 - ▣ Ternary relationship between *student*, *assignment*, and *submission*
 - ▣ Need to allow multiple submissions for a particular assignment, from a particular student
- In this case, it could make sense to represent as a ternary relationship
 - ▣ Doesn't make sense to have only two of these three entities in a relationship



Course Database Example (2)

11

- Other ways to represent students, assignments and submissions?
- Can also represent as two binary relationships

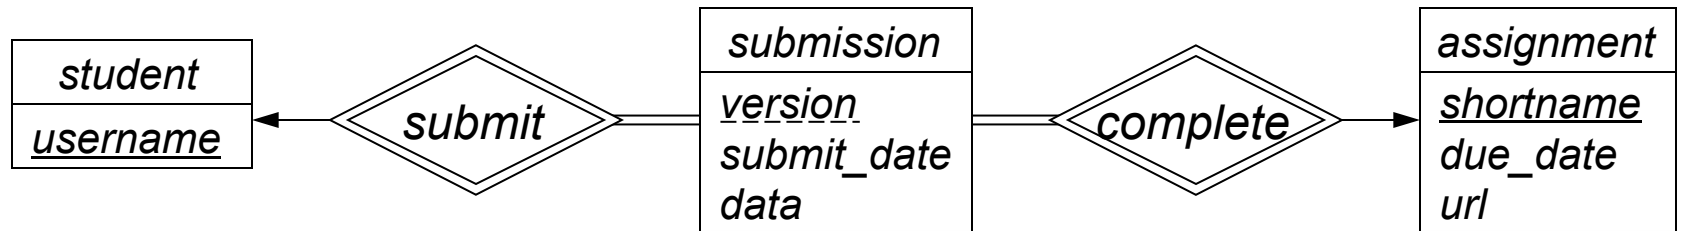


- Note the total participation constraints!
 - ▣ Required to ensure that every *submission* has an associated *student*, and an associated *assignment*
 - ▣ Also, two one-to-many constraints

Course Database Example (3)

12

- Could even make *submission* a weak entity-set
 - ▣ Both *student* and *assignment* are identifying entities!




- Discriminator for *submission* is version number
- Primary key for *submission* ?
 - ▣ Union of primary keys from all owner entity-sets, plus discriminator
 - ▣ (*username*, *shortname*, *version*)

Binary vs. N-ary Relationships

13

- Sometimes ternary relationships are best
 - ▣ Clearly indicates all entities involved in relationship
 - ▣ Only way to represent certain constraints!

- Bank jobs example:



- ▣ Each (employee, branch) pair can have only one job
- ▣ Simply cannot construct the same constraint using only binary relationships
 - (Reason is related to issue identified on slide 8)

E-R Model and Real Databases

14

- For E-R model to be useful, need to be able to convert diagrams into an implementation schema
- Turns out to be very easy to do this!
 - ▣ Big overlaps between E-R model and relational model
 - ▣ Biggest difference is E-R composite/multivalued attributes, vs. relational model atomic attributes
- Three components of conversion process:
 - ▣ Specify schema of the relation itself
 - ▣ Specify primary key on the relation
 - ▣ Specify any foreign key references to other relations

Strong Entity-Sets

15

- Strong entity-set E with attributes a_1, a_2, \dots, a_n
 - ▣ Assume simple, single-valued attributes for now
- Create a relation schema with same name E , and same attributes a_1, a_2, \dots, a_n
- Primary key of relation schema is same as primary key of entity-set
 - ▣ Strong entity-sets require no foreign keys to other things
- Every entity in E is represented by a tuple in the corresponding relation

Entity-Set Examples

16

- Geocache location E-R diagram:

- ▣ Entity-set named *location*

<i>location</i>
<u><i>latitude</i></u>
<u><i>longitude</i></u>
<i>description</i>
<i>last_visited</i>


- Convert to relation schema:

location(*latitude*, *longitude*, *description*, *last_visited*)

Entity-Set Examples (2)

17

- E-R diagram for customers and loans:



- Convert *customer* and *loan* entity-sets:
customer(*cust_id*, *name*, *street_address*, *city*)
loan(*loan_id*, *amount*)

Relationship-Sets

18

- Relationship-set R
 - ▣ For now, assume that all participating entity-sets are strong entity-sets
 - ▣ a_1, a_2, \dots, a_m is the union of all participating entity-sets' primary key attributes
 - ▣ b_1, b_2, \dots, b_n are descriptive attributes on R (if any)
- Relational model schema for R is:
 - ▣ $\{a_1, a_2, \dots, a_m\} \cup \{b_1, b_2, \dots, b_n\}$
- $\{a_1, a_2, \dots, a_m\}$ is a superkey, but not necessarily a candidate key
 - ▣ Primary key of R depends on R 's mapping cardinality

Relationship-Sets: Primary Keys

19

- For binary relationship-sets:
 - ▣ e.g. between strong entity-sets A and B
 - ▣ If many-to-many mapping:
 - Primary key of relationship-set is union of all entity-set primary keys
 - $primary_key(A) \cup primary_key(B)$
 - ▣ If one-to-one mapping:
 - Either entity-set's primary key is acceptable
 - $primary_key(A)$, or $primary_key(B)$
 - Enforce both candidate keys in DB schema!

Relationship-Sets: Primary Keys (2)

20

- For many-to-one or one-to-many mappings:
 - ▣ e.g. between strong entity-sets A and B
 - ▣ Primary key of entity-set on “many” side is primary key of relationship
- Example: relationship R between A and B
 - ▣ One-to-many mapping, with B on “many” side
 - ▣ Schema contains $primary_key(A) \cup primary_key(B)$, plus any descriptive attributes on R
 - ▣ $primary_key(B)$ is primary key of R
 - Each $a \in A$ can map to many $b \in B$
 - Each value for $primary_key(B)$ can appear only once in R


Relationship-Set Foreign Keys

21

- Relationship-sets associate entities in entity-sets
 - ▣ We need foreign-key constraints on relation schema for R !
- For each entity-set E_i participating in R :
 - ▣ Relation schema for R has a foreign-key constraint on E_i relation, for *primary_key*(E_i) attributes
- Relation schema notation doesn't provide mechanism for indicating foreign key constraints
 - ▣ Don't forget about foreign keys and candidate keys!
 - Making notes on your relational model schema is a very good idea
 - ▣ Can specify both foreign key constraints and candidate keys in the SQL DDL

Relationship-Set Example


22



- Relation schema for *borrower*:
 - ▣ Primary key of *customer* is *cust_id*
 - ▣ Primary key of *loan* is *loan_id*
 - ▣ Descriptive attribute *access_date*
 - ▣ *borrower* mapping cardinality is many-to-many
 - ▣ Result: *borrower*(*cust_id*, *loan_id*, *access_date*)

Relationship-Set Example (2)


23



- In cases like this, must use roles to distinguish between the entities involved in the relationship-set
 - ▣ *employee* participates in *works_for* relationship-set twice
 - ▣ Can't create a schema (*employee_id*, *employee_id*) !
- Change names of key-attributes to distinguish roles
 - ▣ e.g. (*manager_employee_id*, *worker_employee_id*)
 - ▣ e.g. (*manager_id*, *employee_id*)

Relationship-Set Example (2)

24



- Relation schema for *employee* entity-set:
 - ▣ (For now, ignore *phone_number* and *num_reports*...)
employee(*employee_id*, *name*)
- Relation schema for *works_for*:
 - ▣ One-to-many mapping from *manager* to *worker*
 - ▣ “Many” side is used for primary key
 - ▣ Result: *works_for*(*employee_id*, *manager_id*)


N-ary Relationship Primary Keys

25

- For degree > 2 relationship-sets:
 - ▣ If no arrows (“many-to-many” mapping), relationship-set primary key is union of all participating entity-sets’ primary keys
 - ▣ If one arrow (“one-to-many” mapping), relationship-set primary key is union of primary keys of entity-sets without an arrow
 - ▣ Don’t allow more than one arrow for relationship-sets with degree > 2

N-ary Relationship-Set Example

26



□ Entity-set schemas:

job(title, level)

employee(employee_id, employee_name)

branch(branch_name, branch_city, assets)

□ Relationship-set schema:

- ▣ Primary key includes entity-sets on non-arrow links

works_on(employee_id, branch_name, title)

Weak Entity-Sets

27

- Weak entity-sets depend on at least one strong entity-set
 - ▣ The identifying entity-set, or owner entity-set
 - ▣ Relationship between the two is called the identifying relationship
- Weak entity-set A owned by strong entity-set B
 - ▣ Attributes of A are $\{a_1, a_2, \dots, a_m\}$
 - Some subset of these attributes comprises the discriminator of A
 - ▣ $primary_key(B) = \{b_1, b_2, \dots, b_n\}$
 - ▣ Relation schema for A : $\{a_1, a_2, \dots, a_m\} \cup \{b_1, b_2, \dots, b_n\}$
 - ▣ Primary key of A is $discriminator(A) \cup primary_key(B)$
 - ▣ A has a foreign key constraint on $primary_key(B)$, to B


Identifying Relationship?

28

- The identifying relationship is many-to-one, with no descriptive attributes
- Relation schema for weak entity-set already includes primary key for strong entity-set
 - ▣ Foreign key constraint is imposed, too
- No need to create relational model schema for the identifying relationship
 - ▣ Would be redundant to the weak entity-set's relational model schema!

Weak Entity-Set Example

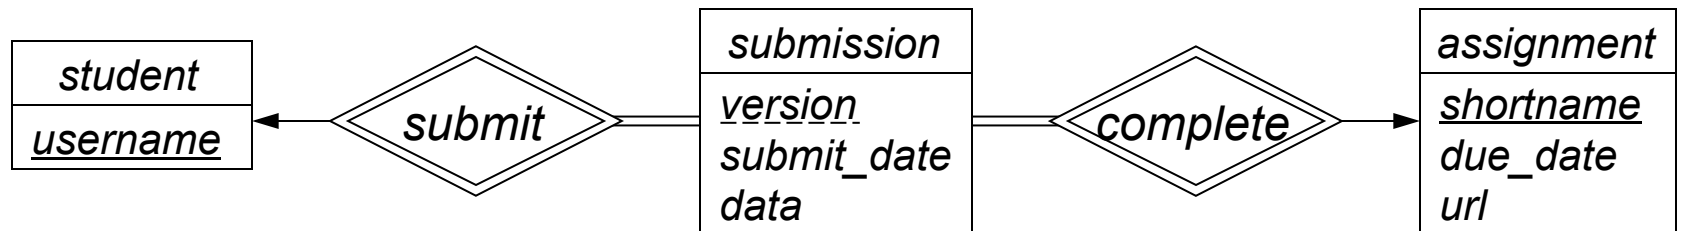
29



- *account* schema:
account(*account_number*, *balance*)
- *check* schema:
 - ▣ Discriminator is *check_number*
 - ▣ Primary key for *check* is: (*account_number*, *check_number*)*check*(*account_number*, *check_number*, *check_date*,
recipient, *amount*, *memo*)

Weak Entity-Set Example (2)

30



- Schemas for strong entity-sets:
student(username)
assignment(shortname, due_date, *url*)
- Schema for *submission* weak entity-set:
 - ▣ Discriminator is *version*
 - ▣ Both *student* and *assignment* are owners!*submission*(username, shortname, version, submit_date, *data*)
 - Two foreign keys in this relation as well

Composite Attributes

31

- Relational model simply doesn't handle composite attributes
 - ▣ All attribute domains are *atomic* in the relational model
- When mapping E-R composite attributes to relation schema: simply flatten the composite
 - ▣ Each component attribute maps to a separate attribute in relation schema
 - ▣ In relation schema, simply can't refer to the composite as a whole
 - ▣ (Can adjust this mapping for databases that support composite types)

Composite Attribute Example

32

- Customers with addresses:

<i>customer</i>
<u><i>cust_id</i></u>
<i>name</i>
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>

- Each component of *address* becomes a separate attribute

customer(*cust_id*, *name*, *street*, *city*, *state*, *zip_code*)

Multivalued Attributes

33

- Multivalued attributes require a separate relation
 - ▣ Again, no such thing as a multivalued attribute in the relational model
 - ▣ E-R constraint on multivalued attributes: in a specific entity's multivalued attribute, each value may only appear once
- For a multivalued attribute M in entity-set E
 - ▣ Create a relation schema R to store M , with attribute(s) A corresponding to the single-valued version of M
 - ▣ Attributes of R are: $primary_key(E) \cup A$
 - ▣ Primary key of R includes all attributes of R
 - Each value in M for an entity e must be unique
 - ▣ Foreign key from R to E , on $primary_key(E)$ attributes

Multivalued Attribute Example

34

- Change our E-R diagram to allow customers to have multiple addresses:

<i>customer</i>
<u><i>cust_id</i></u>
<i>name</i>
{ <i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i> }

- Now, must create a separate relation to store the addresses

customer(*cust_id*, *name*)

cust_addrs(*cust_id*, *street*, *city*, *state*, *zipcode*)

- ▣ Large primary keys aren't ideal – tend to be costly

ADVANCED E-R FEATURES

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 17

Extensions to E-R Model

2

- Basic E-R model is good for many uses
- Several extensions to the E-R model for more advanced modeling
 - ▣ Generalization and specialization
 - ▣ Aggregation
- These extensions can also be converted to the relational model
 - ▣ Introduces a few more design choices
- Will only discuss specialization today
 - ▣ See book §7.8.5 for details on aggregation (material will be included with Assignment 5 too)

Specialization

3

- An entity-set might contain distinct subgroups of entities
 - ▣ Subgroups have some different attributes, not shared by the entire entity-set
- E-R model provides specialization to represent such entity-sets
- Example: bank account categories
 - ▣ Checking accounts
 - ▣ Savings accounts
 - ▣ Have common features, but also unique attributes

Generalization and Specialization

4

- Generalization: a “bottom up” approach
 - ▣ Taking similar entity-sets and unifying their common features
 - ▣ Start with specific entities, then create generalizations from them
- Specialization: a “top down” approach
 - ▣ Creating general purpose entity-sets, then providing specializations of the general idea
 - ▣ Start with the general notion, then refine it
- Terms are basically equivalent
 - ▣ Book refers to generalization as the overarching concept

Bank Account Example


5

- Checking and savings accounts both have:
 - ▣ account number
 - ▣ balance
 - ▣ owner(s)
- Checking accounts also have:
 - ▣ overdraft limit and associated overdraft account
 - ▣ check transactions
- Savings accounts also have:
 - ▣ minimum balance
 - ▣ interest rate

Bank Account Example (2)

6

- Create entity-set to represent common attributes
 - ▣ Called the superclass, or higher-level entity-set
- Create entity-sets to represent specializations
 - ▣ Called subclasses, or lower-level entity-sets
- Join superclass to subclasses with hollow-head arrow(s)



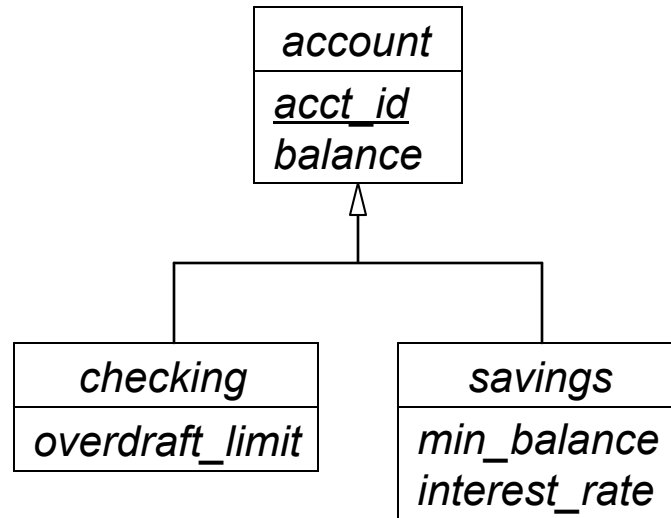
Inheritance

7

- Attributes of higher-level entity-sets are inherited by lower-level entity-sets
- Relationships involving higher-level entity-sets are also inherited by lower-level entity-sets!
 - ▣ Lower-level entity-sets can also participate in *their* own relationship-sets, separate from higher-level entity-set
- Usually, entity-sets inherit from one superclass
 - ▣ Entity-sets form a hierarchy
- Can also inherit from multiple superclasses
 - ▣ Entity-sets form a lattice
 - ▣ Introduces many subtle issues, of course

Specialization Constraints

8



- Can an account be both a savings account and a checking account?
- Can an account be neither a savings account nor a checking account?
- Can specify constraints on specialization
 - ▣ Enforce what “makes sense” for the enterprise

Disjointness Constraints

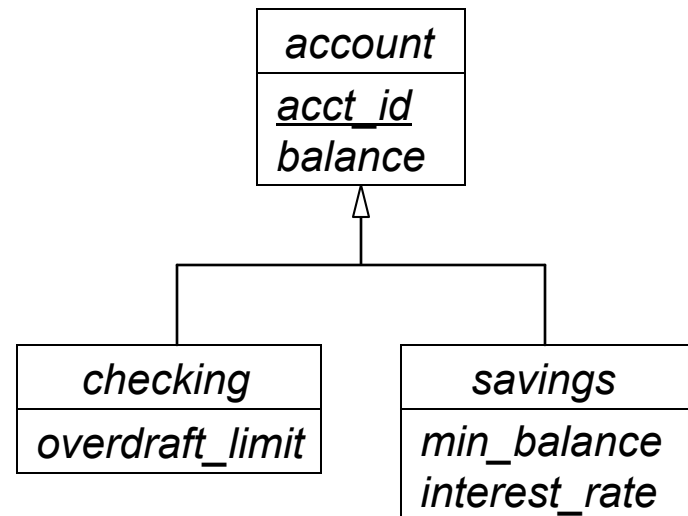
9

- “An account cannot be *both* a checking account and a savings account.”
- An entity may belong to at most one of the lower-level entity-sets
 - ▣ Must be a member of *checking*, or a member of *savings*, but not both!
 - ▣ Called a “disjointness constraint”
 - ▣ A better way to state it: a disjoint specialization
- If an entity can be a member of multiple lower-level entity-sets:
 - ▣ Called an overlapping specialization

Disjointness Constraints (2)

10


- How the arrows are drawn indicates whether the specialization is disjoint or overlapping
- Bank account example:
 - ▣ One arrow split into multiple parts indicates a disjoint specialization
 - ▣ An account may only be a checking account, or a savings account, not both



Disjointness Constraints (3)

11

- Another example from the book:
 - ▣ Specialization hierarchy for people at a university
- Multiple separate arrows indicates an overlapping specialization
 - ▣ A person can be an employee of the university and a student
- One arrow split into multiple parts is a disjoint specialization
 - ▣ An employee can be an instructor or a secretary, but not both



Completeness Constraints


12

- “An account must be a checking account, or it must be a savings account.”
- Every entity in higher-level entity-set must also be a member of at least one lower-level entity-set
 - ▣ Called total specialization
- If entities in higher-level entity-set aren’t required to be members of lower-level entity-sets:
 - ▣ Called partial specialization
- *account* specialization is a total specialization

Completeness Constraints (2)

13

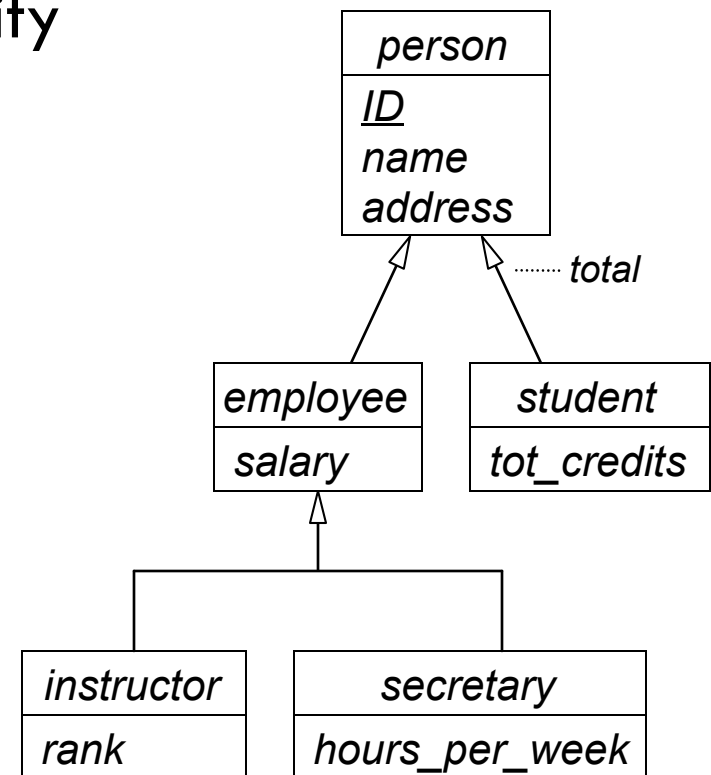
- Default constraint is partial specialization
- Specify total specialization constraint by annotating the specialization arrow(s)
- Updated bank account diagram:



Completeness Constraints (3)

14

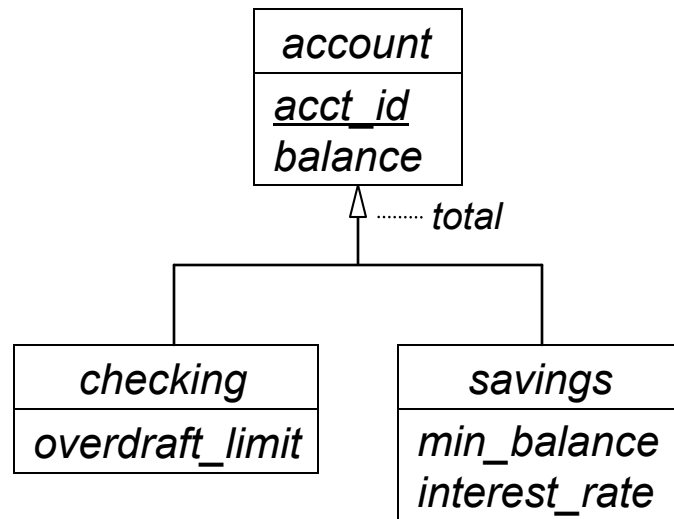
- Same approach with overlapping specialization
- Example: people at a university
 - ▣ Every person is an employee or a student
 - ▣ Not every employee is an instructor or a secretary
- Annotate arrows pointing to person with “total” to indicate total specialization
 - ▣ Every person must be an employee, a student, or both



Account Types?

15

- Our bank schema so far:



- How to tell whether an account is a checking account or a savings account?
 - ▣ No attribute indicates type of account

Membership Constraints

16

- Membership constraints specify which lower-level entity-sets each entity is a member of
 - ▣ e.g. which accounts are checking or savings accounts
- Condition-defined lower-level entity-sets
 - ▣ Membership is specified by a predicate
 - ▣ If an entity satisfies a lower-level entity-set's predicate then it is a member of that lower-level entity-set
 - ▣ If *all* lower-level entity-sets refer to the same attribute, this is called attribute-defined specialization
 - e.g. *account* could have an *account_type* attribute set to “c” for checking, or “s” for savings


Membership Constraints (2)

17

- Entities may simply be assigned to lower-level entity-sets by a database user
 - ▣ No explicit predicate governs membership
 - ▣ Called user-defined membership
- Generally used when an entity's membership could change in the future

Final Bank Account Diagram

18



- Would also create relationship-sets against various entity-sets in hierarchy
 - ▣ associate *customer* with *account*
 - ▣ associate *check* weak entity-set with *checking*

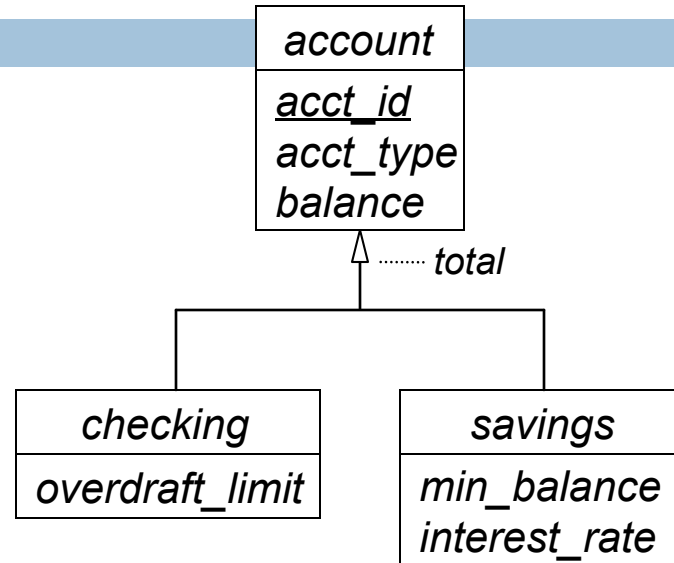
Mapping to Relational Model

19

- Mapping generalization/specialization to relational model is straightforward
- Create relation schema for higher-level entity-set
 - ▣ Including primary keys, etc.
- Create schemas for lower-level entity-sets
 - ▣ Subclass schemas include superclass' primary key attributes!
 - ▣ Primary key is same as superclass' primary key
 - Subclasses can also contain their own candidate keys!
 - Enforce these candidate keys in implementation schema
 - ▣ Foreign key reference from subclass schemas to superclass schema, on primary-key attributes

Mapping Bank Account Schema

20



□ Schemas:

`account(acct_id, acct_type, balance)`

`checking(acct_id, overdraft_limit)`

`savings(acct_id, min_balance, interest_rate)`

- ▣ Could use **CHECK** constraints on SQL tables for membership constraints, other constraints (although it may be expensive)

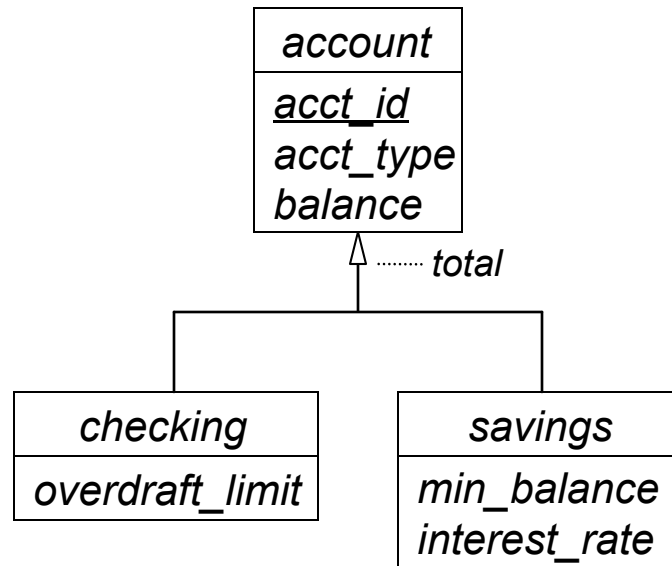
Alternative Schema Mapping

21

- If specialization is disjoint and complete, could convert only lower-level entity-sets to relational schemas
 - ▣ Every entity in higher-level entity-set also appears in lower-level entity-sets
 - ▣ Every entity is a member of *exactly one* lower-level entity-set
- Each lower-level entity-set has its own relation schema
 - ▣ All attributes of superclass entity-set are included on each subclass entity-set
 - ▣ No relation schema for superclass entity-set

Alternative Account Schema

22



□ Schemas, take 2:

checking(acct_id, acct_type, balance, overdraft_limit)

savings(acct_id, acct_type, balance, min_balance, interest_rate)

Alternative Account Schema (2)

23

- Alternative schemas:

checking(acct_id, acct_type, balance, overdraft_limit)

savings(acct_id, acct_type, balance, min_balance, interest_rate)

- Problems?

- ▣ Enforcing uniqueness of account IDs!

- ▣ Representing relationships involving both kinds of accounts

- Can solve by creating a simple relation:

account(acct_id)

- ▣ Contains *all* valid account IDs

- ▣ Relationships involving accounts can use *account*

- ▣ Need foreign key constraints again...

Generating Primary Keys

24

- Generating primary key values is actually the easy part
- Most databases provide sequences
 - ▣ A source of unique, increasing **INTEGER** or **BIGINT** values
 - ▣ Perfect for primary key values
 - ▣ Multiple tables can use the same sequence for their primary keys

- PostgreSQL example:

```
CREATE SEQUENCE acct_seq;
```

```
CREATE TABLE checking (  
    acct_id INT PRIMARY KEY DEFAULT nextval('acct_seq');  
    ...  
);
```

```
CREATE TABLE savings (  
    acct_id INT PRIMARY KEY DEFAULT nextval('acct_seq');  
    ...  
);
```


Alternative Schema Mapping

25

- Alternative mapping has serious drawbacks
 - ▣ Doesn't actually give many benefits in general case
- Fewer drawbacks if:
 - ▣ Total, disjoint specialization
 - ▣ No relationships against superclass entity-set
- If specialization is overlapping, some details will be stored multiple times
 - ▣ Unnecessary redundancy, and consistency issues
- Also limits future schema changes
 - ▣ Should always think about this when creating schemas

Recap: Weak Entity-Set Example

26




- *account* schema:
account(*account_number*, *balance*)
- *check* schema:
 - ▣ Discriminator is *check_number*
 - ▣ Primary key for *check* is: (*account_number*, *check_number*)*check*(*account_number*, *check_number*, *check_date*,
recipient, *amount*, *memo*)

Schema Combination

27


- Relationship between weak entity-set and strong entity-set doesn't need represented separately
 - ▣ Many-to-one relationship
 - ▣ Weak entity-set has total participation
 - ▣ Weak entity-set's schema already captures the identifying relationship
- Can apply this technique to other relationship-sets:
 - ▣ One-to-many mapping, with total participation on the “many” side



Schema Combination (2)

28

- Entity-sets A and B , relationship-set AB
 - ▣ Many-to-one mapping from A to B
 - ▣ A 's participation in AB is total
- Generates relation schemas A , B , AB
 - ▣ Primary key of A is $primary_key(A)$
 - ▣ Primary key of AB is also $primary_key(A)$
 - (A is on “many” side of mapping)
 - ▣ AB has foreign key constraints on both A and B
 - ▣ There is one relationship in AB for every entity in A
- Can combine A and AB relation schemas
 - ▣ Primary key of combined schema still $primary_key(A)$
 - ▣ Only requires one foreign-key constraint, to B




Schema Combination (3)

29


- In this case, when relationship-set is combined into the entity-set, the entity-set's primary key *doesn't change!*

- If A 's participation in AB is partial, can still combine schemas



- Must store *null* values for *primary_key(B)* attributes when an entity in A maps to no entity in B


- If AB is one-to-one mapping:



- Can also combine schemas in this case
- Could incorporate AB into schema for A , or schema for B
- Don't forget that AB has two candidate keys...
 - The combined schema must still enforce both candidate keys

Schema-Combination Example


30



- Manager to worker mapping is one-to-many
- Relation schemas were:
 `employee(employee_id, name)`
 `works_for(employee_id, manager_id)`
- Could combine into:
 `employee(employee_id, name, manager_id)`
 - ▣ (A very common schema combination)
 - ▣ Need to store *null* for employees with no manager

Schema Combination Example (2)


31



- One-to-one mapping between customers and loans
customer(*cust_id*, *name*, *street_address*, *city*)
loan(*loan_id*, *amount*)
borrower(*cust_id*, *loan_id*) – *loan_id* also a candidate key
- Could combine *borrower* schema into *customer* schema or *loan* schema
 - ▣ Does it matter which one you choose?

Schema Combination Example (3)


32



- Participation of *loan* in *borrower* will be total
 - ▣ Combining *borrower* into *customer* would require *null* values for customers without loans
- Better to combine *borrower* into *loan* schema
 - customer*(*cust_id*, *name*, *street_address*, *city*)
 - loan*(*loan_id*, *cust_id*, *amount*)
 - ▣ No *null* values!

Schema Combination Example (4)

33



- Schema:
 - customer*(cust_id, name, street_address, city)
 - loan*(loan_id, cust_id, amount)
- What if, after a while, we wanted to change the mapping cardinality?
 - ▣ Schema changes would be significant
 - ▣ Would need to migrate existing data to a new schema

Schema Combination Notes

34

- Benefits of schema combination:
 - ▣ Usually eliminates one foreign-key constraint, and the associated performance impact
 - Constraint enforcement
 - Extra join operations in queries
 - ▣ Reduces storage requirements
- Drawbacks of schema combination:
 - ▣ May necessitate the use of *null* values to represent the absence of relationships
 - ▣ Makes it harder to change mapping cardinality constraints in the future