

# SQL DDL II

CS121: Introduction to Relational Database Systems  
Fall 2014 – Lecture 8

# Last Lecture

2

- Covered SQL constraints
  - ▣ **NOT NULL** constraints
  - ▣ **CHECK** constraints
  - ▣ **PRIMARY KEY** constraints
  - ▣ **FOREIGN KEY** constraints
  - ▣ **UNIQUE** constraints
- Impact of **NULL** values on constraint enforcement
  - ▣ Specifically, **FOREIGN KEY** and **UNIQUE**...
- Automatic resolution of constraint violation

# Constraint Names

3

- Can assign names to constraints
  - ▣ When constraint is violated, error indicates which constraint
  - ▣ Database usually assigns names to constraints if you don't
  - ▣ Rules on constraint names vary

- Example:

```
CREATE TABLE employee (  
    ...  
    CONSTRAINT emp_pk      PRIMARY KEY (emp_id),  
    CONSTRAINT emp_ssn_ck  UNIQUE (emp_ssn),  
    CONSTRAINT emp_mgr_fk  FOREIGN KEY (manager_id)  
                        REFERENCES employee
```

- Useful for referring to specific constraints

# Temporary Constraint Violation

4

- Constraints take time to enforce
  - ▣ Can dramatically impact performance of large data-import operations
- Some operations may need to temporarily violate constraints
  - ▣ The operation is performed within a larger transaction (i.e. a batch of operations that should be treated as a unit)
  - ▣ During the transaction, constraints are temporarily violated
  - ▣ At end of transaction, constraint is restored
- Defer constraint enforcement to end of transaction
  - ▣ At end of transaction, all changes are checked against deferred constraints

# Deferring Constraint Application

5

- Can mark constraints as deferrable
- In constraint declaration, specify:
  - ▣ **DEFERRABLE** constraints may be deferred to end of transaction
  - ▣ **NOT DEFERRABLE** constraints are *always* applied immediately
- For **DEFERRABLE** constraints:
  - ▣ **INITIALLY IMMEDIATE** is applied immediately by default
  - ▣ **INITIALLY DEFERRED** is applied at end of transaction by default

# Temporarily Removing Constraints

6

- To defer constraints in current transaction:
  - `SET CONSTRAINTS c1, c2, ... DEFERRED;`
  - ▣ Specified constraints must be deferrable
- Not all databases support deferred constraints
  - ▣ Only option is to temporarily remove and then reapply constraints
  - ▣ Will usually affect all users of database! Safest to ensure exclusive access for this.
  - ▣ Remove, then reapply constraints with **ALTER TABLE** syntax

# Date and Time Values

7

- SQL provides data types for dates and times
- **DATE**
  - ▣ A calendar date, including year, month, and day of month
- **TIME**
  - ▣ A time of day, including hour, minute, and second value
  - ▣ Doesn't include fractional seconds
- **TIME (P)**
  - ▣ Just like **TIME**, but includes P digits of fractional seconds
  - ▣ Typically,  $P = [0, 6]$

# Date and Time Values (2)

8

- Can include timezone info as well:
  - ▣ **TIME WITH TIMEZONE**
  - ▣ **TIME (P) WITH TIMEZONE**
- **TIMESTAMP**
  - ▣ A combination of date and time values
  - ▣ Includes fractional seconds by default
  - ▣ Can also specify **TIMESTAMP (P)**
  - ▣ P = 6 by default
  - ▣ Timestamps can also include time zone info
    - **TIMESTAMP WITH TIMEZONE**
    - **TIMESTAMP (P) WITH TIMEZONE**



# Date and Time Values (3)

9

- Often a variety of other non-standard types
  - ▣ **DATETIME** – Like **TIMESTAMP** but  $P = 0$  by default
  - ▣ **YEAR** – Just a 4-digit year value
  - ▣ Nonstandard = not portable

# Microsoft SQLServer Date Types

10

- SQLServer 2005 and earlier provide very different date/time support
  - ▣ **DATETIME** – more like standard **TIMESTAMP** type
    - Represents both date and time
    - Jan 1, 1753 – Dec 31, 9999; precision of 3.33ms (???)
  - ▣ **SMALLDATETIME**
    - Jan 1, 1900 – Jun 6, 2079; precision of 1 minute
  - ▣ No ability to represent only a date, or only a time!
- SQLServer 2008 adds more standard-like support
  - ▣ **DATE, TIME, DATETIME2** – similar to standard types
  - ▣ **DATETIMEOFFSET** – date/time value plus timezone

# Date and Time Formats

11

- Date and time values follow specific formats
  - ▣ Enclosed in single-quotes
- Examples: MER-A “Spirit” launch time
  - ▣ Timestamp value (UT; +0):  
`'2003-06-10 17:58:46.773'`
  - ▣ Date value: `'2003-06-10'`
  - ▣ Time value: `'17:58:47'`
- Can have invalid date/time values:
  - ▣ Invalid time: `'25:14:68'`
  - ▣ Invalid date: `'2001-02-31'`
  - ▣ Some DBMSes can allow partial/invalid dates and times, if required by an application

# Date and Time Formats (2)

12

- Most DBMSes support many date/time formats
- Most widely supported is ISO-8601 date/time format
  - ▣ ISO-8601 format:
    - '2003-06-10 17:58:46.773'
    - year-month-day hour:minutes:seconds.milliseconds
    - Sometimes date and time are separated by “T” character
    - Time is in 24-hour time format
    - Optional timezone specification at end
  - ▣ Other formats:
    - 'June 10, 2003 5:58:46 PM'
    - '10-Jun-2003 17:58:46.773'
  - ▣ Most databases can parse all of these

# “Current Time” Values

13

- Several functions provide current date and time values

**CURRENT\_DATE ( )**

**CURRENT\_TIME ( )**

**CURRENT\_TIMESTAMP ( )**

- ▣ Include time zone information

**LOCALTIME ( )**

**LOCALTIMESTAMP ( )**

- ▣ Don't include time zone information

- Usually many other functions too, e.g. **NOW ( )**

- ▣ Nonstandard, but widely supported

# Components of Dates and Times

14

- Date and time values are *not* atomic
  - ▣ Not really allowed in the Relational Model...
  - ▣ (In reality, many SQL types are not atomic)
- SQL provides a function to extract components of dates and times
  - ▣ **EXTRACT** (*field* FROM *value*)
  - ▣ Can specify:
    - YEAR, MONTH, DAY, HOUR, MINUTE, SECOND
    - TIMEZONE\_HOUR, TIMEZONE\_MINUTE
  - ▣ Many other (nonstandard but common) options too
    - week of year, day of year, day of week, quarter, century, ...

# Example Date Operation

15

## □ Sales records:

```
CREATE TABLE salesrecords (  
    sale_id INTEGER PRIMARY KEY,  
    cust_id INTEGER NOT NULL,  
    sale_time TIMESTAMP NOT NULL,  
    sales_total NUMERIC(8, 2) NOT NULL,  
    ...  
);
```

## □ Compute monthly sales totals:

- Start by finding month of each sale

```
SELECT sale_id,  
    EXTRACT (MONTH FROM sale_time) AS sale_month  
FROM salesrecords;
```

- Build larger query using this information

# Time Intervals

16

- **INTERVAL**
  - ▣ Data type for time intervals
  - ▣ Supports operations on dates and times
  - ▣ Also supports a precision: **INTERVAL (P)**
- If  $x$  and  $y$  are date values:  
 $x - y$  produces an **INTERVAL**
- If  $i$  is an **INTERVAL** value:  
 $x + i$  or  $x - i$  produces a date value
- Can use **INTERVAL** to specify fixed intervals
  - ▣ **INTERVAL 1 WEEK**
  - ▣ **INTERVAL '1 WEEK'**



# Example Date Schema

17

- Event database schema:

```
CREATE TABLE event (  
    event_id    INTEGER          PRIMARY KEY,  
    event_type  VARCHAR(20)     NOT NULL,  
    event_date  DATE             NOT NULL,  
    event_desc  VARCHAR(200)  
);
```

- To generate notices of upcoming events:

```
SELECT * FROM event  
WHERE event_date >= CURRENT_DATE() AND  
      event_date <=  
      (CURRENT_DATE() + INTERVAL 1 WEEK);
```

# Example Date Schema (2)

18

- Can rewrite to use **BETWEEN** syntax:

```
SELECT * FROM event
WHERE event_date BETWEEN
      CURRENT_DATE() AND
      (CURRENT_DATE() + INTERVAL 1 WEEK);
```

- Current date/time functions are evaluated only once during a query! 😊
  - ▣ e.g. query will see one value for **CURRENT\_TIME()** even if it runs for an extended period of time

# “Large Object” Types

19

- **SQL CHAR (N)** and **VARCHAR (N)** types have limited sizes
  - ▣ For **CHAR**, usually  $N < 256$
  - ▣ For **VARCHAR**, usually  $N < 65536$
- **BLOB** and **CLOB** types support larger data sizes
  - ▣ “LOB” = Large Object
  - ▣ Useful for storing images, documents, etc.
  - ▣ Support varies widely across DBMSes
  - ▣ **TEXT** is also rather common
    - Large text fields, e.g. MB or GB of text data

# Example Schema

20

- Schema for storing book reviews:

```
CREATE TABLE bookreview (  
    review_id    INT PRIMARY KEY,  
    book_title   VARCHAR(50) NOT NULL,  
    book_image   BLOB,  
    reviewer     VARCHAR(30) NOT NULL,  
    pub_time     TIMESTAMP NOT NULL,  
    review_text  CLOB NOT NULL,  
    UNIQUE (book_title, reviewer)  
);
```

- Review text can be large
- Can also include a book image, if desired

# Large Object Notes

21

- General support for “large object” types is usually focused on smaller objects
  - ▣ No larger than a few 10s of KBs
  - ▣ A few MBs is *definitely* pushing it
- Most expensive part is moving large objects into and out of database
  - ▣ For simple, general purpose DBMSes, can involve constructing large SQL statements with escaped data
- Databases also don’t store this information very efficiently

# Large Object Notes (2)

22

- For objects larger than  $\sim 100$  KB, should definitely use the filesystem
  - ▣ That's what it's designed for!
  - ▣ Store *filesystem paths* in the database instead
- For smaller objects that are frequently retrieved, storing on filesystem can take load off database
  - ▣ e.g. user icons for a social networking website
  - ▣ Let webserver serve them directly from the filesystem – again, it knows how to do that kind of thing more quickly
- Some DBMSes have specialized support for storing and manipulating very large objects
  - ▣ Just don't expect your application to be easily portable...

# Default Values

23

- Can specify default values for columns
  - *colname type DEFAULT expr*
  - ▣ Can specify an actual value
    - *book\_rating INT DEFAULT 3*
  - ▣ Can specify an expression
    - *pub\_time TIMESTAMP DEFAULT NOW()*
- If unspecified, default value is **NULL**
- Affects **INSERT** statements
  - ▣ Columns with default values don't have to be specified
  - ▣ Columns without a default value *must* be specified at insert-time!

# Serial Primary Key Values

24

- Many databases offer special support for integer primary keys
  - ▣ DB will generate unique values for use as primary keys
- Examples:
  - ▣ PostgreSQL and MySQL:

```
CREATE TABLE employee (  
    emp_id    SERIAL PRIMARY KEY,  
    ...
```
  - ▣ Microsoft SQLServer:

```
CREATE TABLE employee (  
    emp_id    INT IDENTITY PRIMARY KEY,  
    ...
```



# Updated Book Review Schema

25

```
CREATE TABLE bookreview (  
    review_id    SERIAL          PRIMARY KEY,  
    book_title   VARCHAR(50) NOT NULL,  
    book_image   BLOB,  
    reviewer     VARCHAR(30) NOT NULL,  
    pub_time     TIMESTAMP      NOT NULL DEFAULT NOW() ,  
    book_rating  INT             NOT NULL DEFAULT 3,  
    review_text  CLOB            NOT NULL,  
    UNIQUE (book_title, reviewer)  
);
```

- Every new review gets a unique ID value
- Publication time is set to current time when review is added to database
- Default book rating is 3 out of 5

# Altering Table Schemas

26

- ❑ **SQL ALTER TABLE** command allows schema changes
- ❑ Wide variety of operations
  - ▣ Rename a table
  - ▣ Add and remove constraints
  - ▣ Add and remove table columns
  - ▣ Change the type of a column
  - ▣ Change default values for columns
- ❑ Very useful for migrating schema to new version
  - ▣ Migration process must be carefully designed...
- ❑ Again, support varies across DBMSes

# Example Alterations

27

- Rename the *bookreview* table:

```
ALTER TABLE bookreview  
    RENAME TO item_review;
```

- Remove the book image column:

```
ALTER TABLE bookreview  
    DROP COLUMN book_image;
```

- Add a constraint to the *bookreview* table:

```
ALTER TABLE bookreview  
    ADD CHECK (book_rating BETWEEN 1 AND 5);
```

# Table Alteration Notes

28

- Can drop columns from tables
  - ▣ What if the column is a key?
  - ▣ What if the column is referenced by a view?
  - ▣ Can often specify **CASCADE** to delete dependent objects, if desired
- Newly added columns must have a default value
  - ▣ Existing rows in database get default value for new column
- Changing table schema can be very expensive
  - ▣ Some operations can require scanning or rewriting the entire table
    - Some DBs do this for all schema-alteration commands, e.g. MySQL
  - ▣ e.g. adding a new constraint requires a table scan

# Temporary Tables

29

- Sometimes want to generate and store relations temporarily
  - ▣ Complex operations implemented as multiple queries
  - ▣ This is relational algebra assignment operation: ←
- SQL provides temporary tables for these cases
  - ▣ Table's contents are associated with client's session
  - ▣ Clients can't access each others' temp table data
- SQL standard specifies global temporary tables
  - ▣ Temporary table has a global name and schema
  - ▣ Only the contents of the temporary table are per-client
  - ▣ When client disconnects, their temporary data is purged

# Temporary Tables (2)

30

- Many databases also provide local temporary tables
  - ▣ Table's schema is also local to client session
  - ▣ When client disconnects, the table is dropped
  - ▣ Different clients can use same table name with different schemas
- Client can manually purge data from temp tables when needed
  - ▣ In case of local temp tables, can also drop them anytime during session

# Temporary Table Syntax

31

- Simple variation of **CREATE TABLE** syntax
  - ▣ Add **TEMPORARY** (or **GLOBAL TEMPORARY**) to command
- Example:
  - ▣ Make a temporary table to store counts of sales grouped by month

```
CREATE TEMPORARY TABLE salesbymonth (  
    sale_month INT NOT NULL,  
    num_sales  INT NOT NULL  
);
```

# Temporary Table Example

32

- Can populate temp table with computed values

```
INSERT INTO salesbymonth
    SELECT EXTRACT (MONTH FROM sale_time) AS mon,
           COUNT (*)
    FROM salesrecords GROUP BY mon;
```

- Only need to perform computations once
- Improves efficiency of large or multi-step operations
- Temporary results are cleaned up at end of session
- Issue queries against temporary table and use results

```
SELECT sale_month, num_sales, promotion_desc
    FROM salesbymonth
    JOIN promotions USING (sale_month);
```



# Temporary Table Contents

33

- When to flush temporary table contents?
- Two main options:
  - ▣ At end of current transaction
  - ▣ When entire client session ends
- Can specify behavior with **ON COMMIT** clause at end of table declaration
  - ▣ To flush temp table at end of each transaction:  
**ON COMMIT DELETE ROWS**
  - ▣ To flush temp table at end of session:  
**ON COMMIT PRESERVE ROWS**
  - ▣ SQL standard specifies default is **DELETE ROWS !**
    - Not all DBMSes follow this, but some do!

# Example ON COMMIT Clauses

34

- To flush rows after each transaction:

```
CREATE TEMPORARY TABLE salesbymonth (  
    sale_id      INT NOT NULL,  
    sale_month   INT NOT NULL  
) ON COMMIT DELETE ROWS;
```

- To keep rows until end of session:

```
CREATE TEMPORARY TABLE salesbymonth (  
    sale_id      INT NOT NULL,  
    sale_month   INT NOT NULL  
) ON COMMIT PRESERVE ROWS;
```

# Using Temporary Tables

35

- Temporary tables can dramatically improve performance of certain queries
- Approach:
  - ▣ Create temporary table to store useful but costly intermediate results
    - Don't use many (or any) constraints – want to be fast!
  - ▣ Populate temporary table via **INSERT ... SELECT** statement
  - ▣ Use temporary table to compute other results
  - ▣ Temporary table goes away automatically, at end of transaction, or at end of session

# Alternate Temp-Table Syntaxes

36

- Databases frequently support alternate syntaxes for creating and populating temporary tables
  - ▣ Simplify the common case!
- One common syntax (e.g. MySQL, Postgres, Oracle):  
**CREATE TEMPORARY TABLE *tblname* AS  
    *select\_stmt*;**
- Another common syntax (e.g. Postgres, SQLServer):  
**SELECT ... INTO TEMPORARY TABLE ...;**
- Both syntaxes can also create non-temporary tables

# Real-World Example

37

- A query run on a MySQL server:

```
SELECT ident, total_a / total_b AS ratio
FROM (SELECT CONCAT(a1, a2) AS ident,
                SUM(val_a) AS total_a
        FROM t1 GROUP BY ident) AS result1,
     (SELECT CONCAT(a1, a2) AS ident,
                SUM(val_b) AS total_b
        FROM t2 GROUP BY ident) AS result2
WHERE result1.ident = result2.ident;
```

- Overall query takes ~15 mins to execute on fast server
- Inner queries complete in << 1 second by themselves

# Real-World Example (2)

38

## □ MySQL query:

```
SELECT ident, total_a / total_b AS ratio
FROM (SELECT CONCAT(a1, a2) AS ident,
                SUM(val_a) AS total_a
        FROM t1 GROUP BY ident) AS result1,
     (SELECT CONCAT(a1, a2) AS ident,
                SUM(val_b) AS total_b
        FROM t2 GROUP BY ident) AS result2
WHERE result1.ident = result2.ident;
```

## □ Problem is that MySQL cannot *efficiently* join two derived results using a computed column

- ▣ A limitation of MySQL's join processor ☹

# Real-World Example (3)

39

## □ A solution:

- First, create temporary tables to hold intermediate results

```
CREATE TEMPORARY TABLE temp1 AS  
  SELECT CONCAT(a1, a2) AS ident,  
         SUM(val_a) AS total_a  
  FROM t1 GROUP BY ident;
```

- ...same with other inner query...

- Second, create indexes on temporary tables
- Finally, issue outer query against temporary tables

## □ Result:

- Entire process, including create/drop temp tables, takes < 1 second (as opposed to ~15 minutes)