

DATABASE SCHEMA DESIGN

ENTITY-RELATIONSHIP MODEL

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 14

Designing Database Applications

2

- Database applications are large and complex
- A few of the many design areas:
 - ▣ Database schema (physical/logical/view)
 - ▣ Programs that access and update data
 - ▣ Security constraints for data access
- Also requires familiarity with the problem domain
 - ▣ Domain experts *must* help drive requirements

General Approach

3

- Collect user requirements
 - ▣ Information that needs to be represented
 - ▣ Operations to perform on that information
 - ▣ Several techniques for representing this info, e.g. UML
- Develop a conceptual schema of the database
 - ▣ A high-level representation of the database's structure and constraints
 - Physical *and* logical design issues are ignored at this stage
 - ▣ Follows a specific data model
 - ▣ Often represented graphically

Conceptual Schema

4

- Also need to create a specification of functional requirements
 - ▣ “What operations will be performed against the data?”
 - ▣ Updating data, adding data, deleting data, ...
- Designer can use functional requirements to verify the conceptual schema
 - ▣ Is each operation possible?
 - ▣ How complicated or involved is it?
 - ▣ Performance or scalability concerns?

Implementation Phases

5

- Once conceptual schema and functional requirements are verified:
 - ▣ Convert conceptual schema into an implementation data model
 - ▣ Want to have a simple mapping from conceptual model to implementation model
- Finally: any necessary physical design
 - ▣ Not always present!
 - ▣ Smaller applications have few physical design concerns
 - ▣ Larger systems usually need additional design and tuning (e.g. indexes, disk-level partitioning of data)

Importance of Design Phase

6

- Not all changes have the same impact!
- Physical-level changes have the least impact
 - ▣ (Thanks, relational model!)
 - ▣ Typically affect performance, scalability, reliability
 - ▣ Little to no change in functionality
- Logical-level changes are typically *much* bigger
 - ▣ Affects how to interact with the data...
 - ▣ Also affects what is even *possible* to do with the data
- Very important to spend time up front designing the database schema

Design Decisions

7

- Many different ways to represent data
- Must avoid two major problems:
 - ▣ Unnecessary redundancy
 - Redundant information wastes space
 - Greater potential for inconsistency!
 - Ideally: each fact appears in exactly one place
 - ▣ Incomplete representation
 - Schema must be able to fully represent all details and relationships required by the application

More Design Decisions

8

- Even with correct design, usually many other concerns
 - ▣ How easy/hard is it to access useful information? (e.g. reporting or summary info)
 - ▣ How hard is it to update the system?
 - ▣ Performance considerations?
 - ▣ Scalability considerations?
- Schema design requires a good balance between aesthetic and practical concerns
 - ▣ Frequently need to make compromises between conflicting design principles

Simple Design Example

9

- Purchase tracking database
 - ▣ Store details about product purchases by customers
 - ▣ Actual purchases tracked in database
- Can represent sales as relationships between customers and products
 - ▣ What if product price changes? Where to store product sale price? Will this affect other recent purchases?
 - ▣ What about giving discounts to some customers? May want to give different prices to different customers.
- Can also represent sales as separate entities
 - ▣ Gives much more flexibility for special pricing, etc.

The Entity-Relationship Model

10

- A very common model for schema design
 - ▣ Also written as “E-R model”
- Allows for specification of complex schemas in graphical form
- Basic concepts are simple, but can also represent very sophisticated abstractions
 - ▣ e.g. type hierarchies
- Can be mapped very easily to the relational model!
 - ▣ Simplifies implementation phase
 - ▣ Mapping process can be automated by design tools

Entities and Entity-Sets

11

- An entity is any “thing” that can be uniquely represented
 - e.g. a product, an employee, a software defect
 - ▣ Each entity has a set of attributes
 - ▣ Entities are uniquely identified by some set of attributes
- An entity-set is a named collection of entities of the same type, with the same attributes
 - ▣ Can have multiple entity-sets with same entity type, representing different (possibly overlapping) sets

Entities and Entity-Sets (2)

12

- An entity has a set of attributes
 - ▣ Each attribute has a name and domain
 - ▣ Each attribute also has a corresponding value
- Entity-sets also specify a set of attributes
 - ▣ Every entity in the entity-set has the same set of attributes
 - ▣ Every entity in the entity-set has its own value for each attribute


Diagramming an Entity-Set

13

Example: a *customer* entity-set

- Attributes:

- *cust_id*
- *name*
- *street_address*
- *city*



- Entity-set is denoted by a box
- Name of entity-set is given in the top part of box
- Attributes are listed in the lower part of the box

Relationships

14

- A relationship is an association between two or more entities
 - ▣ e.g. a bank loan, and the customer who owns it
- A relationship-set is a named collection of relationships of the same type
 - ▣ i.e. involving the same entities
- Formally, a relationship-set is a mathematical relation involving n entity-sets, $n \geq 2$
 - ▣ E_1, E_2, \dots, E_n are entity sets; e_1, e_2, \dots are entities in E_1, E_2, \dots
 - ▣ A relationship set R is a subset of:
 $\{ (e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n \}$
 - ▣ (e_1, e_2, \dots, e_n) is a specific relationship in R

Relationships (2)

15

- Entity-sets participate in relationship-sets
 - ▣ Specific entities participate in a relationship instance
- Example: bank loans
 - ▣ *customer* and *loan* are entity-sets
 - (555-55-5555, Jackson, Woodside) is a *customer* entity
 - (L-14, 1500) is a *loan* entity
 - ▣ *borrower* is a relationship-set
 - *customer* and *loan* participate in *borrower*
 - *borrower* contains a relationship instance that associates customer “Jackson” and loan “L-14”

Relationships and Roles

16

- An entity's role in a relationship is the function that the entity fills

Example: a *purchase* relationship between a *product* and a *customer*

- the product's role is that it was purchased
- the customer did the purchasing
- Roles are usually obvious, and therefore unspecified
 - Entities participating in relationships are distinct...
 - Names clearly indicate the roles of various entities...
 - In these cases, roles are left unstated.

Relationships and Roles (2)

17

- Sometimes the roles of entities are *not* obvious
 - ▣ Situations where entity-sets in a relationship-set are *not* distinct

Example: a relationship-set named *works_for*, specifying employee/manager assignments

- ▣ Relationship involves two entities, and both are *employee* entities
- Roles are given names to distinguish entities
 - ▣ The relationship is a set of entities ordered by role:
(*manager*, *worker*)
 - ▣ First entity's role is named *manager*
 - ▣ Second entity's role is named *worker*

Relationships and Attributes

18

- Relationships can also have attributes!
 - ▣ Called descriptive attributes
 - ▣ They describe a particular relationship
 - ▣ They *do not* identify the relationship!
- Example:
 - ▣ The relationship between a software defect and an employee can have a *date_assigned* attribute
- Note: this distinction between entity attributes and relationship attributes is not made by relational model
 - ▣ Entity-relationship model is a higher level of abstraction than the relational model

Relationships and Attributes (2)

19

- Specific relationships are identified *only* by the participating entities
 - ▣ ...not by any relationship attributes!
 - ▣ Different relationships are allowed to have the same value for a descriptive attribute
 - ▣ (This is why entities in an entity-set must be uniquely identifiable.)
- Given:
 - ▣ Entity-sets A and B , both participating in a relationship-set R
- Specific entities $a \in A$ and $b \in B$ can only have one relationship instance in R
 - ▣ Otherwise, we would require more than just the participating entities to uniquely identify relationships

Degree of Relationship Set


20

- Most relationships in a schema are binary
 - ▣ Two entities are involved in the relationship
- Sometimes there are ternary relationships
 - ▣ Three entities are involved
 - ▣ Far less common, but still useful at times
- The number of entity-sets that participate in a relationship-set is called its degree
 - ▣ Binary relationship: $\text{degree} = 2$
 - ▣ Ternary relationship: $\text{degree} = 3$

Diagramming a Relationship-Set

21

Example: the *borrower* relationship-set between the *customer* and *loan* entity-sets



- Relationship-set is a diamond
 - ▣ Connected to participating entity-sets by solid lines
- Relationship-set can have descriptive attributes
 - ▣ Listed in another box, connected with a dotted-line

Attribute Structure

22

- Each attribute has a domain or value set
 - ▣ Values come from that domain or value set
- Simple attributes are atomic – they have no subparts
 - ▣ e.g. *amount* attribute is a single numeric value
- Composite attributes have subparts
 - ▣ Can refer to composite attribute as a whole
 - ▣ Can also refer to subparts individually
 - ▣ e.g. *address* attribute, composed of *street*, *city*, *state*, *postal_code* attributes

Attribute Cardinality

23

- Single-valued attributes only store one value
 - ▣ e.g. a *customer*'s *cust_id* only has one value
- Multi-valued attributes store zero or more values
 - ▣ e.g. a *customer* can have multiple *phone_number* values
 - ▣ A multi-valued attribute stores a set of values, not a multiset
 - ▣ Different *customer* entities can have different sets of phone numbers
 - ▣ Lower and upper bounds can be specified too
 - Can set upper bound on *phone_number* to 2

Attribute Source

24

- Base attributes (aka source attributes) are stored in the database
 - ▣ e.g. the *birth_date* of a *customer* entity
- Derived attributes are computed from other attributes
 - ▣ e.g. the *age* of a *customer* entity would be computed from their *birth_date*

Diagramming Attributes

25

- Example: Extend customers with more detailed info
- Composite attributes are shown as a hierarchy of values
 - ▣ Indented values are components of the higher-level value
 - ▣ e.g. *name* is comprised of *first_name*, *middle_initial*, and *last_name*

<i>customer</i>
<i>cust_id</i>
<i>name</i> <ul style="list-style-type: none"><i>first_name</i><i>middle_initial</i><i>last_name</i>
<i>address</i> <ul style="list-style-type: none"><i>street</i><i>city</i><i>state</i><i>zip_code</i>

Diagramming Attributes (2)

26

- Example: Extend customers with more detailed info
- Multivalued attributes are enclosed with curly-braces
 - ▣ e.g. each customer can have zero or more phone numbers

<i>customer</i>
<i>cust_id</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>
{ <i>phone_number</i> }

Diagramming Attributes (3)

27

- Example: Extend customers with more detailed info
- Derived attributes are indicated by a trailing set of parentheses ()
 - ▣ e.g. each customer has a base attribute recording their date of birth
 - ▣ Also a derived attribute that reports the customer's current age

<i>customer</i>
<i>cust_id</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>
{ <i>phone_number</i> }
<i>birth_date</i>
<i>age</i> ()

Representing Constraints

28

- E-R model can represent different kinds of constraints
 - ▣ Mapping cardinalities
 - ▣ Key constraints in entity-sets
 - ▣ Participation constraints
- Allows more accurate modeling of application's data requirements
 - ▣ Constrain design so that schema can only represent valid information
- Enforcing constraints can impact performance...
 - ▣ Still ought to specify them in the design!
 - ▣ Can always leave out constraints at implementation time

Mapping Cardinalities

29

- Mapping cardinality represents:
 - “How many other entities can be associated with an entity, via a particular relationship set?”
- Example:
 - ▣ How many *customer* entities can the *borrower* relationship associate with a single *loan* entity?
 - ▣ How many *loans* can *borrower* relationship associate with a single *customer* entity?
 - ▣ Specific answer depends on what is being modeled
- Also known as the cardinality ratio
- Easiest to reason about with binary relationships

Mapping Cardinalities (2)


30

Given:

- Entity-sets A and B
- Binary relationship-set R associating A and B

One-to-one mapping (1:1)

- An entity in A is associated with *at most* one entity in B
- An entity in B is associated with *at most* one entity in A



Mapping Cardinalities (3)


31

One-to-many mapping (1:M)

- An entity in *A* is associated with *zero or more* entities in *B*
- An entity in *B* is associated with *at most one* entity in *A*

Many-to-one mapping (M:1)

- Opposite of one-to-many
- An entity in *A* is associated with *at most one* entity in *B*
- An entity in *B* is associated with *zero or more* entities in *A*




Mapping Cardinalities (4)

32

Many-to-many mapping

- An entity in A is associated with *zero or more* entities in B
- An entity in B is associated with *zero or more* entities in A



Mapping Cardinalities (5)

33

- Which mapping cardinality is most appropriate for a given relationship?
 - ▣ Answer depends on what you are trying to model!
 - ▣ Could just use many-to-many relationships everywhere, but that would be dumb.
- Goal:
 - ▣ Constrain the mapping cardinality to most accurately reflect what should be allowed
 - ▣ Database can enforce these constraints automatically
 - ▣ Good schema design reduces or eliminates the *possibility* of storing bad data

Example: *borrower* relationship between *customer* and *loan*

34

One-to-one mapping:

- ▣ Each customer can have only one loan
- ▣ Customers can't share loans
(e.g. with spouse or business partner)

One-to-many mapping:

- ▣ A customer can have multiple loans
- ▣ Customers still can't share loans

Many-to-one mapping:

- ▣ Each customer can have only one loan
- ▣ Customers can share loans

Many-to-many mapping:

- ▣ A customer can have multiple loans
- ▣ Customers can share loans too


Best choice for *borrower* :
many-to-many mapping

Handles real-world needs!

Diagramming Cardinalities

35


- In relationship-set diagrams:
 - ▣ an arrow towards an entity represents “one”
 - ▣ a simple line represents “many”
 - ▣ arrow is *always* towards the entity
- Many-to-many mapping between *customer* and *loan*:



Diagramming Cardinalities (2)

36

- One-to-many mapping between *customer* and *loan*:




- ▣ Each customer can have multiple loans
- ▣ A loan is owned by exactly one customer
 - (Actually, this is technically “at most one”. Participation constraints will allow us to say “exactly one.”)

Diagramming Cardinalities (3)

37

- One-to-one mapping between *customer* and *loan*:



- Each customer can have only one loan
- A loan is owned by exactly one customer