

ADVANCED E-R FEATURES

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 17

Extensions to E-R Model

2

- Basic E-R model is good for many uses
- Several extensions to the E-R model for more advanced modeling
 - ▣ Generalization and specialization
 - ▣ Aggregation
- These extensions can also be converted to the relational model
 - ▣ Introduces a few more design choices
- Will only discuss specialization today
 - ▣ See book §7.8.5 for details on aggregation (material will be included with Assignment 5 too)

Specialization

3

- An entity-set might contain distinct subgroups of entities
 - ▣ Subgroups have some different attributes, not shared by the entire entity-set
- E-R model provides specialization to represent such entity-sets
- Example: bank account categories
 - ▣ Checking accounts
 - ▣ Savings accounts
 - ▣ Have common features, but also unique attributes

Generalization and Specialization

4

- Generalization: a “bottom up” approach
 - ▣ Taking similar entity-sets and unifying their common features
 - ▣ Start with specific entities, then create generalizations from them
- Specialization: a “top down” approach
 - ▣ Creating general purpose entity-sets, then providing specializations of the general idea
 - ▣ Start with the general notion, then refine it
- Terms are basically equivalent
 - ▣ Book refers to generalization as the overarching concept

Bank Account Example

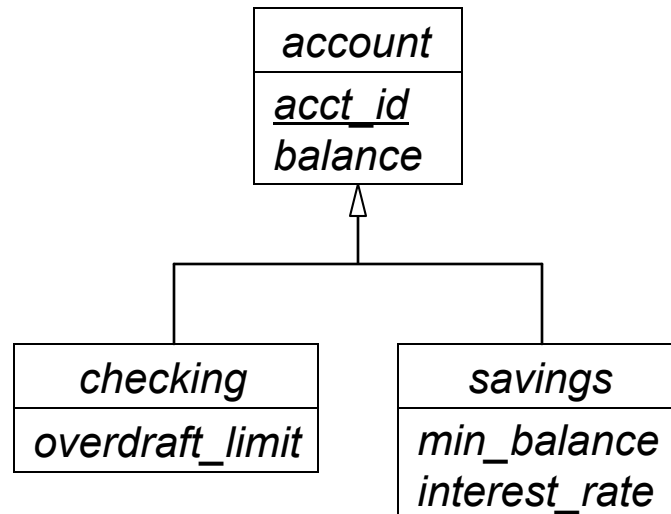
5

- Checking and savings accounts both have:
 - ▣ account number
 - ▣ balance
 - ▣ owner(s)
- Checking accounts also have:
 - ▣ overdraft limit and associated overdraft account
 - ▣ check transactions
- Savings accounts also have:
 - ▣ minimum balance
 - ▣ interest rate

Bank Account Example (2)

6

- Create entity-set to represent common attributes
 - ▣ Called the superclass, or higher-level entity-set
- Create entity-sets to represent specializations
 - ▣ Called subclasses, or lower-level entity-sets
- Join superclass to subclasses with hollow-head arrow(s)



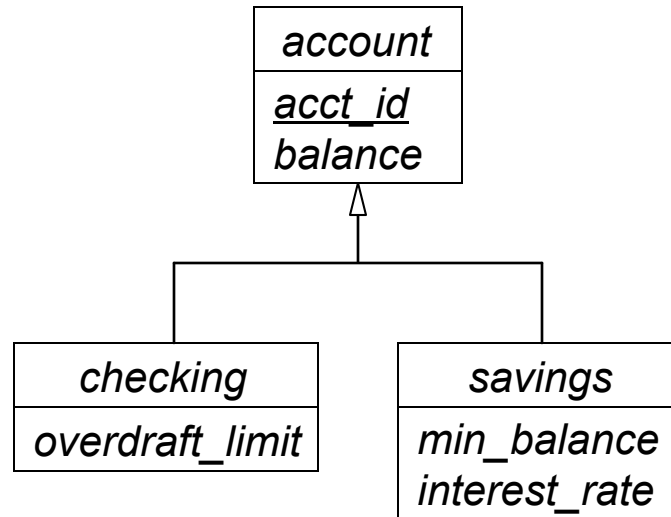
Inheritance

7

- Attributes of higher-level entity-sets are inherited by lower-level entity-sets
- Relationships involving higher-level entity-sets are also inherited by lower-level entity-sets!
 - ▣ Lower-level entity-sets can also participate in *their* own relationship-sets, separate from higher-level entity-set
- Usually, entity-sets inherit from one superclass
 - ▣ Entity-sets form a hierarchy
- Can also inherit from multiple superclasses
 - ▣ Entity-sets form a lattice
 - ▣ Introduces many subtle issues, of course

Specialization Constraints

8



- Can an account be both a savings account and a checking account?
- Can an account be neither a savings account nor a checking account?
- Can specify constraints on specialization
 - ▣ Enforce what “makes sense” for the enterprise

Disjointness Constraints


9

- “An account cannot be *both* a checking account and a savings account.”
- An entity may belong to at most one of the lower-level entity-sets
 - ▣ Must be a member of *checking*, or a member of *savings*, but not both!
 - ▣ Called a “disjointness constraint”
 - ▣ A better way to state it: a disjoint specialization
- If an entity can be a member of multiple lower-level entity-sets:
 - ▣ Called an overlapping specialization

Disjointness Constraints (2)

10

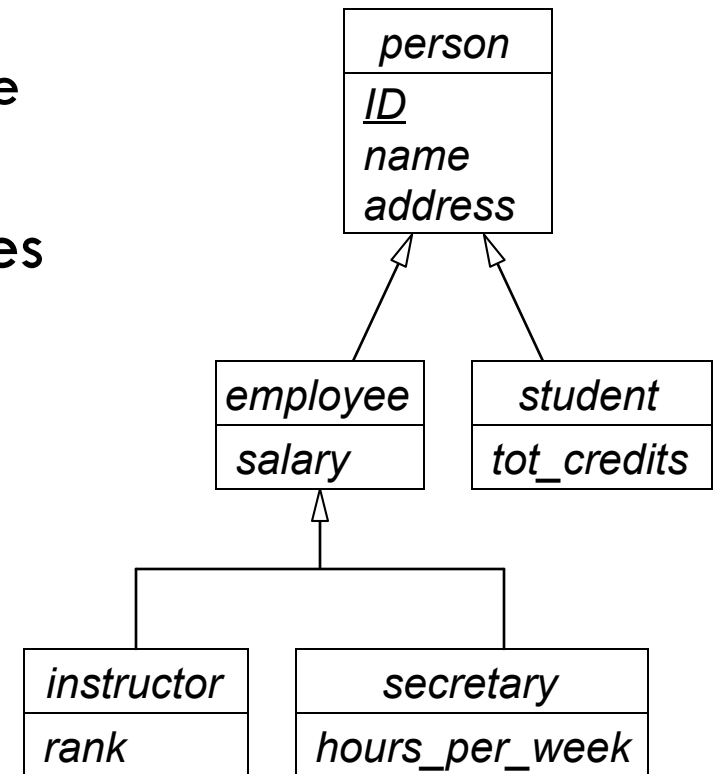
- How the arrows are drawn indicates whether the specialization is disjoint or overlapping
- Bank account example:
 - ▣ One arrow split into multiple parts indicates a disjoint specialization
 - ▣ An account may only be a checking account, or a savings account, not both



Disjointness Constraints (3)

11

- Another example from the book:
 - ▣ Specialization hierarchy for people at a university
- Multiple separate arrows indicates an overlapping specialization
 - ▣ A person can be an employee of the university and a student
- One arrow split into multiple parts is a disjoint specialization
 - ▣ An employee can be an instructor or a secretary, but not both



Completeness Constraints

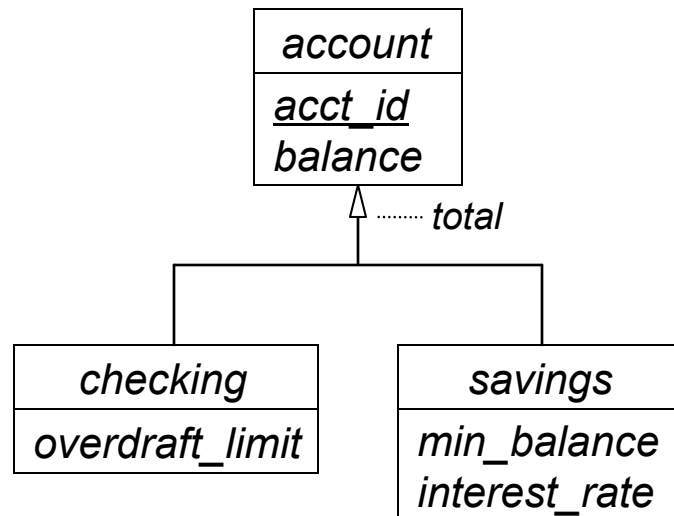
12

- “An account must be a checking account, or it must be a savings account.”
- Every entity in higher-level entity-set must also be a member of at least one lower-level entity-set
 - ▣ Called total specialization
- If entities in higher-level entity-set aren’t required to be members of lower-level entity-sets:
 - ▣ Called partial specialization
- *account* specialization is a total specialization

Completeness Constraints (2)

13


- Default constraint is partial specialization
- Specify total specialization constraint by annotating the specialization arrow(s)
- Updated bank account diagram:



Completeness Constraints (3)

14


- Same approach with overlapping specialization
- Example: people at a university
 - ▣ Every person is an employee or a student
 - ▣ Not every employee is an instructor or a secretary
- Annotate arrows pointing to person with “total” to indicate total specialization
 - ▣ Every person must be an employee, a student, or both



Account Types?

15

- Our bank schema so far:



- How to tell whether an account is a checking account or a savings account?
 - ▣ No attribute indicates type of account

Membership Constraints

16

- Membership constraints specify which lower-level entity-sets each entity is a member of
 - ▣ e.g. which accounts are checking or savings accounts
- Condition-defined lower-level entity-sets
 - ▣ Membership is specified by a predicate
 - ▣ If an entity satisfies a lower-level entity-set's predicate then it is a member of that lower-level entity-set
 - ▣ If *all* lower-level entity-sets refer to the same attribute, this is called attribute-defined specialization
 - e.g. *account* could have an *account_type* attribute set to “c” for checking, or “s” for savings

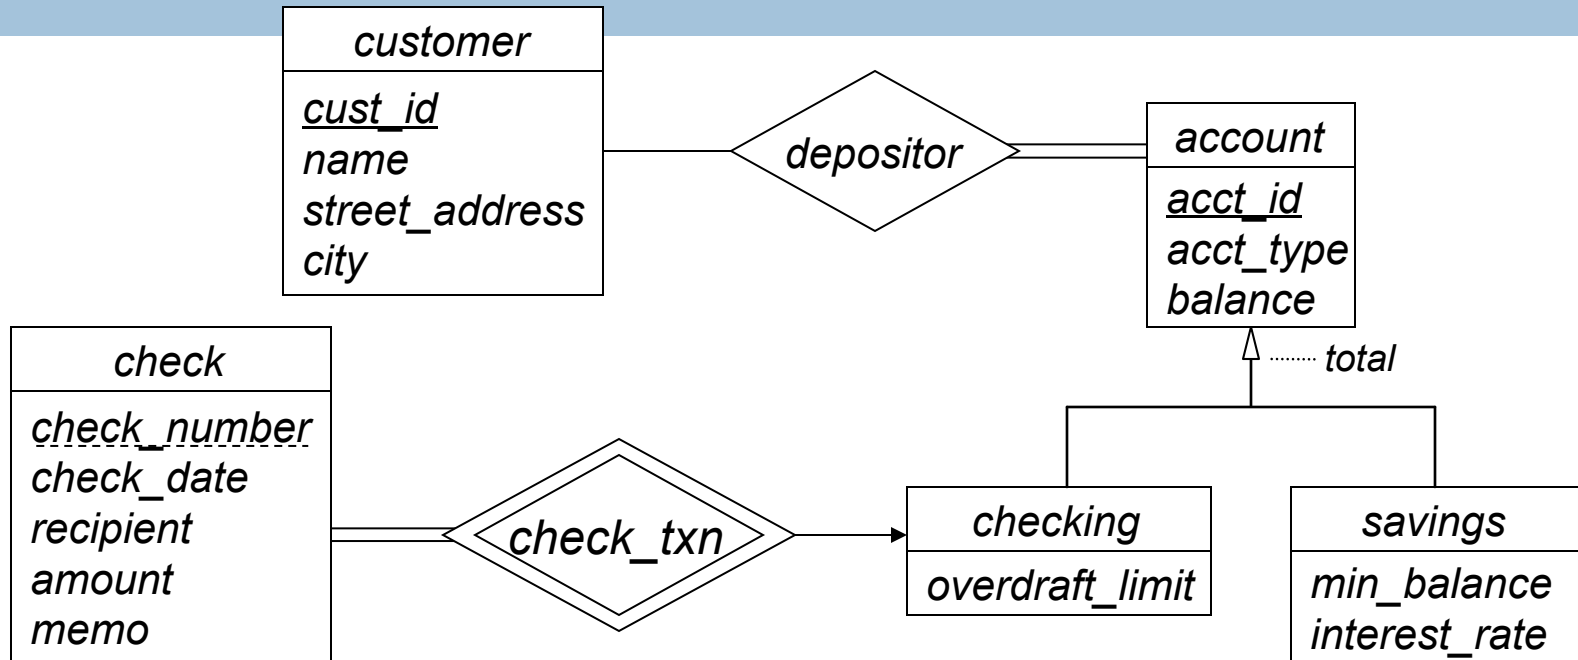
Membership Constraints (2)

17

- Entities may simply be assigned to lower-level entity-sets by a database user
 - ▣ No explicit predicate governs membership
 - ▣ Called user-defined membership
- Generally used when an entity's membership could change in the future

Final Bank Account Diagram

18



- Would also create relationship-sets against various entity-sets in hierarchy
 - ▣ associate *customer* with *account*
 - ▣ associate *check* weak entity-set with *checking*


Mapping to Relational Model

19

- Mapping generalization/specialization to relational model is straightforward
- Create relation schema for higher-level entity-set
 - ▣ Including primary keys, etc.
- Create schemas for lower-level entity-sets
 - ▣ Subclass schemas include superclass' primary key attributes!
 - ▣ Primary key is same as superclass' primary key
 - Subclasses can also contain their own candidate keys!
 - Enforce these candidate keys in implementation schema
 - ▣ Foreign key reference from subclass schemas to superclass schema, on primary-key attributes

Mapping Bank Account Schema

20



□ Schemas:

account(*acct_id*, *acct_type*, *balance*)

checking(*acct_id*, *overdraft_limit*)

savings(*acct_id*, *min_balance*, *interest_rate*)

- ▣ Could use **CHECK** constraints on SQL tables for membership constraints, other constraints (although it may be expensive)

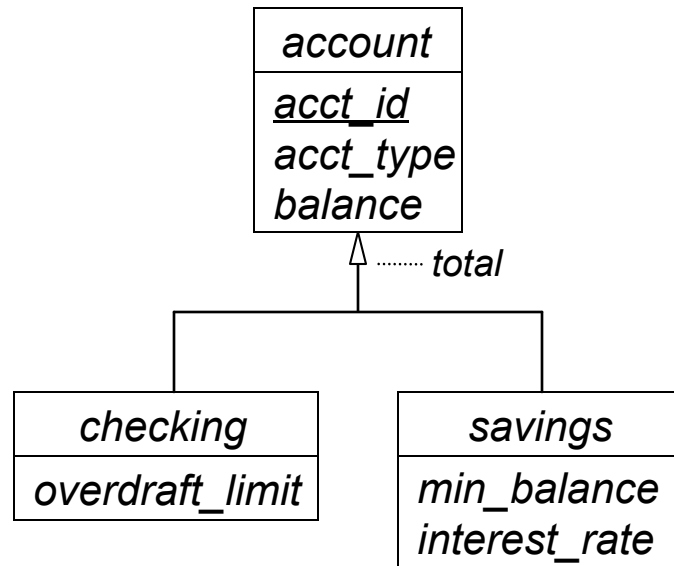
Alternative Schema Mapping

21

- If specialization is disjoint and complete, could convert only lower-level entity-sets to relational schemas
 - ▣ Every entity in higher-level entity-set also appears in lower-level entity-sets
 - ▣ Every entity is a member of *exactly one* lower-level entity-set
- Each lower-level entity-set has its own relation schema
 - ▣ All attributes of superclass entity-set are included on each subclass entity-set
 - ▣ No relation schema for superclass entity-set

Alternative Account Schema

22



□ Schemas, take 2:

checking(*acct_id*, *acct_type*, *balance*, *overdraft_limit*)

savings(*acct_id*, *acct_type*, *balance*, *min_balance*, *interest_rate*)

Alternative Account Schema (2)

23

- Alternative schemas:

checking(acct_id, acct_type, balance, overdraft_limit)

savings(acct_id, acct_type, balance, min_balance, interest_rate)

- Problems?

- ▣ Enforcing uniqueness of account IDs!

- ▣ Representing relationships involving both kinds of accounts

- Can solve by creating a simple relation:

account(acct_id)

- ▣ Contains *all* valid account IDs

- ▣ Relationships involving accounts can use *account*

- ▣ Need foreign key constraints again...

Generating Primary Keys

24

- Generating primary key values is actually the easy part
- Most databases provide sequences
 - ▣ A source of unique, increasing **INTEGER** or **BIGINT** values
 - ▣ Perfect for primary key values
 - ▣ Multiple tables can use the same sequence for their primary keys

- PostgreSQL example:

```
CREATE SEQUENCE acct_seq;
```

```
CREATE TABLE checking (  
    acct_id INT PRIMARY KEY DEFAULT nextval('acct_seq');  
    ...  
);
```

```
CREATE TABLE savings (  
    acct_id INT PRIMARY KEY DEFAULT nextval('acct_seq');  
    ...  
);
```



Alternative Schema Mapping

25

- Alternative mapping has serious drawbacks
 - ▣ Doesn't actually give many benefits in general case
- Fewer drawbacks if:
 - ▣ Total, disjoint specialization
 - ▣ No relationships against superclass entity-set
- If specialization is overlapping, some details will be stored multiple times
 - ▣ Unnecessary redundancy, and consistency issues
- Also limits future schema changes
 - ▣ Should always think about this when creating schemas

Recap: Weak Entity-Set Example

26




- *account* schema:
account(*account_number*, *balance*)
- *check* schema:
 - ▣ Discriminator is *check_number*
 - ▣ Primary key for *check* is: (*account_number*, *check_number*)*check*(*account_number*, *check_number*, *check_date*,
recipient, *amount*, *memo*)

Schema Combination

27


- Relationship between weak entity-set and strong entity-set doesn't need represented separately
 - ▣ Many-to-one relationship
 - ▣ Weak entity-set has total participation
 - ▣ Weak entity-set's schema already captures the identifying relationship
- Can apply this technique to other relationship-sets:
 - ▣ One-to-many mapping, with total participation on the “many” side



Schema Combination (2)

28

- Entity-sets A and B , relationship-set AB
 - ▣ Many-to-one mapping from A to B
 - ▣ A 's participation in AB is total
- Generates relation schemas A , B , AB
 - ▣ Primary key of A is $primary_key(A)$
 - ▣ Primary key of AB is also $primary_key(A)$
 - (A is on “many” side of mapping)
 - ▣ AB has foreign key constraints on both A and B
 - ▣ There is one relationship in AB for every entity in A
- Can combine A and AB relation schemas
 - ▣ Primary key of combined schema still $primary_key(A)$
 - ▣ Only requires one foreign-key constraint, to B




Schema Combination (3)

29


- In this case, when relationship-set is combined into the entity-set, the entity-set's primary key *doesn't change!*

- If A 's participation in AB is partial, can still combine schemas



- Must store *null* values for *primary_key(B)* attributes when an entity in A maps to no entity in B


- If AB is one-to-one mapping:



- Can also combine schemas in this case
- Could incorporate AB into schema for A , or schema for B
- Don't forget that AB has two candidate keys...
 - The combined schema must still enforce both candidate keys

Schema-Combination Example


30



- ❑ Manager to worker mapping is one-to-many
- ❑ Relation schemas were:
 `employee(employee_id, name)`
 `works_for(employee_id, manager_id)`
- ❑ Could combine into:
 `employee(employee_id, name, manager_id)`
 - ▣ (A very common schema combination)
 - ▣ Need to store *null* for employees with no manager

Schema Combination Example (2)


31



- One-to-one mapping between customers and loans
customer(*cust_id*, *name*, *street_address*, *city*)
loan(*loan_id*, *amount*)
borrower(*cust_id*, *loan_id*) – *loan_id* also a candidate key
- Could combine *borrower* schema into *customer* schema or *loan* schema
 - ▣ Does it matter which one you choose?

Schema Combination Example (3)


32



- Participation of *loan* in *borrower* will be total
 - ▣ Combining *borrower* into *customer* would require *null* values for customers without loans
- Better to combine *borrower* into *loan* schema
 - customer*(*cust_id*, *name*, *street_address*, *city*)
 - loan*(*loan_id*, *cust_id*, *amount*)
 - ▣ No *null* values!

Schema Combination Example (4)

33



- Schema:
 - customer*(cust_id, name, street_address, city)
 - loan*(loan_id, cust_id, amount)
- What if, after a while, we wanted to change the mapping cardinality?
 - ▣ Schema changes would be significant
 - ▣ Would need to migrate existing data to a new schema

Schema Combination Notes

34

- Benefits of schema combination:
 - ▣ Usually eliminates one foreign-key constraint, and the associated performance impact
 - Constraint enforcement
 - Extra join operations in queries
 - ▣ Reduces storage requirements
- Drawbacks of schema combination:
 - ▣ May necessitate the use of *null* values to represent the absence of relationships
 - ▣ Makes it harder to change mapping cardinality constraints in the future