

NORMAL FORMS

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 18

Equivalent Schemas

2

- Many different schemas can represent a set of data
 - ▣ Which one is best?
 - ▣ What does “best” even mean?
- Main goals:
 - ▣ Representation must be complete
 - ▣ Data should not be unnecessarily redundant
 - ▣ Should be easy to manipulate the information
 - ▣ Should be easy to enforce [most] constraints

Normal Forms

3

- A “good” pattern for database schemas to follow is called a normal form
- Several different normal forms, with different constraints
- Normal forms can be formally specified
 - ▣ Can test a schema against a normal form
 - ▣ Can transform a schema into a normal form
- Goal:
 - ▣ Design schemas that satisfy a particular normal form
 - ▣ If a schema isn’t “good,” transform it into an appropriate normal form

Example Schema Design

4

- Schema for representing loans and borrowers:
 - ▣ *customer* relation stores customer details, including a *cust_id* primary-key attribute
 - ▣ *loan*(*loan_id*, *amount*)
 - ▣ *borrower*(*cust_id*, *loan_id*)
- Many-to-many mapping
 - ▣ A customer can have multiple loans
 - ▣ A loan can be owned by multiple customers

loan_id	amount
...	...
L-100	10000
...	...

loan

cust_id	loan_id
...	...
23-652	L-100
15-202	L-100
23-521	L-100
...	...

borrower

Larger Schema?

5

- Could replace *loan* and *borrower* relations with a larger, combined relation

bor_loan(*cust_id*, *loan_id*, *amount*)

cust_id	loan_id	amount
...
23-652	L-100	10000
15-202	L-100	10000
23-521	L-100	10000
...

bor_loan

- Rationale:
 - ▣ Eliminates a join when retrieving loan amounts
- Problem: mapping between customers and loans is many-to-many
 - ▣ Multiple redundant copies of *amount* to keep in sync!

Repeated Values

6

- How do we *know* that this is a problem?
 - ▣ “Because we see values that appear multiple times”
 - ▣ *This isn’t a good enough reason!!!*
 - ▣ Could easily have different loans with the same amount
- A repeated value doesn’t *automatically* indicate a problem...

cust_id	loan_id	amount
...
23-652	L-100	10000
19-065	L-205	10000
15-202	L-100	10000
23-521	L-100	10000
20-419	L-205	10000
...

bor_loan

Back to the Enterprise

7

- What are the rules of the enterprise that we are modeling?
 - ▣ “Every loan must have only one amount.”
- In other words:
 - ▣ Every loan ID corresponds to exactly one amount.
 - ▣ If there were a schema (*loan_id*, *amount*) then *loan_id* can be a primary key.
- Specified as a functional dependency
 - ▣ *loan_id* \rightarrow *amount*
 - ▣ *loan_id* functionally determines *amount*

Repeated Values v2.0

8

- *bor_loan* relation has both *loan_id* and *amount* attributes
bor_loan(*cust_id*, *loan_id*, *amount*)
- But, *loan_id* → *amount*, and *loan_id* by itself can't be a primary key in *bor_loan*
 - ▣ Need to support many-to-many mappings between customers and loans
 - ▣ Combination of *cust_id* and *loan_id* must be a primary key, so a particular *loan_id* value can appear multiple times
- In rows with the same *loan_id* value, *amount* will have to be repeated.

Functional Dependencies

9

- Functional dependencies are very important in schema analysis
 - ▣ Have a *lot* to do with keys!
 - ▣ “Good” schema designs are guided by functional dependencies
 - ▣ Frequently helpful to identify them during schema design
- Can formally define functional dependencies, and reason about them
- Can also specify constraints on schemas using functional dependencies

Another Example Schema

10

- A “large” schema for employee information

employee(emp_id, emp_name, phone, title, salary, start_date)

emp_id	emp_name	phone	title	salary	start_date
...
123-45-6789	Jeff	555-1234	CTO	120000	1996-03-15
314-15-9265	Mary	555-3141	CFO	120000	1997-08-02
987-65-4321	Helen	555-9876	Developer	90000	1996-05-23
101-01-0101	Marcus	555-1010	Tester	70000	1995-11-04
...

employee

- Employee ID is unique, but other attributes could have duplicate values

Smaller Schemas?

11

- Could represent this with two smaller schemas:

emp_ids(emp_id, emp_name)

emp_details(emp_name, phone, title, salary, start_date)

emp_id	emp_name
...	...
123-45-6789	Jeff
314-15-9265	Mary
987-65-4321	Helen
101-01-0101	Marcus
...	...

emp_ids

emp_name	phone	title	salary	start_date
...
Jeff	555-1234	CTO	120000	1996-03-15
Mary	555-3141	CFO	120000	1997-08-02
Helen	555-9876	Developer	90000	1996-05-23
Marcus	555-1010	Tester	70000	1995-11-04
...

emp_details

- Generate original employee data with a join:
emp_ids ⋈ *emp_details*
- Any problems with this?

emp_name is not unique!

12


- Joins using *emp_name* can generate invalid tuples!

emp_id	emp_name
...	...
314-15-9265	Mary
161-80-3398	Mary
...	...

emp_ids

emp_name	phone	title	salary	start_date
...
Mary	555-3141	CFO	120000	1997-08-02
Mary	555-1618	Gofer	25000	1998-01-07
...

emp_details



emp_ids ⋈ *emp_details*

emp_id	emp_name	phone	title	salary	start_date
...
314-15-9265	Mary	555-1618	Gofer	25000	1998-01-07
314-15-9265	Mary	555-3141	CFO	120000	1997-08-02
161-80-3398	Mary	555-3141	CFO	120000	1997-08-02
161-80-3398	Mary	555-1618	Gofer	25000	1998-01-07
...

Bad Decompositions

13

- This decomposition is clearly broken
 - ▣ It can't represent the information correctly!
- Problem: enterprise needs to support different employees with the same name
- Lossy decompositions cannot accurately represent all facts about an enterprise
- Lossless decompositions can accurately represent all facts
- “Good” schema designs will avoid lossy decompositions

First Normal Form

14

- A schema is in first normal form (1NF) if all attribute domains are atomic
 - ▣ An atomic domain has values that are indivisible units
- E-R model supports non-atomic attributes
 - ▣ Multivalued attributes
 - ▣ Composite attributes
- Relational model specifies atomic domains for attributes
 - ▣ Schemas are automatically in 1NF
 - ▣ Mapping from E-R model to relational model changes composite/multivalued attributes into an atomic form

1NF Example

15

- E-R diagram for magazine subscribers

- *address* is composite
- *email_addr* is multivalued

<i>subscriber</i>
<u><i>sub_id</i></u>
{ <i>email_addr</i> }
<i>address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zip_code</i>

- Converts to a 1NF schema:

subscriber(*sub_id*, *street*, *city*, *state*, *zip_code*)

sub_emails(*sub_id*, *email_addr*)

- The conversion rules we have discussed, *automatically* convert E-R schemas into 1NF

1 NF and Non-Atomic Attributes

16

- Many, but not all, SQL DBs have non-atomic types
 - ▣ Some offer support for composite attributes
 - ▣ Some offer support for multivalued attributes
 - ▣ These are SQL extensions – not portable
- As long as you steer clear of using non-atomic attributes in primary/foreign keys, can sometimes be quite useful
 - ▣ Will likely encounter them very rarely in practice, though
 - ▣ Biggest reason: DB support for list/vector column-types isn't terribly widespread, or always very easy to use

1 NF and Non-Atomic Attributes (2)

17

- Composite types:
 - ▣ e.g. defining an “address” composite type
 - ▣ Can definitely be useful for making a schema clearer, as long as they aren’t used in a key!
- Multivalued types:
 - ▣ e.g. arrays, lists, sets, vectors
 - ▣ Can sometimes be useful for storing pre-computed values that aren’t expected to change frequently
 - ▣ If you are regularly issuing queries that search through or change these values, you may need to revise your schema!
 - Should probably factor non-atomic data out into a separate table

Other Normal Forms

18

- Other normal forms relate to functional dependencies
- Analysis of functional dependencies shows if a schema needs decomposed
- Keys are functional dependencies too!
- Formally define functional dependencies, and reason about them
- Define normal forms in terms of functional dependencies

Schemas and Constraints

19

- Keys and functional dependencies are constraints that a database must satisfy
 - Legal relations satisfy the required constraints
 - Relation doesn't contain any tuples that violate the specified constraints
- More terminology:
 - Relation schema R , relation $r(R)$
 - A set of functional dependencies F
 - Relation r satisfies F if r is legal
 - When we say “ F holds on R ”, specifies the set of relations with R as their schema, that are legal with respect to F

Functional Dependencies

20

- Formal definition of a functional dependency:
 - ▣ Given a relation schema R with attribute-sets $\alpha, \beta \subseteq R$
 - ▣ The functional dependency $\alpha \rightarrow \beta$ holds on $r(R)$ if
$$\langle \forall t_1, t_2 \in r : t_1[\alpha] = t_2[\alpha] : t_1[\beta] = t_2[\beta] \rangle$$
- In other words:
 - ▣ For all pairs of tuples t_1 and t_2 in r ,
if $t_1[\alpha] = t_2[\alpha]$ then $t_1[\beta] = t_2[\beta]$
 - ▣ α functionally determines β

Dependencies and Superkeys

21

- Given relation schema R , a subset K of R can be a superkey
 - ▣ In a relation $r(R)$, no two tuples can share the same values for attributes in K
- Can also say: K is a superkey if $K \rightarrow R$
 - ▣ The functional dependency $K \rightarrow R$ holds if
$$\langle \forall t_1, t_2 \in r(R) : t_1[K] = t_2[K] : t_1[R] = t_2[R] \rangle$$
 - ▣ $t_1[R] = t_2[R]$ (or $t_1 = t_2$) means t_1 and t_2 are the same tuple
 - ▣ The superkey K functionally determines the whole relation R
- Functional dependencies are a more general form of constraint than superkeys are.

The *bor_loan* Relation

22

- *bor_loan*(*cust_id*, *loan_id*, *amount*)
 - ▣ Functional dependency: $\textit{loan_id} \rightarrow \textit{amount}$
 - ▣ “Every loan has exactly one amount.”
 - ▣ Every tuple in *bor_loan* with a given *loan_id* value must have the same *amount* value
- *bor_loan* also has a primary key
 - ▣ Specifies another functional dependency
 - ▣ $\textit{cust_id}, \textit{loan_id} \rightarrow \textit{cust_id}, \textit{loan_id}, \textit{amount}$
 - ▣ This is not a functional dependency *specifically required* by what the enterprise needs to model
 - Can be inferred from other functional dependencies in the schema

Trivial Dependencies

23

- A trivial functional dependency is satisfied by *all* relations!
 - ▣ For a relation R containing attributes A and B ,
 $A \rightarrow A$ is a trivial dependency
 $\langle \forall t_1, t_2 \in r : t_1[A] = t_2[A] : t_1[A] = t_2[A] \rangle$
 - Well, duh!
 - ▣ $AB \rightarrow A$ is also a trivial dependency
 - If $t_1[AB] = t_2[AB]$, then of course $t_1[A] = t_2[A]$ too!
- In general: $\alpha \rightarrow \beta$ is trivial if $\beta \subseteq \alpha$

Closure

24

- Given a set of functional dependencies, we can infer other dependencies
 - ▣ Given relation schema $R(A, B, C)$
 - ▣ If $A \rightarrow B$ and $B \rightarrow C$, holds on R , then $A \rightarrow C$ also holds on R
- Given a set of functional dependencies F
 - ▣ F^+ denotes the closure of F
 - ▣ F^+ includes F , and all dependencies that can be inferred from F . ($F \subseteq F^+$)

Boyce-Codd Normal Form

25

- Eliminates all redundancy that can be discovered using functional dependencies
- Given:
 - ▣ Relation schema R
 - ▣ Set of functional dependencies F
- R is in BCNF with respect to F if:
 - ▣ For all functional dependencies $\alpha \rightarrow \beta$ in F^+ , where $\alpha \in R$ and $\beta \in R$, at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial dependency
 - α is a superkey for R
- A database design is in BCNF if all schemas in the design are in BCNF

BCNF Examples

26

- The *bor_loan* schema isn't in BCNF

bor_loan(*cust_id*, *loan_id*, *amount*)

- ▣ *loan_id* → *amount* holds on *bor_loan*

- ▣ This is not a trivial dependency, and *loan_id* isn't a superkey for *bor_loan*

- The *borrower* and *loan* schemas are in BCNF

borrower(*cust_id*, *loan_id*)

- ▣ No nontrivial dependencies hold

loan(*loan_id*, *amount*)

- ▣ *loan_id* → *amount* holds on *loan*

- ▣ *loan_id* is the primary key of *loan*

BCNF Decomposition

27

- If R is a schema not in BCNF:
 - ▣ There is at least one nontrivial functional dependency $\alpha \rightarrow \beta$ such that α is not a superkey for R
- Replace R with two schemas:
 - $(\alpha \cup \beta)$
 - $(R - (\beta - \alpha))$
 - (stated this way in case α and β overlap; usually they don't)
- *The new schemas might also not be in BCNF!*
 - ▣ Repeat this decomposition process until all schemas are in BCNF

Undoing the Damage

28

- For *bor_loan*, $\alpha = \text{loan_id}$, $\beta = \text{amount}$
 $R = (\text{cust_id}, \text{loan_id}, \text{amount})$
 $(\alpha \cup \beta) = (\text{loan_id}, \text{amount})$
 $(R - (\beta - \alpha)) = (\text{cust_id}, \text{loan_id})$
- Rules successfully decompose *bor_loan* back into *loan* and *borrower* schemas

Review

29

- Normal forms are guidelines for what makes a database design “good”
 - ▣ Can formally specify them
 - ▣ Can transform schemas into normal forms
- Functional dependencies specify constraints between attributes in a schema
 - ▣ A more general kind of constraint than key constraints
- Covered 1NF and BCNF
 - ▣ 1NF requires all attributes to be atomic
 - ▣ BCNF uses functional dependencies to eliminate redundant data

Next Time!

30

- A big question to explore:
 - ▣ Given a set of functional dependencies F , we need to know what dependencies can be inferred from it!
 - i.e. given F , how to compute F^+
 - ▣ BCNF needs this information, as do other normal forms
- Does Boyce-Codd Normal Form have drawbacks?
 - ▣ (yes.)
 - ▣ Motivates the development of 3rd Normal Form