

FUNCTIONAL DEPENDENCY THEORY II

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 20

Last Time: Canonical Cover

2

- Last time, introduced concept of canonical cover
- A canonical cover F_c for F is a set of functional dependencies such that:
 - ▣ F logically implies all dependencies in F_c
 - ▣ F_c logically implies all dependencies in F
 - ▣ Can't infer any functional dependency in F_c from other dependencies in F_c
 - ▣ No functional dependency in F_c contains an extraneous attribute
 - ▣ Left side of all functional dependencies in F_c are unique
 - There are no two dependencies $\alpha_1 \rightarrow \beta_1$ and $\alpha_2 \rightarrow \beta_2$ in F_c such that $\alpha_1 = \alpha_2$

Extraneous Attributes

3

- Given a set F of functional dependencies
 - ▣ An attribute in a functional dependency is extraneous if it can be removed from F without affecting closure of F
- Formally: given F , and $\alpha \rightarrow \beta$
 - ▣ If $A \in \alpha$, and F logically implies $(F - \{\alpha \rightarrow \beta\}) \cup \{(\alpha - A) \rightarrow \beta\}$, then A is extraneous
 - ▣ If $A \in \beta$, and $(F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ logically implies F , then A is extraneous
 - i.e. generate a new set of functional dependencies F' by replacing $\alpha \rightarrow \beta$ with $\alpha \rightarrow (\beta - A)$
 - See if F' logically implies F

Testing Extraneous Attributes

4

- Given relation schema R , and a set F of functional dependencies that hold on R
- Attribute A in $\alpha \rightarrow \beta$
- If $A \in \alpha$ (i.e. A is on left side of the dependency), then let $\gamma = \alpha - \{A\}$
 - ▣ See if $\gamma \rightarrow \beta$ can be inferred from F
 - ▣ Compute γ^+ under F
 - ▣ If $\beta \subseteq \gamma^+$ then A is extraneous in α

Testing Extraneous Attributes (2)

5

- Given relation schema R , and a set F of functional dependencies that hold on R
- Attribute A in $\alpha \rightarrow \beta$
- If $A \in \beta$ (on right side of the dependency), then try the altered set F'
 - ▣ $F' = (F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$
 - ▣ See if $\alpha \rightarrow A$ can be inferred from F'
 - ▣ Compute α^+ under F'
 - ▣ If α^+ includes A then A is extraneous in β

Computing Canonical Cover

6

- A simple way to compute the canonical cover of F

$$F_c = F$$

repeat

 apply union rule to replace dependencies in F_c of form

$$\alpha_1 \rightarrow \beta_1 \text{ and } \alpha_1 \rightarrow \beta_2 \text{ with } \alpha_1 \rightarrow \beta_1\beta_2$$

 find a functional dependency $\alpha \rightarrow \beta$ in F_c with an
 extraneous attribute

 /* Use F_c for the extraneous attribute test, not F !!! */

 if an extraneous attribute is found, delete it from $\alpha \rightarrow \beta$

until F_c stops changing

Canonical Cover Example

7

- Functional dependencies F on schema (A, B, C)
 - ▣ $F = \{ A \rightarrow BC, B \rightarrow C, A \rightarrow B, AB \rightarrow C \}$
 - ▣ Find F_c
- Apply union rule to $A \rightarrow BC$ and $A \rightarrow B$
 - ▣ Left with: $\{ A \rightarrow BC, B \rightarrow C, AB \rightarrow C \}$
- A is extraneous in $AB \rightarrow C$
 - ▣ $B \rightarrow C$ is logically implied by F (obvious)
 - ▣ Left with: $\{ A \rightarrow BC, B \rightarrow C \}$
- C is extraneous in $A \rightarrow BC$
 - ▣ Logically implied by $A \rightarrow B, B \rightarrow C$
- $F_c = \{ A \rightarrow B, B \rightarrow C \}$

Canonical Covers

8

- A set of functional dependencies can have multiple canonical covers
- Example:
 - ▣ $F = \{ A \rightarrow BC, B \rightarrow AC, C \rightarrow AB \}$
 - ▣ Has several canonical covers:
 - $F_c = \{ A \rightarrow B, B \rightarrow C, C \rightarrow A \}$
 - $F_c = \{ A \rightarrow B, B \rightarrow AC, C \rightarrow B \}$
 - $F_c = \{ A \rightarrow C, C \rightarrow B, B \rightarrow A \}$
 - $F_c = \{ A \rightarrow C, B \rightarrow C, C \rightarrow AB \}$
 - $F_c = \{ A \rightarrow BC, B \rightarrow A, C \rightarrow A \}$

Another Example

9

- Functional dependencies F on schema (A, B, C, D)
 - $F = \{ A \rightarrow B, BC \rightarrow D, AC \rightarrow D \}$
 - Find F_c
- In this case, it may look like $F_c = F \dots$
- However, can infer $AC \rightarrow D$ from $A \rightarrow B, BC \rightarrow D$ (pseudotransitivity), so $AC \rightarrow D$ is extraneous in F
 - Therefore, $F_c = \{ A \rightarrow B, BC \rightarrow D \}$
- Alternately, can argue that D is extraneous in $AC \rightarrow D$
 - With $F' = \{ A \rightarrow B, BC \rightarrow D \}$, we see that $\{AC\}^+ = ACD$, so D is extraneous in $AC \rightarrow D$
 - (If you eliminate the entire RHS of a functional dependency, it goes away)

Lossy Decompositions

10

- Some schema decompositions lose information

- Example:

employee(emp_id, emp_name, phone, title, salary, start_date)

- ▣ Decomposed into:

emp_ids(emp_id, emp_name)

emp_details(emp_name, phone, title, salary, start_date)

- Problem:

- ▣ *emp_name* doesn't uniquely identify employees
- ▣ This is a lossy decomposition

Lossless Decompositions

11

- Given:
 - ▣ Relation schema R , relation $r(R)$
 - ▣ Set of functional dependencies F
- Let R_1 and R_2 be a decomposition of R
 - ▣ $R_1 \cup R_2 = R$
- The decomposition is lossless if, for all legal instances of r :
$$\Pi_{R_1}(r) \bowtie \Pi_{R_2}(r) = r$$
- A simple definition...

Lossless Decompositions (2)

12

- Can define with functional dependencies:
 - ▣ R_1 and R_2 form a lossless decomposition of R if at least one of these dependencies is in F^+ :
$$R_1 \cap R_2 \rightarrow R_1$$
$$R_1 \cap R_2 \rightarrow R_2$$
- $R_1 \cap R_2$ forms a superkey of R_1 and/or R_2
 - ▣ Test for superkeys using attribute-set closure

Decomposition Examples (1)

13

- The *employee* example:

employee(*emp_id*, *emp_name*, *phone*, *title*, *salary*,
start_date)

- Decomposed into:

emp_ids(*emp_id*, *emp_name*)

emp_details(*emp_name*, *phone*, *title*, *salary*, *start_date*)

- *emp_name* is not a superkey of *emp_ids* or *emp_details*, so the decomposition is lossy

Decomposition Examples (2)

14

- The *bor_loan* example:

bor_loan(*cust_id*, *loan_id*, *amount*)

- Decomposed into:

borrower(*cust_id*, *loan_id*)

loan(*loan_id*, *amount*) (*loan_id* \rightarrow *loan_id*, *amount*)

- *loan_id* is a superkey of *loan*, so the decomposition is lossless

BCNF Decompositions

15

- If R is a schema not in BCNF:
 - ▣ There is at least one nontrivial functional dependency $\alpha \rightarrow \beta$ such that α is not a superkey for R
 - ▣ For simplicity, also require that $\alpha \cap \beta = \emptyset$
 - (if $\alpha \cap \beta \neq \emptyset$ then $(\alpha \cap \beta)$ is extraneous in β)
- Replace R with two schemas:
 - $R_1 = (\alpha \cup \beta)$
 - $R_2 = (R - \beta)$
 - (was $R - (\beta - \alpha)$, but $\beta - \alpha = \beta$, since $\alpha \cap \beta = \emptyset$)
- BCNF decomposition is lossless
 - ▣ $R_1 \cap R_2 = \alpha$
 - ▣ α is a superkey of R_1
 - ▣ α also appears in R_2

Dependency Preservation

16

- Some schema decompositions are not dependency-preserving
 - ▣ Functional dependencies that span multiple relation schemas are hard to enforce
 - ▣ e.g. BCNF may require decomposition of a schema for one dependency, and make it hard to enforce another dependency
- Can test for dependency preservation using functional dependency theory

Dependency Preservation (2)

17

- Given:
 - ▣ A set F of functional dependencies on a schema R
 - ▣ R_1, R_2, \dots, R_n are a decomposition of R
- The restriction of F to R_i is the set F_i of functional dependencies in F^+ that only has attributes in R_i
 - ▣ Each F_i contains functional dependencies that can be checked efficiently, using only R_i
- Find *all* functional dependencies that can be checked efficiently
 - ▣ $F' = F_1 \cup F_2 \cup \dots \cup F_n$
 - ▣ If $F'^+ = F^+$ then the decomposition is dependency-preserving

Third Normal Form Schemas

18

- Can generate a 3NF schema from a set of functional dependencies F
- Called the 3NF synthesis algorithm
 - ▣ Instead of decomposing an initial schema, generates schemas from a set of dependencies
- Given a set F of functional dependencies
 - ▣ Uses the canonical cover F_c
 - ▣ Ensures that resulting schemas are dependency-preserving

3NF Synthesis Algorithm

19

□ Inputs: set of functional dependencies F , on a schema R

let F_c be a canonical cover for F ;

$i := 0$;

for each functional dependency $\alpha \rightarrow \beta$ in F_c **do**

if none of the schemas $R_i, i = 1, 2, \dots, i$ contains $(\alpha \cup \beta)$ **then**

$i := i + 1$;

$R_i := (\alpha \cup \beta)$

end if

done

if no schema $R_i, i = 1, 2, \dots, i$ contains a candidate key for R **then**

$i := i + 1$;

$R_i :=$ any candidate key for R

end if

return (R_1, R_2, \dots, R_i)

BCNF vs. 3NF

20

- Boyce-Codd Normal Form:
 - ▣ Eliminates more redundant information than 3NF
 - ▣ Some functional dependencies become expensive to enforce
 - The conditions to enforce involve multiple relations
 - ▣ Overall, a very desirable normal form!
- Third Normal Form:
 - ▣ All [more] dependencies are [probably] easy to enforce...
 - ▣ Allows more redundant information, which must be kept synchronized by the database application!
 - ▣ Personal banker example:
 - works_in(emp_id, branch_name)*
 - cust_banker_branch(cust_id, branch_name, emp_id, type)*
 - Branch names must be kept synchronized between these relations!

BCNF and 3NF vs. SQL

21

- SQL constraints:
 - ▣ Only key constraints are fast and easy to enforce!
 - ▣ Only easy to enforce functional dependencies $\alpha \rightarrow \beta$ if α is a key on some table!
 - ▣ Other functional dependencies (even “easy” ones in 3NF) may require more expensive constraints, e.g. **CHECK**
- For SQL databases with materialized views:
 - ▣ Can decompose a schema into BCNF
 - ▣ For dependencies $\alpha \rightarrow \beta$ not preserved in decomposition, create materialized view joining all relations in dependency
 - ▣ Enforce **unique**(α) constraint on materialized view
- Impacts both space and performance, but it works...

Multivalued Attributes

22

- E-R schemas can have multivalued attributes
- 1NF requires only atomic attributes
 - ▣ Not a problem; translating to relational model leaves everything atomic

- Employee example:

employee(*emp_id*, *emp_name*)

emp_deps(*emp_id*, *dependent*)

emp_nums(*emp_id*, *phone_num*)

<i>employee</i>
<u><i>emp_id</i></u>
<i>emp_name</i>
{ <i>phone_num</i> }
{ <i>dependent</i> }

- What are the requirements on these schemas for what tuples must appear?

Multivalued Attributes (2)

23

□ Example data:

<i>emp_id</i>	<i>emp_name</i>
125623	Rick

employee

<i>emp_id</i>	<i>dependent</i>
125623	Jeff
125623	Alice

emp_deps

<i>emp_id</i>	<i>phone_num</i>
125623	555-8888
125623	555-2222

emp_nums

- Every distinct value of multivalued attribute requires a separate tuple, including associated value of *emp_id*
- A consequence of 1NF, in fact!
 - If attributes could be nonatomic, could just store list of values in the appropriate column!
 - 1NF requires extra tuples to represent multivalues

Independent Multivalued Attributes

24

- Question is trickier when a schema stores several *independent* multivalued attributes
- Proposed combined schema:
 `employee(emp_id, emp_name)`
 `emp_info(emp_id, dependent, phone_num)`
- What tuples must appear in `emp_info` ?
 - ▣ `emp_info` is a relation
 - ▣ If an employee has M dependents and N phone numbers, `emp_info` must contain $M \times N$ tuples
 - Exactly what we get if we natural-join `emp_deps` and `emp_nums`
 - ▣ Every combination of the employee's dependents and their phone numbers

Independent Multivalued Attributes

25

□ Example data:

<i>emp_id</i>	<i>emp_name</i>
125623	Rick

employee

<i>emp_id</i>	<i>dependent</i>	<i>phone_num</i>
125623	Jeff	555-8888
125623	Jeff	555-2222
125623	Alice	555-8888
125623	Alice	555-2222

emp_info

- Clearly has unnecessary redundancy
- Can't formulate functional dependencies to represent multivalued attributes
- Can't use BCNF or 3NF decompositions to eliminate redundancy in these cases

Multivalued Attributes Example

26

- Two employees: Rick and Bob
 - ▣ Both share a phone number at work
 - ▣ Both have two kids
 - ▣ Both have a kid named Alice
- Can't use functional dependencies to reason about this situation!
 - ▣ $emp_id \rightarrow phone_num$ doesn't hold since an employee can have several phone numbers
 - ▣ $phone_num \rightarrow emp_id$ doesn't hold either, since several employees can have the same phone number
 - ▣ Same with emp_id and $dependent...$

<i>emp_id</i>	<i>emp_name</i>
125623	Rick
127341	Bob

employee

<i>emp_id</i>	<i>phone_num</i>
125623	555-8888
125623	555-2222
127341	555-2222

emp_nums

<i>emp_id</i>	<i>dependent</i>
125623	Jeff
125623	Alice
127341	Alice
127341	Clara

emp_deps

Dependencies

27

- Functional dependencies rule out what tuples can appear in a relation
 - ▣ If $A \rightarrow B$ holds, then tuples cannot have same value for A but different values for B
 - ▣ Also called equality-generating dependencies
- Multivalued dependencies specify what tuples must be present
 - ▣ To represent a multivalued attribute's values properly, a certain set of tuples *must* be present
 - ▣ Also called tuple-generating dependencies

Multivalued Dependencies

28

- Given a relation schema R
 - Attribute-sets $\alpha \in R, \beta \in R$
 - $\alpha \twoheadrightarrow \beta$ is a multivalued dependency
 - “ α multidetermines β ”
- A multivalued dependency $\alpha \twoheadrightarrow \beta$ holds on R if, in any legal relation $r(R)$:

For all pairs of tuples t_1 and t_2 in r such that $t_1[\alpha] = t_2[\alpha]$,
There also exists tuples t_3 and t_4 in r such that:

 - $t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$
 - $t_1[\beta] = t_3[\beta]$ and $t_2[\beta] = t_4[\beta]$
 - $t_1[R - \beta] = t_4[R - \beta]$ and $t_2[R - \beta] = t_3[R - \beta]$

Multivalued Dependencies (2)

29

- Multivalued dependency $\alpha \twoheadrightarrow \beta$ holds on R if, in any legal relation $r(R)$:

For all pairs of tuples t_1 and t_2 in r such that $t_1[\alpha] = t_2[\alpha]$,

There also exists tuples t_3 and t_4 in r such that:

- $t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$
- $t_1[\beta] = t_3[\beta]$ and $t_2[\beta] = t_4[\beta]$
- $t_1[R - \beta] = t_4[R - \beta]$ and $t_2[R - \beta] = t_3[R - \beta]$

- Pictorially:

	α	β	$R - (\alpha \cup \beta)$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

Multivalued Dependencies (3)

30

- Multivalued dependency:

	α	β	$R - (\alpha \cup \beta)$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

- If $\alpha \twoheadrightarrow \beta$ then $R - (\alpha \cup \beta)$ is independent of this fact
 - Every distinct value of β must be associated once with every distinct value of $R - (\alpha \cup \beta)$
- Let $\gamma = R - (\alpha \cup \beta)$
 - If $\alpha \twoheadrightarrow \beta$ then also $\alpha \twoheadrightarrow \gamma$
 - $\alpha \twoheadrightarrow \beta$ implies $\alpha \twoheadrightarrow \gamma$
 - Sometimes written $\alpha \twoheadrightarrow \beta \mid \gamma$

Trivial Multivalued Dependencies

31

- $\alpha \twoheadrightarrow \beta$ is a trivial multivalued dependency on R if all relations $r(R)$ satisfy the dependency
- Specifically, $\alpha \twoheadrightarrow \beta$ is trivial if $\beta \subseteq \alpha$, or if $\alpha \cup \beta = R$
- Employee examples:
 - For schema $\text{emp_deps}(\text{emp_id}, \text{dependent})$, $\text{emp_id} \twoheadrightarrow \text{dependent}$ is trivial
 - For $\text{emp_info}(\text{emp_id}, \text{dependent}, \text{phone_num})$, $\text{emp_id} \twoheadrightarrow \text{dependent}$ is not trivial

Inference Rules

32

- Can reason about multivalued dependencies, just like functional dependencies
 - ▣ There is a set of complete, sound inference rules for MVDs
- Example inference rules:
 - ▣ Complementation rule:
 - If $\alpha \twoheadrightarrow \beta$ holds on R , then $\alpha \twoheadrightarrow R - (\alpha \cup \beta)$ holds
 - ▣ Multivalued augmentation rule:
 - If $\alpha \twoheadrightarrow \beta$ holds, and $\gamma \subseteq R$, and $\delta \subseteq \gamma$, then $\gamma\alpha \twoheadrightarrow \delta\beta$ holds
 - ▣ Multivalued transitivity rule:
 - If $\alpha \twoheadrightarrow \beta$ and $\beta \twoheadrightarrow \gamma$ holds, then $\alpha \twoheadrightarrow \gamma - \beta$ holds
 - ▣ Coalescence rule:
 - If $\alpha \twoheadrightarrow \beta$ holds, and $\gamma \subseteq \beta$, and there is a δ such that $\delta \subseteq R$, and $\delta \cap \beta = \emptyset$, and $\delta \rightarrow \gamma$, then $\alpha \rightarrow \gamma$ holds

Functional Dependencies

33

- Functional dependencies are also multivalued dependencies
- Replication rule:
 - ▣ If $\alpha \rightarrow \beta$, then $\alpha \twoheadrightarrow \beta$ too
 - ▣ Note there is an additional constraint from $\alpha \rightarrow \beta$: each value of α has *at most* one associated value for β
- Usually, functional dependencies are not stated as multivalued dependencies
 - ▣ The extra caveat is *important*, but not obvious in notation
 - ▣ Also, functional dependencies are easier to reason about!

Closures and Restrictions

34

- For a set D of functional and multivalued dependencies, can compute closure D^+
 - ▣ Use inference rules for both functional and multivalued dependencies to compute closure
- Sometimes need the restriction of D^+ to a relation schema R , too
- The restriction of D to a schema R_i includes:
 - ▣ All functional dependencies in D^+ that include only attributes in R_i
 - ▣ All multivalued dependencies of the form $\alpha \twoheadrightarrow \beta \cap R_i$, where $\alpha \subseteq R_i$, and $\alpha \twoheadrightarrow \beta$ is in D^+

Fourth Normal Form

35

- Given:
 - ▣ Relation schema R
 - ▣ Set of functional and multivalued dependencies D
- R is in 4NF with respect to D if:
 - ▣ For all multivalued dependencies $\alpha \twoheadrightarrow \beta$ in D^+ , where $\alpha \subseteq R$ and $\beta \subseteq R$, at least one of the following holds:
 - $\alpha \twoheadrightarrow \beta$ is a trivial multivalued dependency
 - α is a superkey for R
 - ▣ Note: If $\alpha \rightarrow \beta$ then $\alpha \twoheadrightarrow \beta$
- A database design is in 4NF if all schemas in the design are in 4NF

4NF and BCNF

36

- Main difference between 4NF and BCNF is use of multivalued dependencies instead of functional dependencies
- Every schema in 4NF is also in BCNF
 - ▣ If a schema is not in BCNF then there is a nontrivial functional dependency $\alpha \rightarrow \beta$ such that α is not a superkey for R
 - ▣ If $\alpha \rightarrow \beta$ then $\alpha \twoheadrightarrow \beta$

4NF Decompositions

37

- Decomposition rule very similar to BCNF
- If schema R is not in 4NF with respect to a set of multivalued dependencies D :
 - ▣ There is some nontrivial dependency $\alpha \twoheadrightarrow \beta$ in D^+ where $\alpha \subseteq R$ and $\beta \subseteq R$, and α is not a superkey of R
 - Also constrain that $\alpha \cap \beta = \emptyset$
 - ▣ Replace R with two new schemas:
 - $R_1 = (\alpha \cup \beta)$
 - $R_2 = (R - \beta)$

Employee Information Example

38

- Combined schema:

employee(emp_id, emp_name)

emp_info(emp_id, dependent, phone_num)

- Also have these dependencies:

- $emp_id \rightarrow emp_name$

- $emp_id \twoheadrightarrow dependent$

- $emp_id \twoheadrightarrow phone_num$

- *emp_info* is not in 4NF

- Following the rules for 4NF decomposition produces:

(emp_id, dependent)

(emp_id, phone_num)

- Note: Each relation's candidate key is the entire relation. The multivalued dependencies are trivial.

Lossless Decompositions

39

- Can also define lossless decomposition with multivalued dependencies
 - ▣ R_1 and R_2 form a lossless decomposition of R if at least one of these dependencies is in D^+ :

$$R_1 \cap R_2 \twoheadrightarrow R_1$$

$$R_1 \cap R_2 \twoheadrightarrow R_2$$

Beyond Fourth Normal Form?

40

- Additional normal forms with various constraints
- Example: join dependencies
- Given R , and a decomposition R_1 and R_2 where $R_1 \cup R_2 = R$:
 - ▣ The decomposition is lossless if, for all legal instances of $r(R)$,
 $\Pi_{R_1}(r) \bowtie \Pi_{R_2}(r) = r$
- Can state this as a join dependency: $*(R_1, R_2)$
 - ▣ This is actually identical to a multivalued dependency!
 - ▣ $*(R_1, R_2)$ is equivalent to $R_1 \cap R_2 \twoheadrightarrow R_1 \mid R_2$

Join Dependencies and 5NF

41

- Join dependencies (JD) are a generalization of multivalued dependencies (MVD)
 - ▣ Can specify JDs involving N relation schemas, $N \geq 2$
 - ▣ JDs are equivalent to MVDs when $N = 2$
 - ▣ Can easily construct JDs where $N > 2$, with no equivalent set of MVDs
- Project-Join Normal Form (a.k.a. PJNF or 5NF):
 - ▣ A relation schema R is in PJNF with respect to a set of join dependencies D if, for all JDs in D^+ of the form $*(R_1, R_2, \dots, R_n)$ where $R_1 \cup R_2 \cup \dots \cup R_n = R$, at least one of the following holds:
 - $*(R_1, R_2, \dots, R_n)$ is a trivial join dependency
 - Every R_i is a superkey for R

Join Dependencies and 5NF (2)

42

- If a schema is in Project-Join Normal Form then it is also in 4NF (and thus, in BCNF)
 - ▣ Every multivalued dependency is also a join dependency
 - ▣ (Every functional dependency is also a multivalued dependency)
- One small problem:
 - ▣ There isn't a complete, sound set of inference rules for join dependencies!
 - ▣ Can't reason about our set of join dependencies $D...$
 - ▣ This limits PJNF's real-world usefulness

Domain-Key Normal Form

43

- Domain-key normal form (DKNF) is an even more general normal form, based on:
 - ▣ **Domain constraints:** what values may be assigned to attribute A
 - Usually inexpensive to test, even with **CHECK** constraints
 - ▣ **Key constraints:** all attribute-sets K that are a superkey for a schema R (i.e. $K \rightarrow R$)
 - Almost always inexpensive to test
 - ▣ **General constraints:** other predicates on valid relations in a schema
 - Could be very expensive to test!
- A schema R is in DKNF if the domain constraints and key constraints logically imply the general constraints
 - ▣ An “ideal” normal form difficult to achieve in practice...