

OBJECTS AND DATABASES

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 21

Relational Model and 1NF

2

- Relational model specifies that all attribute domains must be atomic
 - ▣ A database schema is in 1NF if all attribute domains are atomic
- Not always preferred approach in real world use
- In relational model:
 - employee(emp_id, emp_name)*
 - emp_phone(emp_id, phone_num)*
 - emp_deps(emp_id, dependent)*
 - ▣ Need two joins just to get all data for an employee!

<i>employee</i>
<i><u>emp_id</u></i>
<i>emp_name</i>
<i>{ phone_num }</i>
<i>{ dependent }</i>

Composite Types

3

- Also, frequently have composite types that are reused
- Example:
 - Add home/work addresses to design
- In relational model, composite types are decomposed
*employee(emp_id, emp_name,
work_street, work_city, ...)*
- ...but programming languages typically provide structures or classes for composite types!

<i>employee</i>
<i><u>emp_id</u></i>
<i>emp_name</i>
<i>work_address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zipcode</i>
<i>home_address</i>
<i>street</i>
<i>city</i>
<i>state</i>
<i>zipcode</i>
{ <i>phone_num</i> }
{ <i>dependent</i> }

Database Applications

4

- Programming languages have support for non-atomic types
 - Address class or structure:
 - street, city, state, zipcode
 - Arrays of phone numbers and dependents
- Application has to translate between relational model version and programming language representation
 - Annoying to deal with, at the least...
 - At worst, can have substantial application quality and performance impacts!

SQL User-Defined Types

5

- SQL:1999 includes User-Defined Types
 - ▣ Allows users to define non-atomic types where appropriate
 - ▣ (Make sure it's actually appropriate!)
 - ▣ Frequently abbreviated as UDT
- Multivalued types – arrays, sets, lists, etc.
 - ▣ Elements are all the same type
- Structured types – composite attributes
 - ▣ Elements may be different types

Non-Atomic Types for Employees

6

- Declare new UDT for addresses:

```
CREATE TYPE Address AS (  
    street VARCHAR(60),  
    city    VARCHAR(40),  
    state   CHAR(2),  
    zipcode CHAR(9)  
);
```

 - Only specify types, not constraints!
 - Defines a new structured type within the database schema
- For arrays, just add **ARRAY [n]** to column type
 - **n** is optional
 - Array elements have indexes 1 to **n**

Using Non-Atomic Types

7

- Employee table:

```
CREATE TABLE employee (  
    emp_id          INTEGER          PRIMARY KEY,  
    emp_name        VARCHAR(100)    NOT NULL,  
    work_address    Address          NOT NULL,  
    home_address    Address          NOT NULL,  
    phone_nums      CHAR(12)         ARRAY[],  
    dependents      VARCHAR(100)     ARRAY[]  
);
```

- Now all details of an employee are contained within a single table
 - E-R model maps directly into this design
- Retrieving all details of an employee will be fast

Structured Types in DML

8

- Accessing elements of a structured type:

```
SELECT emp_id, emp_name FROM employee  
WHERE work_address.city = home_address.city;
```
- Specifying all values of a structured type:

```
UPDATE employee SET work_address =  
    ('123 Main St.', 'Springfield', 'OH', '45505')  
WHERE emp_id = 5352;
```
- Specifying individual values of a structured type:

```
UPDATE employee SET work_address.city = 'Akron',  
    work_address.zipcode = '44310'  
WHERE emp_id = 5352;
```


Array Types in DML

9

- Specifying all values of an array type:

```
UPDATE employee SET phone_nums =  
    ARRAY['800-555-1234', '800-555-5678']  
WHERE emp_id = 5352;
```

- Specifying individual values of an array type:

```
UPDATE employee  
    SET phone_nums[1] = '800-555-2345'  
WHERE emp_id = 5352;
```

- Order of elements in array is preserved!
 - ▣ Useful when order of values is meaningful
 - ▣ e.g. author-list in a database of research papers

Array Types in DML (2)

10

- Array columns are like nested relations
 - ▣ A nested relation is stored within a single column
- SQL:1999 provides nesting and unnesting operations for arrays
- To unnest an array value:

```
SELECT emp_id, emp_name, p.phone
FROM employee AS e,
      UNNEST(e.phone_nums) AS p(phone)
WHERE emp_id = 5352;
```

<i>emp_id</i>	<i>emp_name</i>	<i>phone</i>
5352	Bob Smith	800-555-2345
5352	Bob Smith	800-555-5678

Array Types in DML (3)

11

- Can also retrieve element ordering details

```
SELECT emp_id, emp_name, p.phone, p.p_index
FROM employee AS e,
     UNNEST(e.phone_nums) WITH ORDINALITY
     AS p(phone, p_index);
```

<i>emp_id</i>	<i>emp_name</i>	<i>phone</i>	<i>p_index</i>
5352	Bob Smith	800-555-2345	1
5352	Bob Smith	800-555-5678	2

- Can use **COLLECT** to combine values into an array

```
SELECT emp_id, COLLECT(phone_num) AS phone_nums
FROM raw_employee_data GROUP BY emp_id;
```

 - ▣ Very similar to grouping and aggregation operation!
- Can also pass subquery to **ARRAY()** fn. to populate an array

SQL:1999 User Defined Types

12

- SQL:1999 user-defined types help with composite and multivalued attributes...
 - ▣ Can create schemas that don't incur join overheads for multivalued attributes
 - ▣ Can represent composite attributes more naturally within the SQL schema
- Still not quite the same as what programming languages can provide

Objects and Relations

13

- Many programming languages are object-oriented
 - ▣ Objects: encapsulation of state and behavior
 - ▣ Classes: specifications of objects' state and behavior
 - ▣ Also other features, such as class inheritance
 - A class can derive from a parent class
 - Child class has specialized capabilities and additional state
 - ▣ C++, Java, C#, Python, PHP, etc. All widely used.
- Typical approach for storing objects in a relational database:
 - ▣ Classes usually map to tables
 - ▣ Objects map to individual rows in a table

Objects and Relations (2)

14

- Relational databases aren't designed to store objects!
- “Object-relational impedance mismatch”
 - ▣ A number of serious issues arise when storing objects into a relational database
- Relational databases cannot enforce the same constraints that OOP languages can enforce!
 - ▣ Objects encapsulate and manage state very carefully
 - ▣ In a relational database, all values can be manipulated very easily
 - ▣ Storing object-data in a relational database increases potential for data corruption

Objects and Relations (3)

15

- Objects can reference each other
 - More akin to the network data model, which preceded the relational model
 - Objects are accessed by following specific object-references
 - Tuples are retrieved en masse via a query language
- Representing object identities and object references in a relational database can be tricky
 - In OOP languages, object-references are usually opaque to the user, and not manipulated directly
 - In relational model, identifying values are intentionally very visible and meaningful
 - Part of Codd's original intent with relational model

Objects and Relations (4)

16

- OOP languages also provide features not present in relational model
 - ▣ Ability to specify methods to be called on objects
 - ▣ Class inheritance and class hierarchies
 - ▣ Require careful modeling in a relational database, if it can be implemented at all!
 - Frequently have multiple choices for modeling, with different performance and space implications
- A number of other issues as well...
 - ▣ Some are more esoteric than others
- Can definitely live with most of these issues, but be aware of the mismatch in capabilities!

Addressing The Mismatch

17

- Two main approaches to object-relational mismatch
- Object-Oriented Databases (ODBMS/OODBMS)
 - Further extend SQL's type system to support basic object-oriented constructs
 - Database supports object-oriented abstractions directly and internally
- Persistent Programming Languages
 - Hide (R)DBMS storage operations from programmers
 - Automate the translation between programming-language objects and database storage

Object-Oriented Database Systems

18

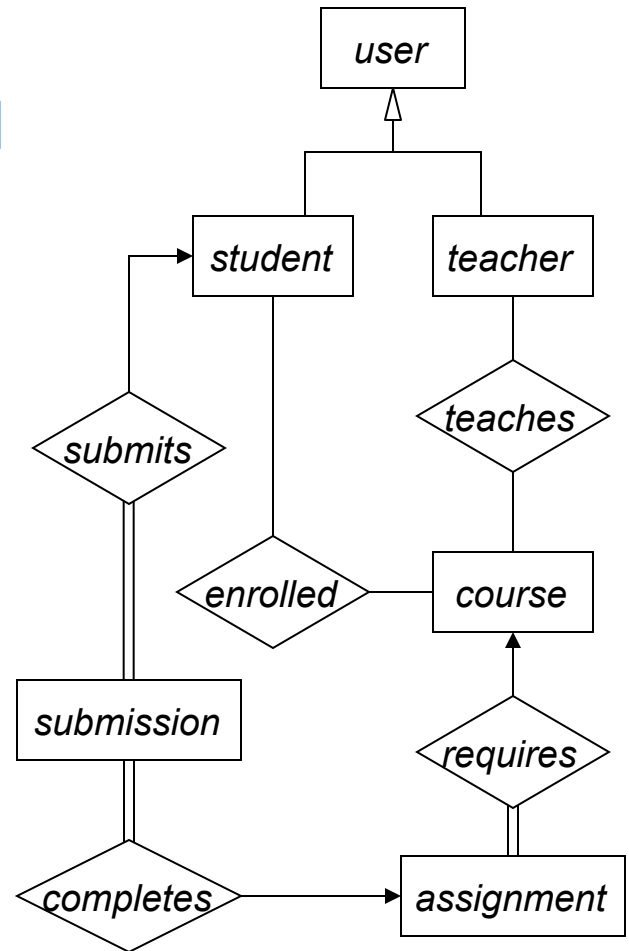
- ODBMSes provide direct support for classes and objects
 - ▣ Define constructors and methods for classes
 - ▣ Define type hierarchies for classes
 - ▣ Provide object-reference support
- Inclusion of objects requires significant changes to the query language
 - ▣ Objects can refer to collections of objects
 - ▣ Must support path-based queries

ODBMS Example

19

- Course management system database schema
- Entities are objects in the ODBMS
 - Relationships specify object-references
- Retrieve names of students enrolled in CS121:

```
SELECT s.name FROM student s
WHERE s.enrolled.course.name = 'cs121';
```



ODBMS Operations

20

- ODBMSes must provide capabilities for:
 - Object-definition, similar to SQL DDL
 - Queries on objects, similar to SQL DML
- Object Data Management Group
 - Consortium founded in 1991, to create ODBMS specifications
- Standardization effort has had limited success
 - Several DB vendors offer ODBMS capabilities, but syntax and feature-sets vary pretty widely.
 - (not unlike SQL standard...)

ODBMS Operations (2)

21

- Object Description Language (ODL)
 - ▣ For specifying object-database schemas
 - ▣ Specify classes, class-members, and class inheritance hierarchies
- Object Query Language (OQL)
 - ▣ A SQL-like query language for querying object-databases
 - ▣ Most significant change is ability to specify “path expressions” that follow relationships between objects

Object-References

22

- All objects in an ODBMS have a unique identifier of some kind
 - Object Identifier (OID)
- Objects reference other objects using their OIDs
 - Akin to a pointer or reference to the object
- ODBMSes generally load referenced objects from disk, as needed.
 - Lazy loading, not eager loading
 - Objects tend to reference many other objects...
 - Object-loading must be done carefully, to avoid unnecessary resource usage!

ODBMS Summary

23

- Hasn't been widespread adoption of ODBMSes
 - ▣ Cost of switching is very high
 - A company's data is extremely valuable
 - Relational model is satisfactory for most needs, so why change?
- Many commercial databases provide a hybrid data model now
 - ▣ Object-Relational Database Management System (ORDBMS)
 - Blends object capabilities and relational database capabilities
 - ▣ Typically provide capabilities for simple type hierarchies, and simple class-method declarations

Persistent Programming Languages

24

- Most popular approach to object-relational impedance mismatch:
 - Create or enhance OOP languages to provide persistent objects directly in the language itself
- Normally, when a program terminates, all objects it created go away
 - These objects are transient
- Persistent objects are stored before termination
 - (in a database of some kind...)
 - Next time the program runs, persistent objects can be retrieved and used

Persistent Programming Languages (2)

25

- Persistent programming languages usually store objects in relational databases
 - ▣ PostgreSQL, MySQL, SQLite, Oracle, etc.
 - ▣ Also called “object-relational mapping layers” or simply “object-relational mappers” (abbrev. ORM)
- Type specification is entirely in the OO programming language itself
 - ▣ Able to leverage most types and OO capabilities of the programming language itself
 - ▣ Usually very few differences in capabilities between transient and persistent objects

Database Operations

26

- Database operations are usually entirely obscured from the programmer
 - ▣ Persistent object storage and retrieval is handled entirely by the framework itself
- Many ORM layers also provide automatic data-definition capabilities
 - ▣ Given a set of persistent objects, ORM layer can generate a SQL schema for those objects
 - ▣ Persistent objects are typically annotated to indicate “primary key” values, etc.

Automatic DDL Generation

27

- Persistence frameworks are becoming quite sophisticated with auto-DDL generation...
- Two main issues to consider!
- Database schema migration:
 - It's easy to change classes, or add new ones
 - Absolutely essential to have a migration path for existing data!
- Database performance:
 - ORM layers don't usually generate a schema tuned for high-performance and scalability

ORM Layers and Schema Migration

28

- Most ORM layers don't yet provide schema/data migration capabilities directly...
- Typically, external libraries/tools are available, to add schema-migration support to ORM layers
- General approach:
 - ▣ Take a snapshot of every stable version of data model
 - ▣ When the schema changes in *simple* ways, tool can generate the needed SQL DDL to migrate the schema
 - ▣ Tools also support manual data-migration steps for more involved changes
- Always back up data before using these tools! 😊

ORM Layers and Performance

29

- Most ORM layers also provide ability to run custom SQL on database before/after schema-generation
 - ▣ Add stored procedures and user-defined functions
 - ▣ Add indexes
 - ▣ Populate database with initial data
- To facilitate this, ORM layers frequently document exactly how table/column names are generated
- DDL that isn't specifically managed by the ORM can make data models significantly more fragile!
 - ▣ May need to change names/structure in multiple places

Manual DDL Creation

30

- Because of these issues, ORM layers also frequently support mapping objects to an existing schema
- You design the schema with specific needs in mind
 - Specify table indexes, partition large tables, etc.
- When a schema needs to change, it's easier to design a migration plan for your data
 - You have the “old” schema definition...
 - You provide the new schema definition yourself...
 - You can design the migration process for upgrading the database and preserving your data.

Persistence Framework Limitations

31

- Many persistence frameworks impose limitations on the kinds of schemas they support
 - Make sure to understand these limitations before designing schemas for these frameworks!
- Multiple-column primary keys
 - Supported by more advanced ORM layers...
 - ...but they may not support multi-column foreign keys!
- Ternary relationships
 - Many ORM layers only support binary relationships
 - Need to model ternary relationships as a combination of binary relationships

Persistence Framework Challenges


32

- Many persistence frameworks also have limited support for database constraints
- Virtually all can handle referential integrity constraints
- Not all ORM layers can handle objects with multiple candidate keys
 - ▣ ...but these days, most of them can.
- Be careful about general **CHECK** constraints!
- ORM layer must identify the cause of database errors generated by violating these constraints
 - ▣ Hard to do for a wide range of database vendors

Loading Persistent Objects

33


- Like ODBMSes, persistent object frameworks must carefully manage object retrieval
- Example:
 - Retrieve all students in CS121
- Step 1: retrieve the **Course** object with name of “cs121”
 - The **Course** object will have a set of **Student**-references, a set of **Assignment**-refs, etc.



Loading Persistent Objects (2)

34

- Should the **Course** object eagerly or lazily load **Student** objects from the database?
 - ▣ We said we want all students...
 - ▣ Makes most sense to get all of them in one query.
- **Student** objects will have a collection of **Submission** objects, each of which has an **Assignment** object, ...
 - ▣ ORM layer must get exactly what is needed, and no more!



Persistent-Object References

35

- OO programming language already has a way of referencing objects
 - ▣ e.g. pointers in C++, or references in Java/Python/...
- ORM layers must map between in-memory reference type and database reference type
 - ▣ A persistent object may not be loaded into memory yet, but other in-memory persistent objects refer to it
- Two kinds of persistent-object references:
 - ▣ A database-reference for when object isn't loaded yet
 - ▣ An in-memory reference for when the object is already in memory

Persistent-Object References (2)

36

- When a DB-reference is followed:
 - ORM layer loads object into memory from database
 - Then, ORM layer switches out the DB-reference for an in-memory reference
- In compiled languages, often implemented with pointer-swizzling
 - ORM layer uses special pointer-values for database-references to objects
 - When pointer is accessed, the ORM layer is notified (e.g. via a page-fault signal)
 - ORM layer loads the object, then directly changes the pointer value to point to the loaded object instead

Persistent-Object References (3)

37

- In interpreted (or VM-based) languages, often implemented with hollow objects
 - ▣ Before a persistent object is loaded, the reference actually points to a proxy
 - ▣ When the proxy is accessed, ORM layer retrieves the object from the DB
 - ▣ Proxy is replaced with the loaded object
- In-memory objects must also track state changes!
 - ▣ Writes to the object must flag the object as “dirty”
 - ▣ ORM layer ensures that dirty objects are saved to DB

Persistent Objects – Summary

38

- Much more popular solution to object-relational impedance mismatch
 - ▣ But, an admittedly incomplete solution to the mismatch
- Obscures much of the pain of moving objects to and from the database
- Also frequently provides ability to manage schema design directly (but watch out for limitations!)
- Some Java persistence frameworks:
 - ▣ Java Persistence API, Hibernate
- Some Python persistence frameworks:
 - ▣ Django models (+ South for migration), SQLAlchemy