

DATA STREAMS AND DATABASES

CS121: Introduction to Relational Database Systems
Fall 2014 – Lecture 22

Static and Dynamic Data Sets

2

- So far, have discussed relatively static databases
 - ▣ Data may change slowly over time...
 - ▣ Queries and updates operate against a static data-set
 - ▣ Especially true in context of transactions: within each txn, database appears to be unchanged by other txns
- Increasingly common to have data streams generated by various sources
 - ▣ An infinite, time-ordered sequence of tuples
 - ▣ Tuples have a particular schema, as before
 - ▣ One attribute in the tuple schema is a timestamp τ

Data Streams: Examples

3

- Many different examples of data streams
- Stock market data!
- Example: Volume-Weighted Average Price (VWAP)
 - ▣ Value computed over a time-window of stock trades
 - ▣ Window is fixed size, contains n stock trades
 - ▣ Trade i has price P_i , with Q_i shares changing hands
 - ▣ $P_{\text{VWAP}} = \sum P_i Q_i \div \sum Q_i$

Data Streams: Examples (2)

4

- Traffic-flow sensor networks generate data streams
 - ▣ Various estimates of traffic-flow characteristics, e.g. time/space mean speed, density, flow
- Seismographic networks generate many data streams!
 - ▣ Individual seismometers generate multiple data streams along different axes (vertical, N/S, E/W)
 - ▣ Other software consumes seismic data streams, identifies earthquakes, produces other data streams
- Computer network security monitoring systems
- Sensor networks in plants, factories, etc.

Data Streams in Relational Databases?

5

- How to implement Volume-Weighted Average Price (VWAP) computation with a relational database?
- Need to store the stream of stock prices into a table
 - ▣ Each stock price must have a timestamp
 - ▣ Example schema:
 - *stock_sales(sale_time, ticker_symbol, num_shares, price_per_share)*
- As new stock sales occur, store them in the table
- Need to eventually remove records from this table when we don't care anymore

Data Streams in Relational Databases?

6

- Periodically compute the VWAP against this table
 - ▣

```
SELECT ticker_symbol,  
       SUM(price_per_share * num_shares) /  
       SUM(num_shares)  
FROM stock_sales  
WHERE sale_time >= NOW() - INTERVAL 5 MINUTE  
GROUP BY ticker_symbol;
```
 - ▣ An index on (*sale_time*) will help select the rows in the window of interest
- Issues?
 - ▣ Performance: we throw away a significant amount of work every time the query is recomputed!

Data Streams in Relational Databases?

7

- Our query:
 - ▣

```
SELECT ticker_symbol,  
       SUM(price_per_share * num_shares) /  
       SUM(num_shares)  
FROM stock_sales  
WHERE sale_time >= NOW() - INTERVAL 5 MINUTE  
GROUP BY ticker_symbol;
```
 - ▣ *Can we do this incrementally instead?*
- Can easily apply same concepts as with materialized views to save and update intermediate state
 - ▣ As rows enter and leave the window of interest, we can update our rolling averages for each stock
 - ▣ Should make it very fast to generate the desired results!

Data Streams and Queries

8

- What if we want to update the output of our query every time the input data stream changes?
 - ▣ e.g. every time more stock values arrive, we update the corresponding VWAP immediately
 - ▣ (e.g. could implement this with triggers on *stock_sales*)
- What if we want to join the data stream against one or more relations?
- What if we want to generate a data stream based on the changes made to a relation?

Data Stream Management Systems

9

- Over the last few decades, many efforts to build Data Stream Management Systems (DSMS)
 - ▣ Like DataBase Management Systems (DBMS), but able to handle data streams as well as static tables
- Also called Complex Event Processing (CEP) systems
- Several approaches to building these systems...
- A popular approach:
 - ▣ Extend the relational model to add stream processing capabilities
 - ▣ (i.e. reuse the existing work of relational databases!)

Data Stream Management Systems (2)

10

- Many major research projects on relation-oriented stream databases:
 - ▣ Aurora/Borealis (Brown, Brandeis, MIT)
 - ▣ STREAM (Stanford)
 - ▣ TelegraphCQ (Berkeley)
- Commercial stream databases:
 - ▣ StreamBase (commercial version of Aurora/Borealis)
 - ▣ Truviso (commercial version of TelegraphCQ; acquired May 2012 by Cisco)
 - ▣ TIBCO Business Events, Oracle BAM

Stream Database Implementations

11

- Some stream databases were built by extending existing relational databases
 - ▣ e.g. TelegraphCQ and Truviso are extensions of PostgreSQL that incorporate data streams
 - ▣ Existing SQL syntax is extended to handle stream declarations, windowed queries on streams, etc.
- Many others are built from scratch
 - ▣ Custom stream-processing engines that are database-agnostic, or that don't use a relational database
 - ▣ Usually have a specific focus, e.g. high-volume data streams, low-latency results, specific kinds of queries, etc.
- Today: focus on relation-oriented stream databases

Data Stream Management Systems (3)

12

- In a relational database:
 - ▣ Data is static! (As long as no DML is issued...)
 - ▣ Issue queries against DB whenever we need results.
- In a stream database:
 - ▣ Stream data changes continually! Thus, results of queries against a stream also change continually.
 - ▣ Such queries are called continuous queries.
 - ▣ Register continuous queries with the database server.
 - ▣ As stream data changes, DB can incrementally update and output query results efficiently.

Stream Data Model

13

- A data stream S is an infinite, time-ordered multiset of tuples, one attribute being a timestamp τ
 - ▣ τ denotes the logical arrival time of the tuple into the system
- A relation R is an unordered multiset of tuples that varies over time
 - ▣ $R(\tau)$ is the version of the relation at a point in time τ
- A continuous query Q is constructed from a tree of operators against streams S_i and relations R_i

Continuous Query Operators

14

- Relation-to-relation operators take one or two relations and produce a relation
 - ▣ Exactly like standard relational algebra, except that relations in stream databases have a notion of time
 - ▣ Uses most recent versions of involved relations $R_i(\tau)$
- Relation-to-stream operators take a relation and produce a data stream
 - ▣ Typically defined to produce a stream containing changes made to a relation R_i
 - ▣ e.g. ISTREAM(R) produces tuples inserted into R at time τ
 - A stream of tuples $\langle s, \tau \rangle$ where $s \in R(\tau) - R(\tau - 1)$
 - ▣ e.g. RSTREAM(R) produces a stream of all tuples in R at τ

Continuous Query Operators (2)

15

- Stream-to-relation operators take a data stream S and produce a relation R
 - ▣ Most continuous queries only care about the *recent* tuples in a data stream...
 - ▣ Define a sliding window on a stream that ends at time τ
- Tuple-based sliding windows contain the N most recent tuples from the data stream S
- Time-based sliding windows contain tuples from S that fall in a range of timestamps
 - ▣ $R(\tau)$ contains all tuples from S with timestamp in $[\tau - \omega, \tau]$
 - ▣ Also provide support for “now” tuples, where $\omega = 0$

Continuous Queries

16


- When a continuous query is added to the database, a plan is constructed from these operators
- Example: compute volume-weighted average stock prices from stream of stock sales
 - ▣ Convert stock-sale data stream into a relation using a time-based sliding-window operator over last 5 minutes of data
 - ▣ Compute a relation containing aggregates using standard grouping/aggregation operations
 - ▣ Convert entire relation into another stream containing results
- Problem: as stated, this is still expensive
 - ▣ Every time, results are computed entirely from scratch!

Continuous Queries (2)

17

- A continuous query operator can maintain a synopsis of its most recent results
 - A relation containing the rows used to generate results for most recent time τ
 - When rows enter or exit the sliding window, synopsis can be incrementally updated very efficiently


- Example:



Continuous Queries (3)

18

- Can construct very complex queries, even with these [relatively] simple operators
- Example: windowed join of two data streams
 - ▣ Shared synopses across different nodes



Continuous Query Languages

19

- Some continuous query languages (CQL) are extensions of SQL
 - ▣ Particularly in implementations built on relational DBs
- Include ability to create and remove streams
 - ▣ `CREATE STREAM stock_sales (
 sale_time TIMESTAMP,
 ticker_symbol VARCHAR(10),
 num_shares INTEGER,
 price_per_share NUMERIC(7, 2)
);`
 - ▣ `DROP STREAM stock_sales;`
- (Actual stream data usually arrives over network and must be in a specific format for the database to use it)

Continuous Query Languages (2)

20

- Need a way to indicate what column is timestamp τ
- Example: Truviso syntax:
 - ▣ `CREATE STREAM stock_sales (
 sale_time TIMESTAMP CQTIME
 USER GENERATED,
 ticker_symbol VARCHAR(10),
 num_shares INTEGER,
 price_per_share NUMERIC(7, 2)
);`
 - ▣ Can also have system-generated timestamps
- Other stream databases use similar mechanisms, e.g. TelegraphCQ has **TIMESTAMPCOLUMN** modifier

Continuous Query Languages (3)

21

- Can also specify if a stream is archived or not
 - ▣ An archived stream can be queried for historical data; an unarchived stream cannot.
 - ▣ Most CQLs have **TYPE [ARCHIVED | UNARCHIVED]**
- To preserve historical data:
 - ▣ **CREATE STREAM stock_sales (
 sale_time TIMESTAMP CQTIME
 USER GENERATED,
 ticker_symbol VARCHAR(10) ,
 num_shares INTEGER,
 price_per_share NUMERIC(7, 2)
) TYPE ARCHIVED;**

Continuous Queries

22

- Can issue queries against both relations and streams
- If a stream is involved, need to specify the window!
- Stanford STREAM CQL annotates streams within the query:
 - ▣

```
SELECT ticker_symbol,  
       SUM(num_shares * price_per_share) /  
       SUM(num_shares)  
FROM stock_sales [RANGE 5 MINUTES]  
GROUP BY ticker_symbol;
```
- Truviso has a similar annotation:
 - ▣

```
SELECT ticker_symbol,  
       SUM(num_shares * price_per_share) /  
       SUM(num_shares)  
FROM stock_sales < VISIBLE '5 MINUTES' >  
GROUP BY ticker_symbol;
```

Continuous Queries (2)

23

- TelegraphCQ has a separate window specification:
 - ▣

```
SELECT ticker_symbol,  
        SUM(num_shares * price_per_share) /  
        SUM(num_shares)  
FROM stock_sales  
GROUP BY ticker_symbol  
WINDOW stock_sales ['5 MINUTES'];
```

Derived Streams

24

- Continuous queries run once and produce their results, just like standard queries
- Can create derived streams from continuous queries
 - ▣ Query is persistent, and produces a data stream of results
 - ▣ Results in output data stream have timestamps, as expected
- Truviso syntax:
 - ▣

```
CREATE STREAM vol_weighted_avg_price AS
    (SELECT ticker_symbol,
             SUM(num_shares * price_per_share) /
             SUM(num_shares)
     FROM stock_sales < VISIBLE '5 MINUTES'
                      ADVANCE '5 SECONDS' >
     GROUP BY ticker_symbol);
```
 - ▣ **ADVANCE** keyword specifies how often to generate results

Stream-Processing Challenges

25

- Many challenges in implementing stream databases
- Frequently, data streams are bursty and can have very high volumes
 - ▣ Peak message rate \gg average message rate
 - ▣ Latency between input data and results skyrockets
- Most common approach: load-shedding
 - ▣ Simply drop some of the input tuples out of the stream!
 - ▣ Stream DBs use statistics gathered about stream data to drop/summarize tuples to maximize accuracy of results
 - ▣ May involve dropping tuples from multiple sources in a query to ensure accuracy is maximized

Stream-Processing Challenges (2)

26

- Other databases use windowed/aggregate operators to achieve a similar result:
 - ▣ Use windowed aggregate operations to generate multiple granularities of a given data stream
 - ▣ Different granularities will produce different volumes of tuples...
- As system load varies, can react very easily by changing the granularity of data being used
 - ▣ Choose the finest granularity of input data that the current system load will allow

Stream-Processing Challenges (3)

27

- Tuples in a stream don't always arrive in order!
 - ▣ A tuple's timestamp can be set by database when the tuple arrives, or it can be an externally specified field
 - ▣ e.g. stock-sale records might already include timestamp
- Generally a characteristic of a specific stream...
 - ▣ Can specify a stream's slack: the maximum “out-of-order”-ness that will be allowed
 - ▣ Tuples that arrive later than the specified slack are ignored!

Stream-Processing Challenges (4)

28

- Truviso example:
 - ```
CREATE STREAM stock_sales (
 sale_time TIMESTAMP CQTIME
 USER GENERATED SLACK '1 MINUTE',
 ticker_symbol VARCHAR(10),
 num_shares INTEGER,
 price_per_share NUMERIC(7, 2)
);
```
- Problem: if tuples can arrive out of order, a query may have already generated invalid results...
  - ▣ Could simply delay outputting results by the specified amount of slack. Then we know results won't change...

# Stream-Processing Challenges (5)

29

- Another approach: output revisions to answers!
- Some input data streams can also include revisions to previous records
  - ▣ (commercial stock ticker feeds, for example)
  - ▣ e.g. correct invalid values, add extra rows, remove rows
- A few stream databases (e.g. Borealis) can handle revisions on data streams
- If all tuples affected by revision are still in memory:
  - ▣ Recompute affected results incrementally, using old and new versions of dataset for the affected timestamp(s)
  - ▣ Only output new records that actually changed
  - ▣ (Revised outputs may force other revisions to be made...)

# Stream-Processing Challenges (6)

30

- If tuples affected by revision no longer in memory:
  - ▣ Must scan and replay archived data-stream tuples to find all relevant tuples
- Also likely that query nodes' synopses will no longer contain data for the affected timestamps ☹️
  - ▣ Must recompute all work from scratch; can't use incremental updates
- As before, only issue revised output records for results that actually change
- If revisions are far enough in past, may cause *many* results to be recomputed and revised ☹️

# DSMS Summary

31

- Possible to use relational databases to continuously process data streams, but not very efficient!
- Relation-oriented data stream management systems (DSMS) blend together standard relational model databases with stream-processing capabilities
- Very powerful solution for problems involving data stream processing!

# References

32

- **STREAM: The Stanford Data Stream Management System** [Arasu et al]
  - ▣ Introduction to how the relational model is extended to include data stream processing
- **Models and Issues in Data Stream Systems** [Babcock et al]
  - ▣ Discussion of data models and query languages used in relational data stream management systems
- **The 8 Requirements of Real-Time Stream Processing** [Stonebraker, Çetintemel, Zdonik]
  - ▣ A good summary of the major requirements of DSMSes
- **Revision Processing in a Stream Processing Engine** [Ryvkina, Maskey, Cherniack, Zdonik]
  - ▣ Description of Borealis revision processing



# References (2)

33

- Processing Flows of Information: From Data Stream to Complex Event Processing [Cugola, Margara]
  - ▣ Excellent survey of many different data stream engines!
- Stanford STREAM research project:
  - ▣ <http://infolab.stanford.edu/stream/>
- Aurora and Borealis research projects:
  - ▣ <http://www.cs.brown.edu/research/aurora/>
  - ▣ <http://www.cs.brown.edu/research/borealis/>
- TelegraphCQ research project:
  - ▣ <http://telegraph.cs.berkeley.edu/>