

PASSWORDS TREES AND HIERARCHIES

Account Password Management

2

- Mentioned a retailer with an online website...
- Need a database to store user account details
 - ▣ Username, password, other information
- How to store a user's password?
- What if the database application's security is compromised?
 - ▣ Can an attacker get a list of all user passwords?
 - ▣ Can the DB administrator be trusted?
- **Do we actually need to store the original password??**

A Naïve Approach

3

- A simple solution:

- ▣ Store each password as plaintext

```
CREATE TABLE account (  
    username VARCHAR(20) PRIMARY KEY,  
    password VARCHAR(20) NOT NULL,  
    ...  
);
```

- Benefits:

- ▣ If user forgets their password, we can email it to them

- Drawbacks:

- ▣ Email is unencrypted – passwords can be acquired by eavesdropping
 - ▣ Users tend to use the same password for many different accounts
 - ▣ If database security is compromised, attacker gets all users' passwords
 - ▣ Of course, an unreliable administrator can also take advantage of this

Hashed Passwords

4

- A safer approach is to hash user passwords
 - ▣ Store hashed password, not the original
 - ▣ For authentication check:
 1. User enters password
 2. Database application hashes the password
 3. If hash matches value stored in DB, authentication succeeds
- Example using MD5 hash:

```
CREATE TABLE account (  
    username VARCHAR(20) PRIMARY KEY,  
    pw_hash CHAR(32) NOT NULL,  
    ...  
);
```

 - ▣ To store a password:

```
UPDATE account SET pw_hash = md5('new password')  
WHERE username = 'joebob';
```

Hashed Passwords (2)

5

- Want a cryptographically secure hash function:
 - ▣ Easy to compute a hash value from the input text
 - ▣ Even small changes in input text result in very large changes in the hash value
 - ▣ Hard to get a specific hash value by choosing input carefully
 - ▣ Should be collision resistant: hard to find two different messages that generate the same hash function
- MD5 is not collision resistant ☹️
 - ▣ “[MD5] should be considered cryptographically broken and unsuitable for further use.” – US-CERT
- SHA-1 was also discovered to not be very good
- Most people use SHA-2/3 hash algorithms now

Hashed Passwords (3)

6

- Benefits:
 - ▣ Passwords aren't stored in plaintext anymore
- Drawbacks:
 - ▣ Handling forgotten passwords is a bit trickier
 - Need alternate authentication mechanism for users
 - ▣ Isn't entirely secure! Still prone to dictionary attacks.
- Attacker computes a dictionary of common passwords, and each password's hash value
 - ▣ If attacker gets the hash values from the database, can crack some subset of accounts

Hashed, Salted Passwords

7

- Solution: salt passwords before hashing

- Example:

```
CREATE TABLE account (  
    username VARCHAR(20) PRIMARY KEY,  
    pw_hash CHAR(32) NOT NULL,  
    pw_salt CHAR(6) NOT NULL,  
    ...  
);
```

- ▣ Each account is assigned a random salt value
 - Salt is always a specific length, e.g. 6 to 16 characters
- ▣ Concatenate plaintext password with salt, before hashing
- ▣ Attacker would have to compute a dictionary of hashes for each salt value... Prohibitively expensive!

Password Management


8

- Basically no reason to store passwords in plaintext!!
 - ▣ Users almost always use the same passwords in multiple places!
 - ▣ Only acceptable in the simplest circumstances
 - ▣ (You don't want to end up on the news because your system got hacked and millions of passwords leaked...)
- Almost always want to employ a secure password storage mechanism
 - ▣ Hashing is insufficient! Still need to protect against dictionary attacks by applying salt
 - ▣ Also need a good way to handle users that forget their passwords

Trees and Hierarchies

9


- Many DB schemas need to represent trees or hierarchies of some sort
- Example: parts-explosion diagrams
 - ▣ “How are parts and subsystems assembled?”
 - ▣ “How much does a subsystem weigh?”
 - ▣ Other computations based on parts in subsystems



Trees and Hierarchies (2)

10

- Many kinds of relationships between people
 - ▣ Employee/manager relationships within an organization
 - ▣ Social graph data from a social network
 - ▣ Family trees of people, or pedigree data for animals



Trees and Hierarchies (3)

11

- Most common way of representing trees in the DB is an adjacency list model
 - ▣ Each node in the hierarchy specifies its parent node
 - ▣ Can represent arbitrary tree depths
- Example: employee database
 - ▣ *employee*(*emp_name*, *address*, *salary*)
 - ▣ *manages*(*emp_name*, *manager_name*)
 - ▣ Both attributes of *manages* are foreign keys referencing *employee* relation

Trees and Hierarchies (4)

12

- Adjacency list model is only one of several ways to represent trees and hierarchies
- Different approaches have different strengths and weaknesses
- Some approaches to consider:
 - ▣ Adjacency list models
 - ▣ Nested set models
 - ▣ Path enumeration models

General Design Questions

13

- How hard is it to access the tree?
 - ▣ Retrieve a specific node
 - ▣ Find the parent/children/siblings of a particular node
 - ▣ Retrieve all leaf nodes in the tree
 - ▣ Retrieve all nodes at a particular level in the tree
 - ▣ Retrieve a node and its entire subtree
- Also, path-based queries:
 - ▣ Retrieve a node corresponding to a particular path from the root
 - ▣ Retrieve nodes matching a path containing wildcards
 - ▣ Is a particular path in the hierarchy?

General Design Questions (2)




14

- How hard is it to modify the tree?
 - ▣ Add a single node
 - ▣ Delete a leaf node
 - ▣ Delete a non-leaf node
 - What to do with subtree of deleted node?
 - ▣ Move a subtree within the tree
- How to enforce constraints in the schema?
 - ▣ Enforce only one root
 - ▣ Disallow cycles in a tree
 - Simplest example: “my parent can’t be myself”
 - ▣ Disallow multiple parents (tree vs. directed acyclic graph)
 - ▣ Enforce a maximum child-count, maximum depth, etc.

Deleting Nodes

15

- What happens when a non-leaf node is deleted?
- Option 1:
 - ▣ Delete entire subtree as well
- Option 2:
 - ▣ Promote one child node into removed node's position
 - ▣ (Or, replace with placeholder)
- Option 3:
 - ▣ Make all child nodes into children of deleted node's parent



Adjacency List Model

16

- Very common approach for modeling trees
- Each node specifies its parent node
- Relationship between nodes are frequently stored in a separate table
 - ▣ e.g. *employee* and *manages*
 - ▣ Can represent multiple trees without *null* values
 - ▣ Can have employees that are not part of the hierarchy

Adjacency List Model (2)

17

□ Strengths:

- ▣ A very flexible model for representing and manipulating trees
- ▣ Easy to add a new node anywhere in the tree
- ▣ Easy to move a whole subtree

□ Weaknesses:

- ▣ Deleting a node often requires extra steps (to relocate children)
- ▣ Operations on entire subtrees are expensive
- ▣ Operations applied to a particular level across a tree are expensive
- ▣ Looking for a node at a specific path is expensive
- ▣ (All of these operations require multiple queries to identify the nodes affected by the operation.)

Queries on Subtrees

18

- Example query using subtrees:
 - ▣ For every manager reporting to a given Vice President, find the total number of employees under that manager, and their total salaries
 - ▣ Computing an aggregate over individual subtrees
- Another example:
 - ▣ Database containing parts in a mechanical assembly
 - Parts combined into sub-assemblies
 - Sub-assemblies and parts combined into larger assemblies
 - Top level assembly is the entire system
 - ▣ Find number of parts, the total cost, and the total weight of each sub-assembly in the system

Finding Nodes in Subtrees

19

- To find all nodes in a specific subtree, must iterate a query using a temporary table
 - ▣ Example: Find all employees under manager “Jones”

```
CREATE TEMPORARY TABLE emps
  ( emp_name VARCHAR(20) NOT NULL );

INSERT INTO emps VALUES ('Jones') ;
INSERT INTO emps SELECT emp_name FROM manages
  WHERE manager_name IN (SELECT * FROM emps)
    AND emp_name NOT IN (SELECT * FROM emps) ;
```
 - ▣ Iterate last statement until no new rows added to temp table
 - ▣ Databases often report a count of how many rows are inserted/modified/deleted by each DML operation

Finding Nodes in Subtrees (2)

20

- Can also store a “depth” value in the result

```
CREATE TEMPORARY TABLE emps (  
    emp_name VARCHAR(20) NOT NULL,  
    depth     INTEGER     NOT NULL  
);
```

```
INSERT INTO emps VALUES ('Jones', 1);
```

```
-- Some variable i, where i = 1 initially:
```

```
INSERT INTO emps SELECT emp_name, i + 1 FROM manages  
    WHERE manager_name IN  
        (SELECT * FROM emps WHERE depth = i);
```

- Each level of the hierarchy has the same depth value
- Slightly more efficient than the previous version
 - Each iteration has fewer rows to consider

Finding Nodes in Subtrees (3)

21

- Best to implement this as a stored procedure
 - ▣ Don't involve command/response round-trips with application code, if possible!
 - ▣ Perform processing entirely within the DB for best performance
- Still acceptable if application has to drive the iteration process
 - ▣ Most DB connectivity APIs let you create temporary tables, get number of rows changed, etc.

Other Adjacency List Notes

22

- Must manually order siblings under a node
 - ▣ Add another column to the table for ordering siblings
- Adjacency list model is also good for representing graphs
 - ▣ Actually easier than using for trees, because less constraints are required
 - ▣ Traversing the graph requires temporary tables and iterative stored procedures
 - ▣ (Other representations we will discuss today aren't well-suited for graphs at all!)

Aside: Recursive Queries

23

- Briefly looked at this Datalog query last time:
 - ▣ Manager relation:
manager(employee_name, manager_name)
 - ▣ Define view relation “empl” for all employees directly or indirectly managed by a particular manager.
empl(X, Y) :- manager(X, Y)
empl(X, Y) :- manager(X, Z), empl(Z, Y)
 - ▣ To find all employees managed by Jones:
? empl(X, “Jones”)
 - ▣ Datalog iterates on the rule definitions until the result doesn’t change (fixpoint)
- Datalog supports traversing hierarchies very easily

Recursive Queries in SQL

24

- For a long time, SQL did not support recursive queries
 - ▣ Only way to traverse adjacency-list data model was to use a temporary table and repeated queries, as described
- Now, most databases also support some form of recursive SQL query
 - ▣ e.g. PostgreSQL 8.4+, SQLServer, Oracle (not MySQL ☹)
 - ▣ If available, makes it much easier to traverse adjacency-list datasets
- Still requires repeated queries against the database...
 - ▣ Even though query is easier to write, it's still slow to execute!

Recursive Queries in PostgreSQL

25

- Example using SQL99/Postgres 8.4 syntax
- Find all employees directly or indirectly managed by Jones:

```
WITH RECURSIVE empl AS (Initial Subquery
    SELECT employee_name, manager_name
    FROM manager
    UNION ALLRecursive Subquery
    SELECT e.employee_name, m.manager_name
    FROM empl AS e JOIN manager AS m
        ON e.manager_name = m.employee_name)
SELECT * FROM empl
WHERE manager_name = 'Jones';
```

Recursive Queries in PostgreSQL (2)

26

- Can compute the depth in the hierarchy, too:

```
WITH RECURSIVE empl AS (  
    SELECT employee_name, manager_name,  
           1 as depth  
    FROM manager  
    UNION ALL  
    SELECT e.employee_name, m.manager_name,  
           e.depth + 1 AS depth  
    FROM empl AS e JOIN manager AS m  
         ON e.manager_name = m.employee_name)  
SELECT * FROM empl  
WHERE manager_name = 'Jones';
```

SQLServer Recursive Queries

27

- Microsoft SQLServer 2005/2008 also uses **WITH** clause for recursive queries
 - ▣ Similar approach, using a non-recursive subquery and a recursive subquery, combined with **UNION [ALL]**
 - ▣ Doesn't use a **RECURSIVE** modifier
- Neither Postgres nor SQLServer recursive queries can handle cycles in the data...
 - ▣ Introducing a cycle into the data causes the query to infinite-loop

Oracle Recursive Queries

28

- Oracle has much more sophisticated recursive query support
- A simple example:

```
SELECT employee_name, manager_name, LEVEL
FROM manager
CONNECT BY PRIOR employee_name = manager_name;
```

 - ▣ **PRIOR** modifier specifies parent in parent/child relationship
 - ▣ **LEVEL** is a pseudocolumn specifying level in the hierarchy
- Can also do *many* other things, such as:
 - ▣ Specify the initial data-set to iterate on
 - ▣ Specify the order of siblings in each level of hierarchy
 - ▣ Detect and report which nodes are leaves in the hierarchy
 - ▣ Detect and report cycles in the dataset
 - ▣ Generate the full path from root to each node in the result

Final Notes on Recursive SQL Queries

29

- As stated earlier:
 - ▣ Recursive SQL queries make it much easier to write the query that traverses adjacency-list data...
 - ▣ Still requires the DB engine to repeatedly issue queries until the entire hierarchy is traversed!
- If an application needs to store hierarchical or graph data, and must select entire subtrees quickly:
 - ▣ Adjacency-list model is *the most expensive* model to use!

Nested Set Model

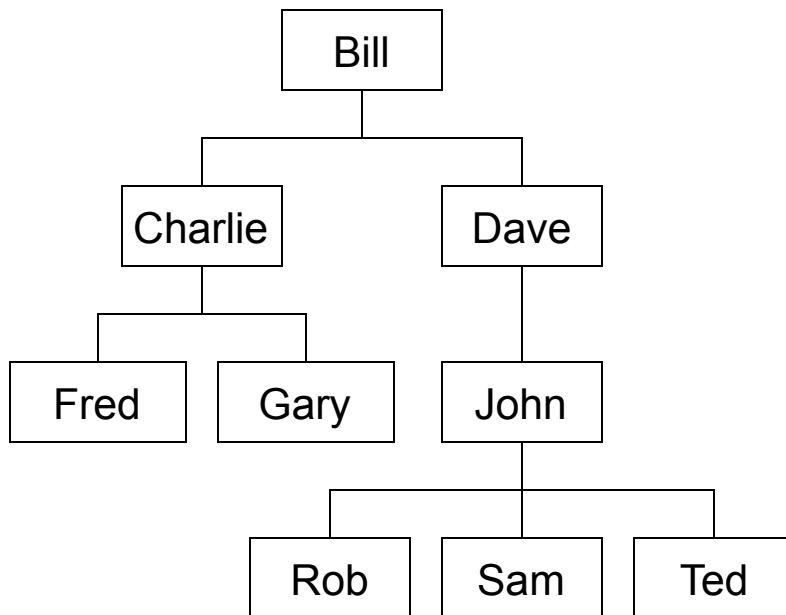
30

- A model optimized for selecting subtrees
- Represents hierarchy using containment
 - ▣ A node “contains” its children
- Give each node a “low” and “high” value
 - ▣ Specifies a range
 - ▣ Always: $\text{low} < \text{high}$
- Use this to represent trees:
 - ▣ A parent node's [low, high] range contains the ranges of all child nodes
 - ▣ Sibling nodes have non-overlapping ranges


Example using Nested Sets

31

□ Manager hierarchy:



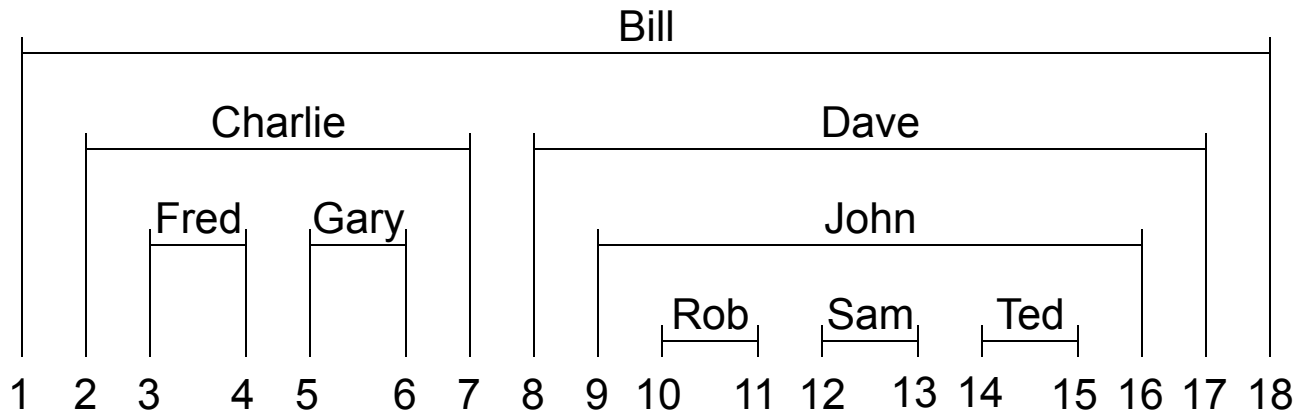
□ Nested set hierarchy:



Selecting Subtrees

32

- Each parent node contains its children:



- Easy to select an entire subtree
 - ▣ Select all nodes with low (or high) value within node's range
- Can also select all leaf nodes [relatively] easily
 - ▣ If all range values separated by 1, find nodes with $\text{low} + 1 = \text{high}$
 - ▣ For arbitrary range sizes, find nodes that contain no other node's range values (requires self-join)

Nested Set Model

33

□ Strengths:

- ▣ Very easy to select a whole subtree
 - Very important for generating reports against subtrees
- ▣ Tracking the order of siblings is built in

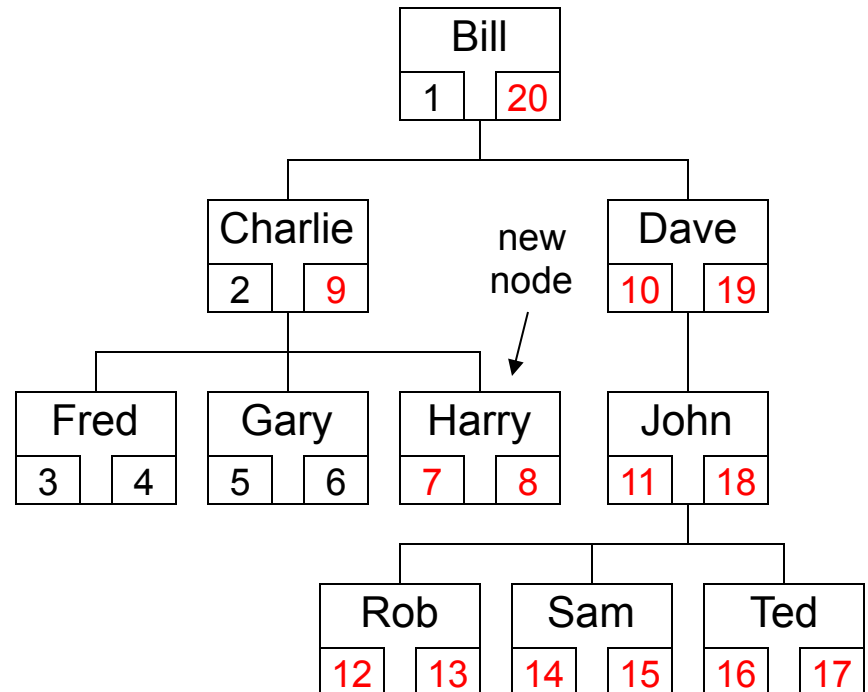
□ Weaknesses:

- ▣ Some selections are more expensive
 - e.g. finding all leaf nodes in tree requires self-join
- ▣ Constraints on range values are expensive
 - **CHECK** constraint to ensure $low < high$ is cheap...
 - **CHECK** constraint to verify other interval characteristics is expensive!
- ▣ Pretty costly to insert nodes or restructure trees
 - Have to update node bounds properly

Adding Nodes In Nested Set Model

34

- If adding a node:
 - ▣ Must choose range values for new node correctly
 - ▣ Often need to update range values of many other nodes, even for simple updates
- Example:
 - ▣ Add new employee Harry under Charlie
 - ▣ Must update ranges of most nodes in the tree!



Supporting Faster Additions

35

- Can separate range values by larger amounts
 - ▣ e.g. spacing of 100 instead of 1
 - ▣ Situations requiring range-adjustments of many nodes will be far less frequent
- Can implement tree-manipulation operations as stored procedures
 - ▣ “Add a node,” “move a subtree,” etc.

Path Enumeration Models

36

- For each node in hierarchy:
 - ▣ Represent node's position in hierarchy as the path from the root to that node
 - ▣ Entire path is stored as a single string value
- Node enumeration:
 - ▣ Each node has some unique ID or name
 - ▣ A path contains the IDs of all nodes between root and a particular node
 - ▣ If ID values are fixed size, don't need a delimiter
 - ▣ If ID values are variable size, choose a delimiter that won't appear in node IDs or names

Path Enumeration Model (2)

37

- Path enumeration model is fastest when nodes are retrieved using full path
 - ▣ “Is a specified node in the hierarchy?”
 - ▣ “What are the details for a specified node?”
 - ▣ Adjacency list model and nested set model simply can’t answer these queries quickly!
- Example:
 - ▣ A database-backed directory service (e.g. LDAP or Microsoft Active Directory)
 - ▣ Objects and properties form a hierarchy
 - ▣ Properties are accessed using full path names
 - “sales.printers.colorprint550.queue”

Strengths and Weaknesses

38

- Optimized for read performance
 - ▣ Retrieving a specific node using its path is *very* fast
 - ▣ Retrieving an entire subtree is also pretty fast
 - Requires text pattern-matching, but matching a prefix is fast, especially with a suitable index on the string
 - Example: Find all sales print servers
 - Use a condition: `path LIKE 'sales.printers.%'`
- Adding leaf nodes is fast
- Restructuring a tree can be very slow
 - ▣ Have to reconstruct many paths...
- Operations rely on string concatenation and string manipulation

Edge Enumeration

39

- Paths can enumerate edges instead of nodes
 - ▣ Each level of path specifies index of node to select
- Primary method used in books

▣ Example:

Edge Enumeration	Section Name
3	SQL
3.1	Background
3.2	Data Definition
3.2.1	Basic Domain Types
3.2.2	Basic Schema Definition in SQL
...	...

- Like node enumeration, requires string manipulation for most operations

Summary: Trees and Hierarchies

40

- Can represent trees and hierarchies in several different ways
 - ▣ Adjacency list model
 - ▣ Nested set model
 - ▣ Path enumeration model
- Each approach has different strengths
 - ▣ Each is optimized for different kinds of usage
- When designing schemas that require hierarchy:
 - ▣ Consider functional and non-functional requirements
 - ▣ Choose a model that best suits the problem